

Advanced Data Structures and Algorithm Analysis

丁尧相
浙江大学

Fall & Winter 2025
Lecture 4

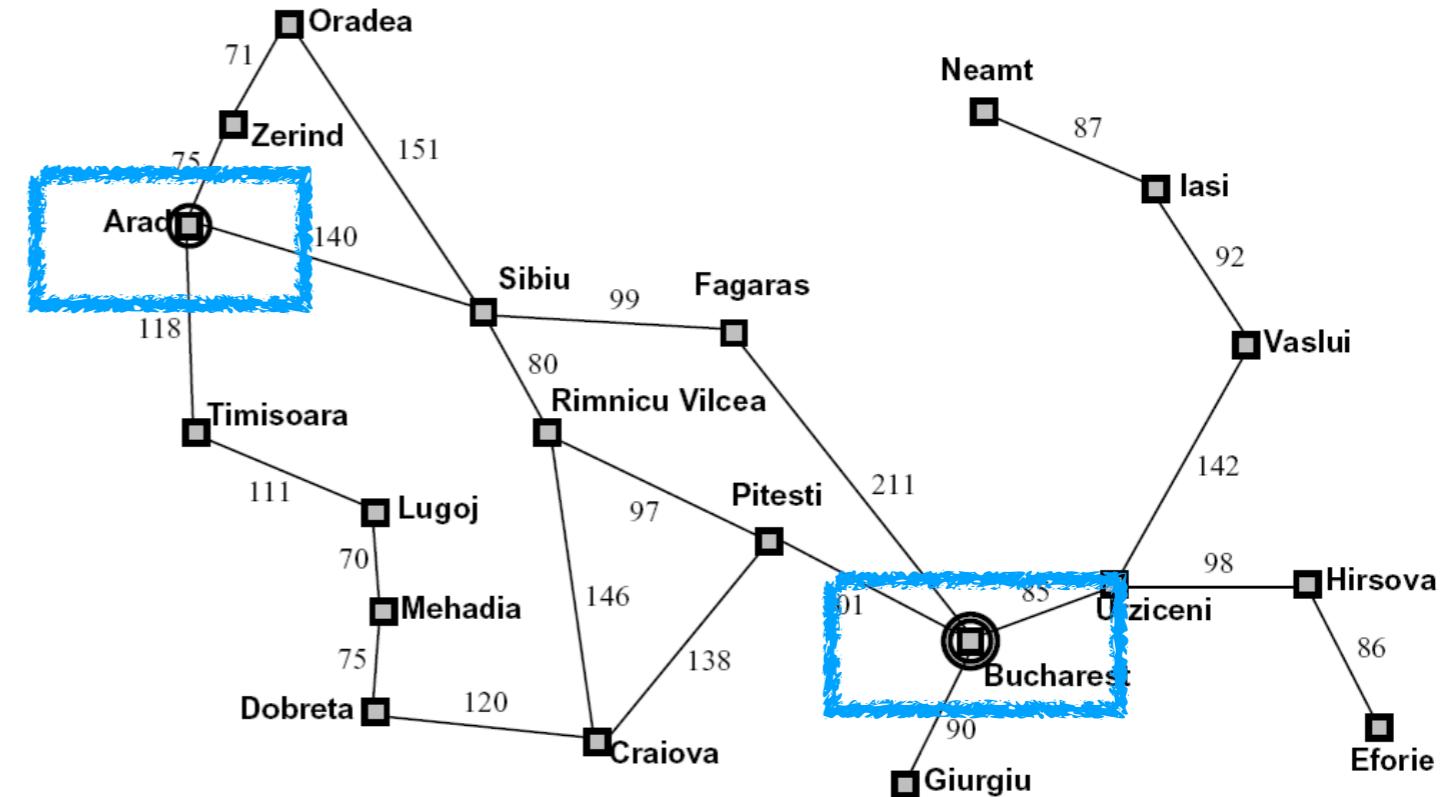
Search & Backtracking

- Fundamental search
- Search with backtracking
- Search by simulation
- Take-home messages

Search & Backtracking

- Fundamental search
- Search with backtracking
- Search by simulation
- Take-home messages

Search Problem

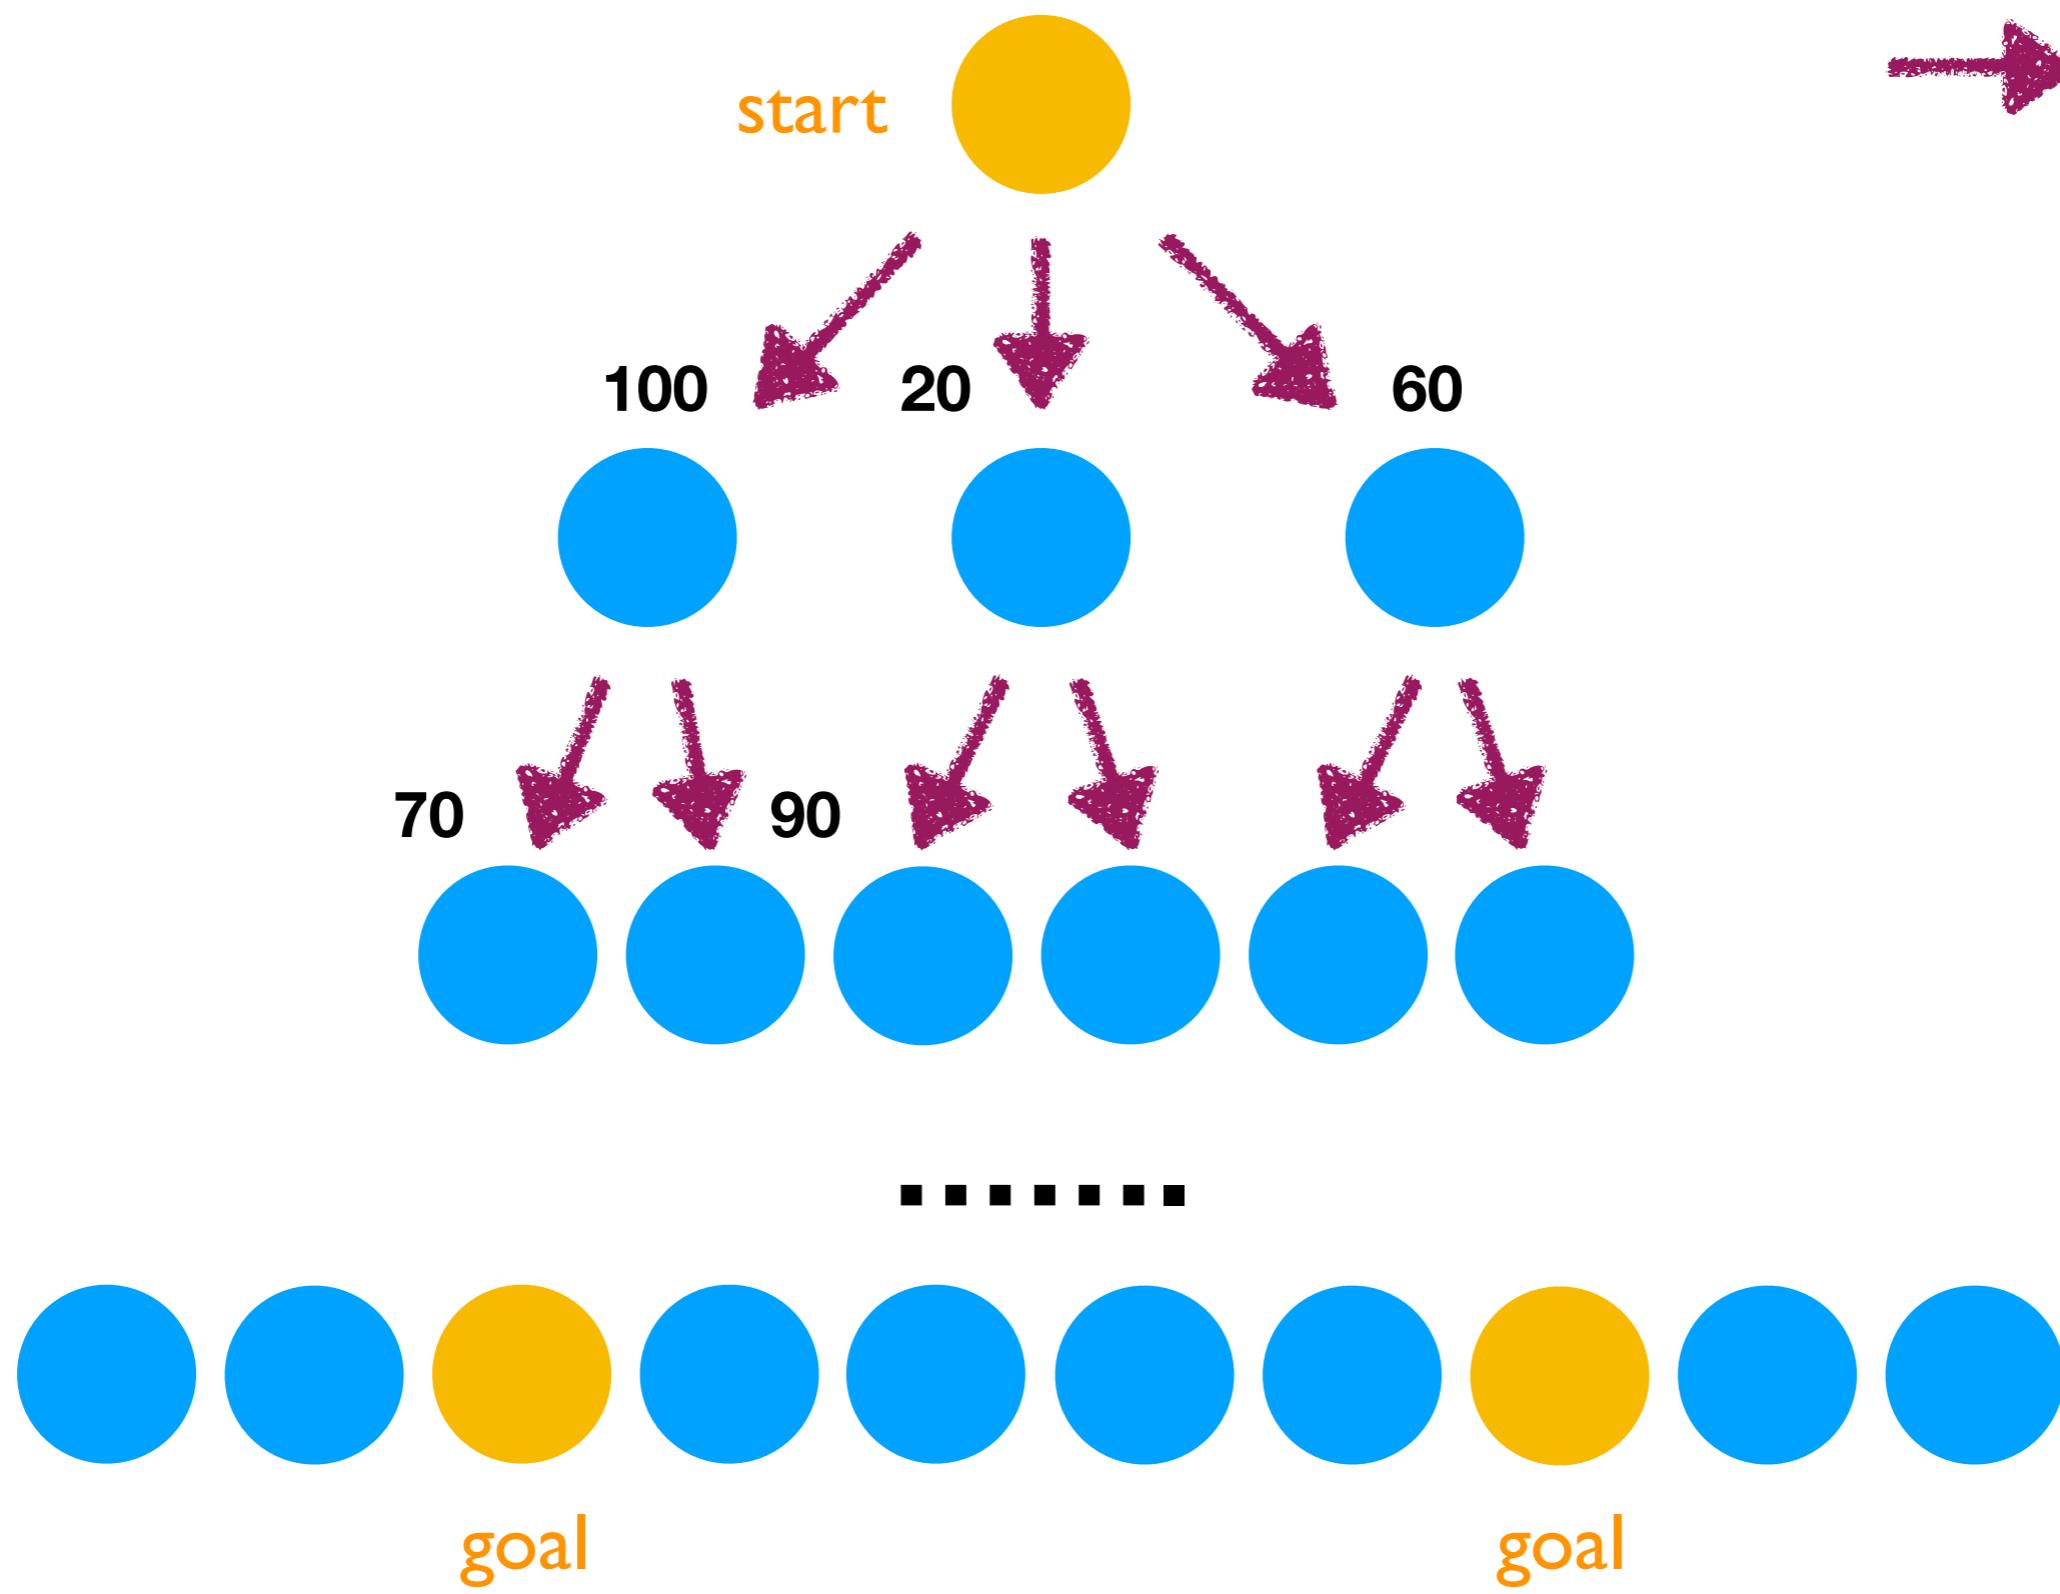


“Trees sprout up just about everywhere in computer science.”

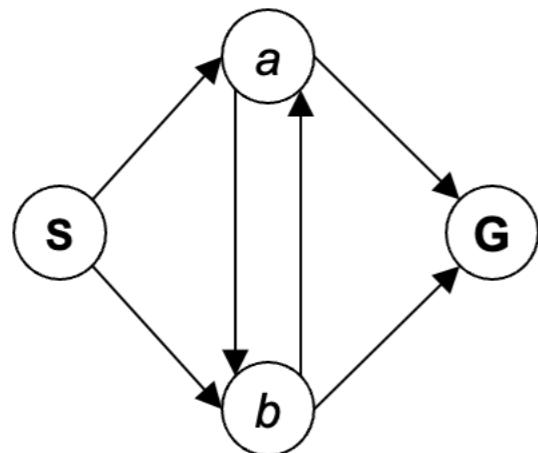
— Donald Knuth

Search Tree

state
choice



What is the Depth of the Search Tree for the Graph?



There may be a lot of repeated structures in a search tree.

Search Algorithms

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
    initialize the search tree using the initial state of problem
    loop do
        if there are no candidates for expansion then return failure
        choose a leaf node for expansion according to strategy
        if the node contains a goal state then return the corresponding solution
        else expand the node and add the resulting nodes to the search tree
    end
```

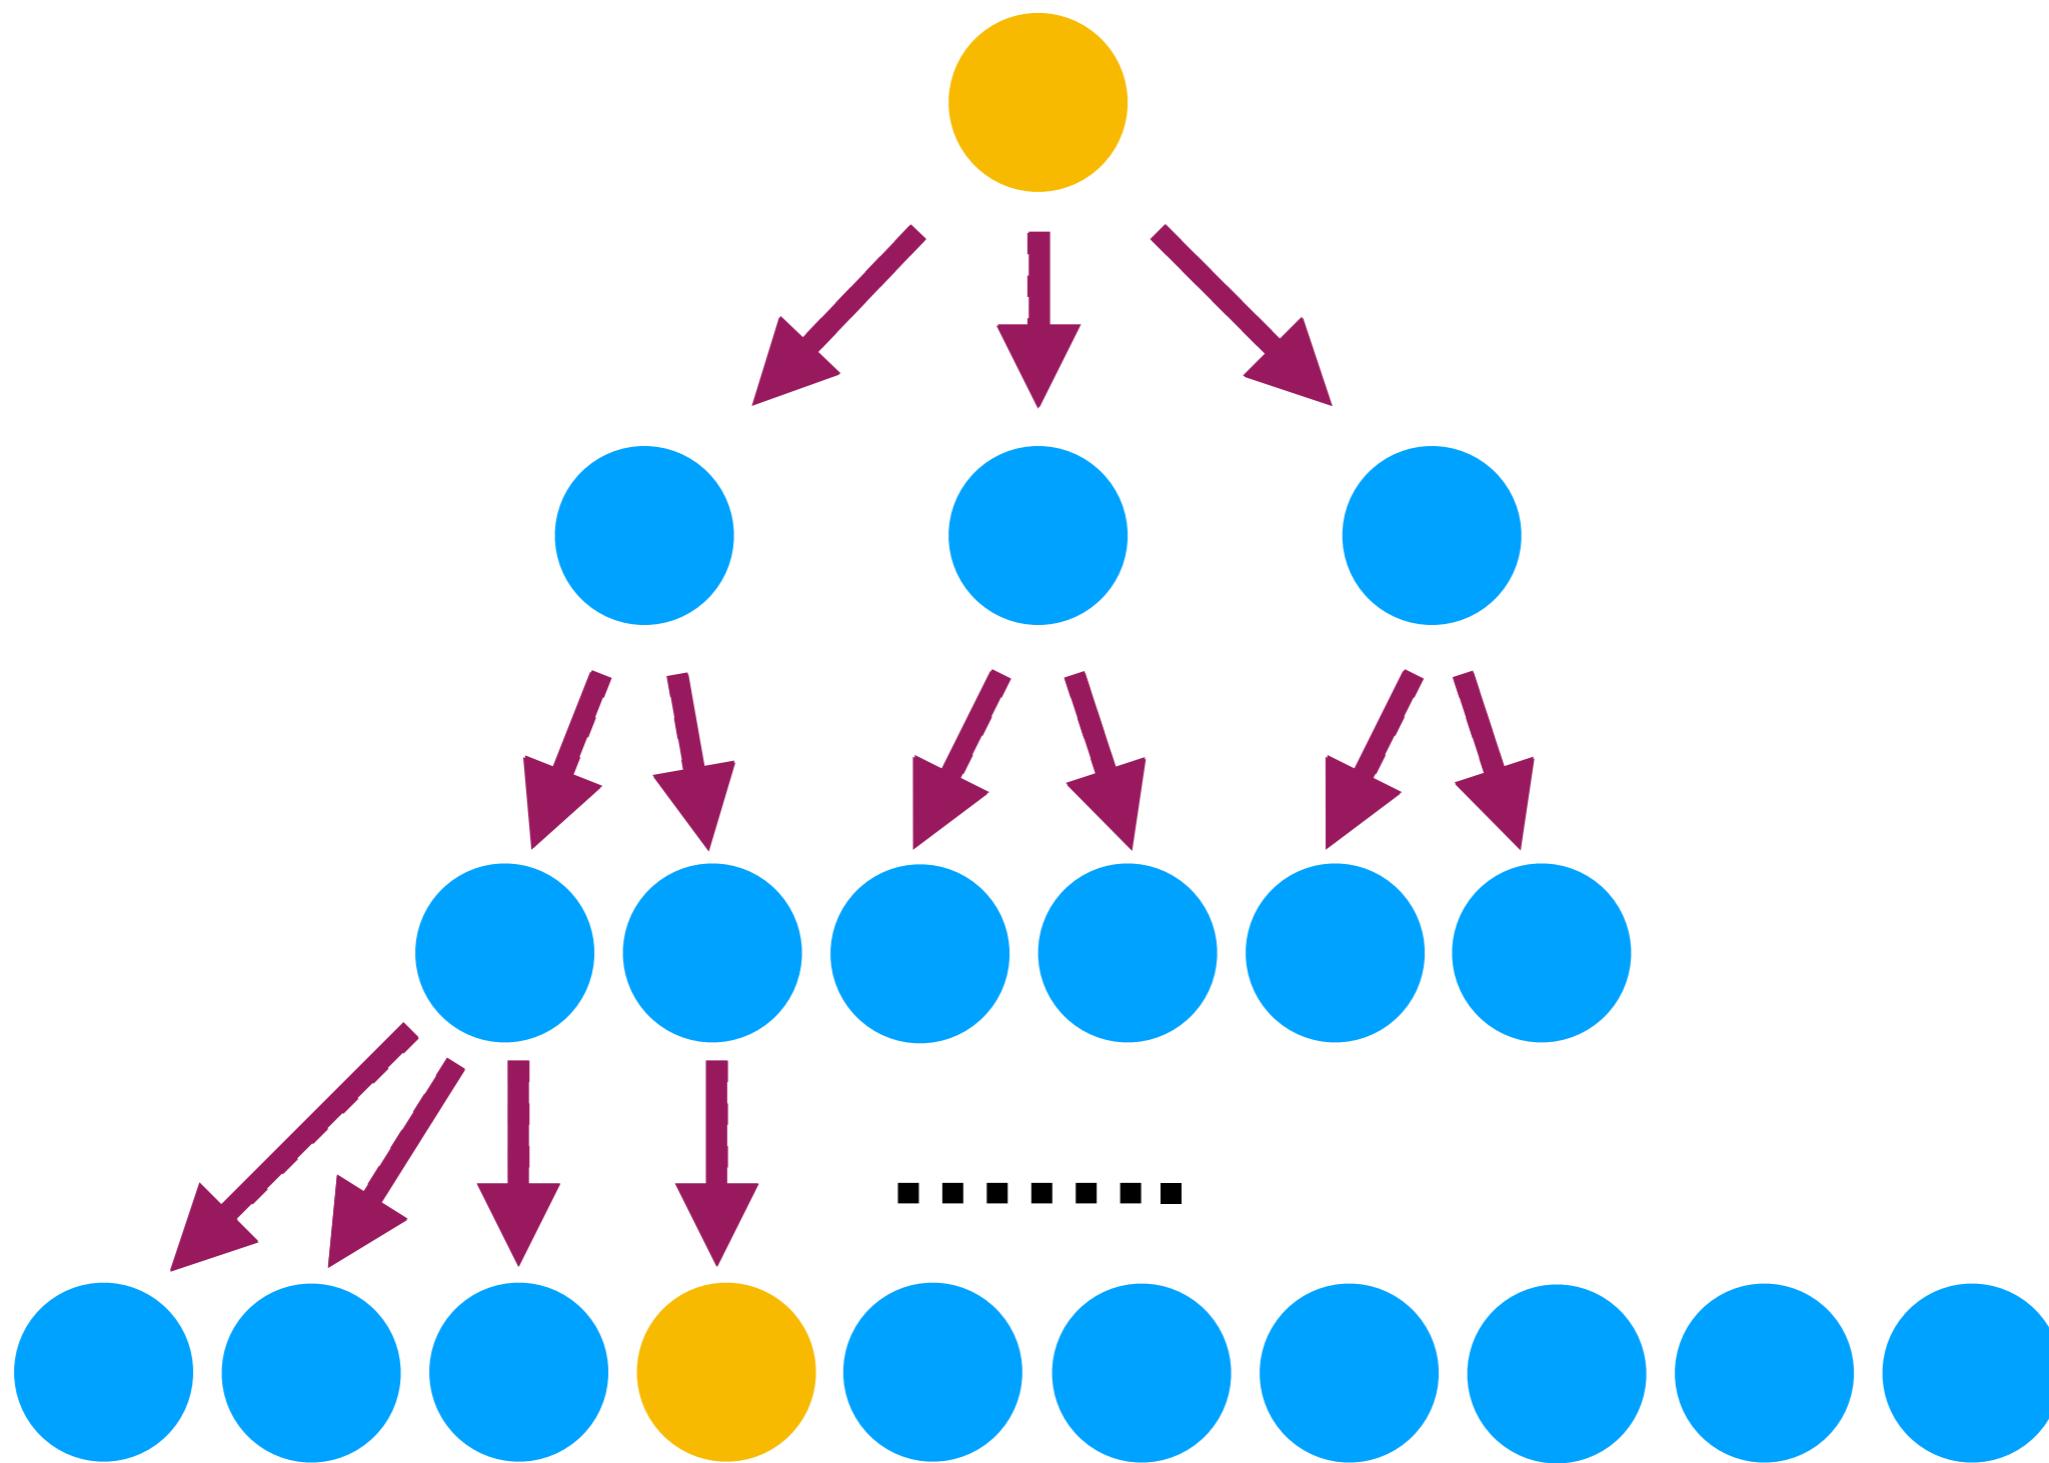
The strategy of choosing which leaf to expand is essential.

Breath-First Search



Notice: Temporally assume uniform cost for any action.

Breath-First Search (BFS)



Implemented with FIFO queue.

How Good is a Search Algorithm?

- Completeness:
 - Guarantee to find a solution if exists.
- Optimality:

caution: not about complexity

 - Guarantee to find the least cost (largest utility) path.
- Time complexity
- Space complexity

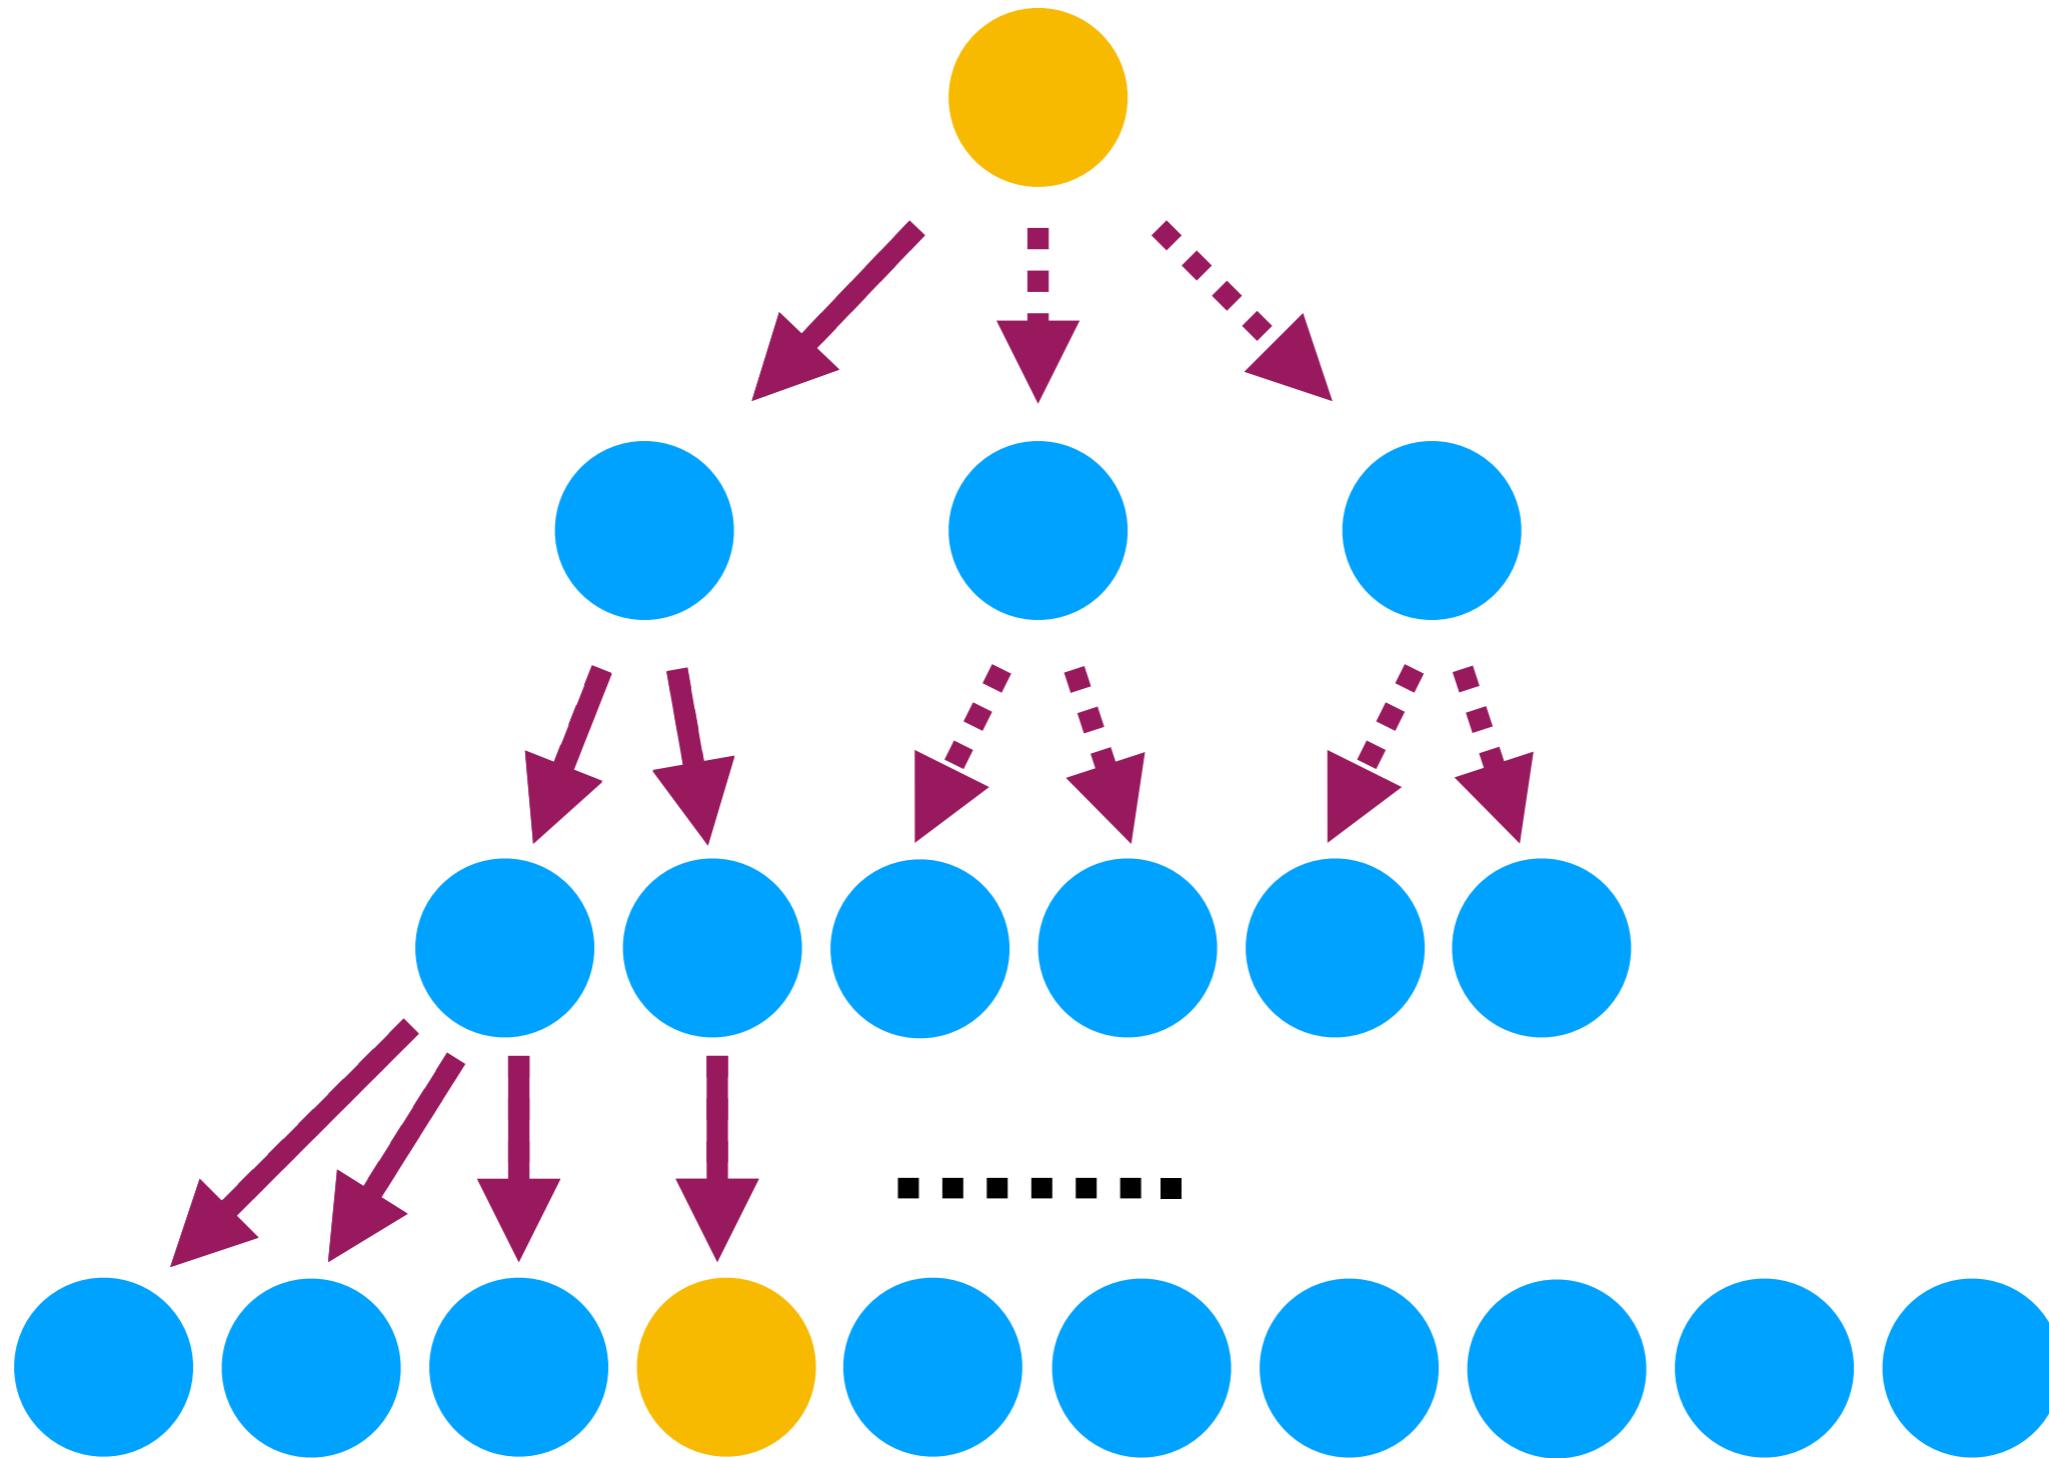
How Good is BFS?

- Completeness:
 - Yes
 - Optimality:
 - Yes when costs on all edges are equal
 - Time complexity
 - $O(b^d)$
 - Space complexity
 - $O(b^d)$
- b : the maximum #edge on a node
 d : the maximum depth of the search tree
- Difficult to improve if no more information about the search tree is given
- Big issue for real-world application!
Anyway to improve?

Depth-First Search



Depth-First Search (DFS)

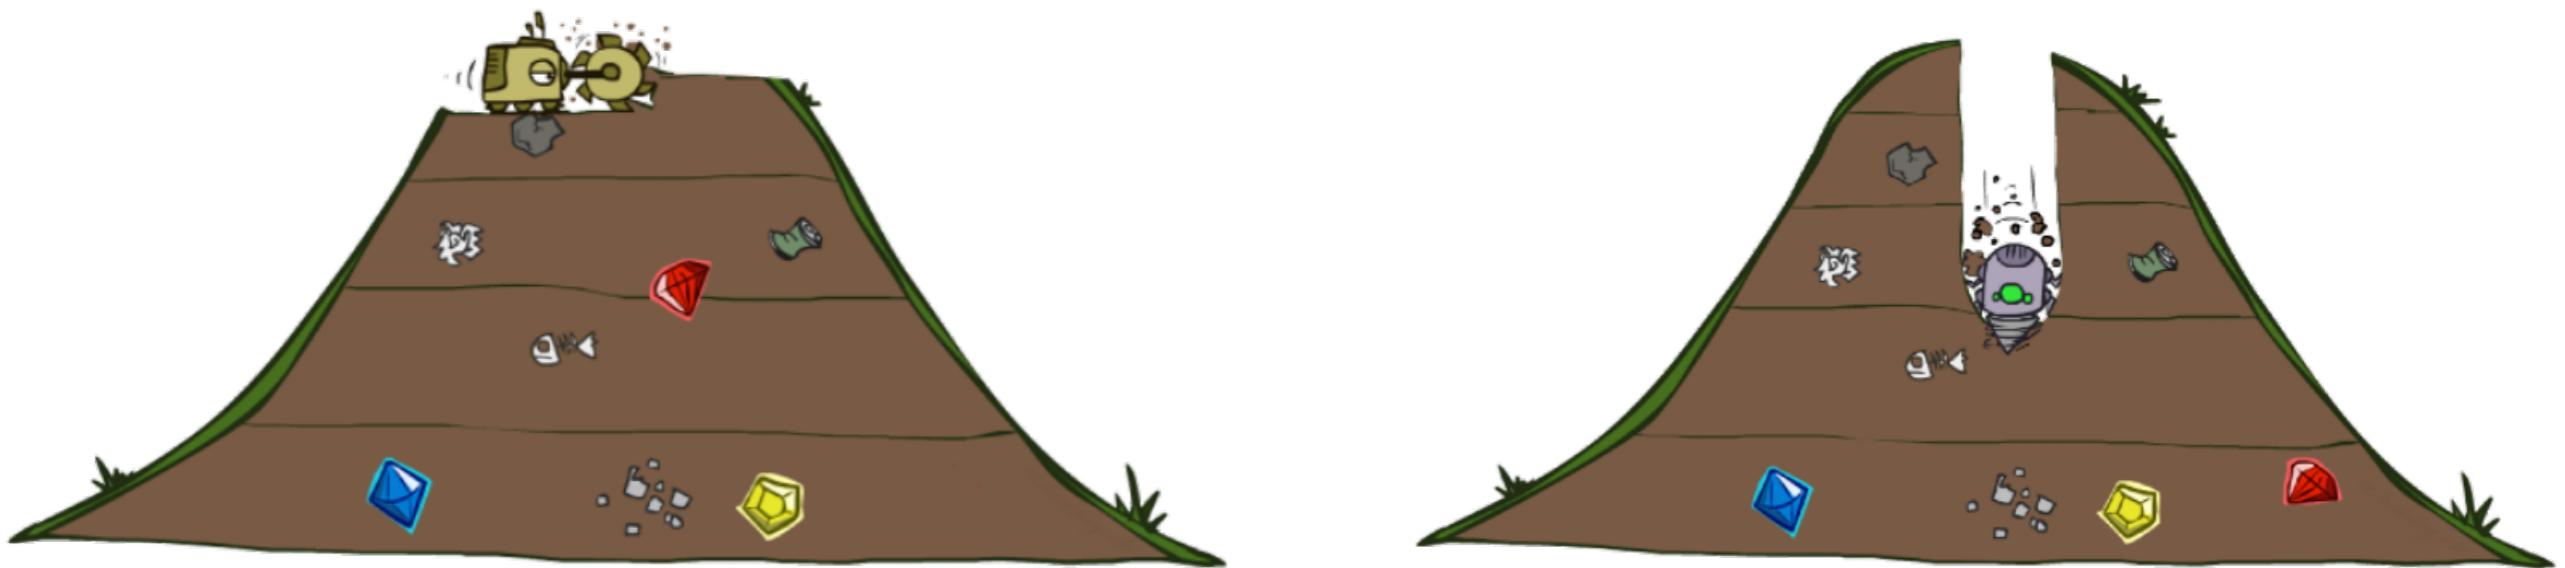


Implemented with LIFO queue (stack).

How Good is DFS?

- Completeness:
 - NO! (why?)
 - Optimality:
 - NO! (why?)
 - Time complexity
 - $O(b^d)$
 - Space complexity
 - $O(bd)$
- b : the maximum #edge on a node
 d : the maximum depth of the search tree
- Much more reasonable than BFS.
Making DFS more widely used in practice.

BFS vs. DFS



Can we get the benefits of both BFS and DFS?

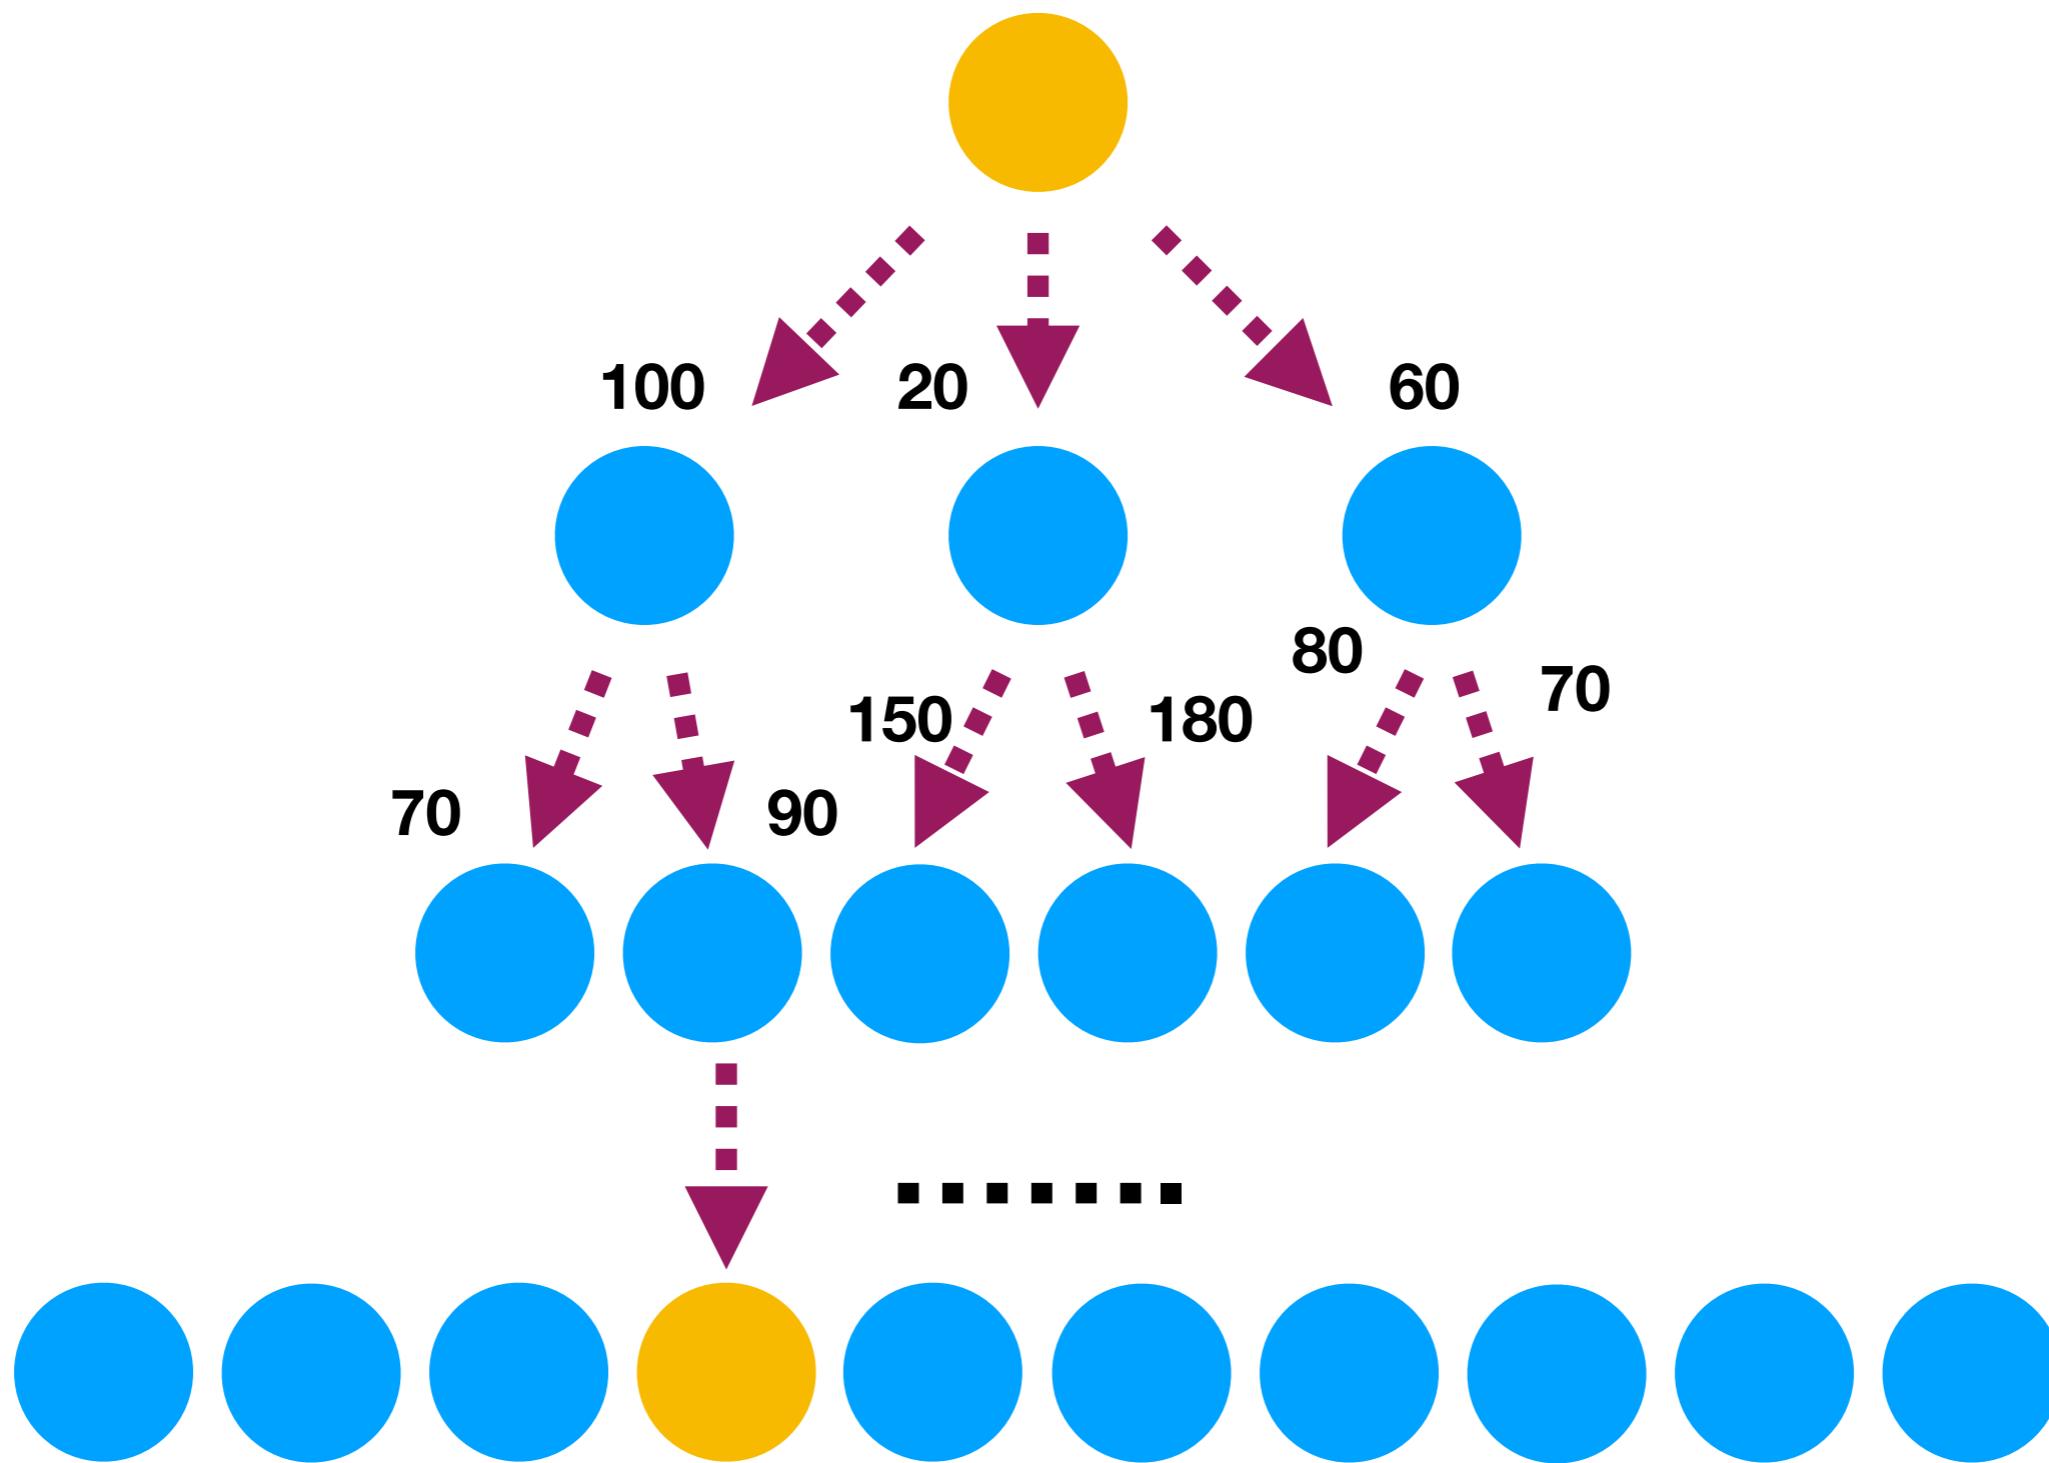
Iterative Deepening Search

- Run a DFS with depth limit 1. If no solution...
- Run a DFS with depth limit 2. If no solution...
- Run a DFS with depth limit 3. ...

Why can we do this? Aren't many nodes searched repeatedly?

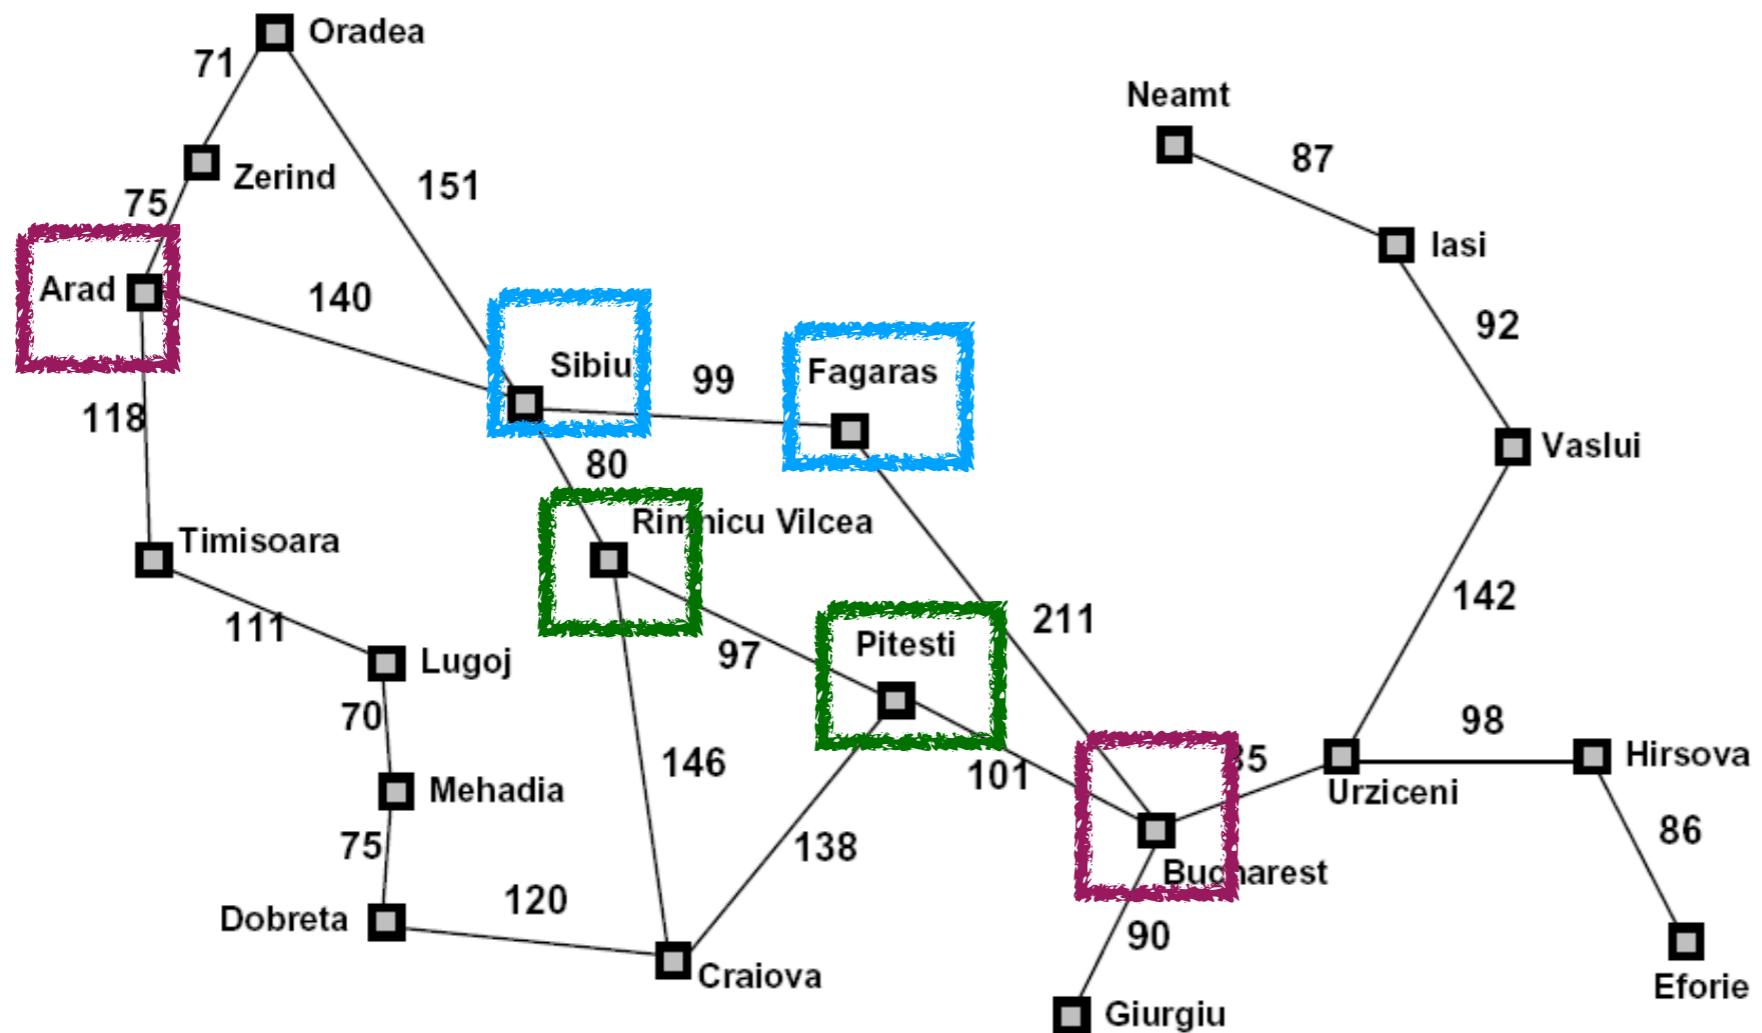
Actually, the time complexity is dominated by the latest DFS!

Cost-Sensitive Search



Neither BFS nor DFS works.

Best-First Search (Greedy Search)

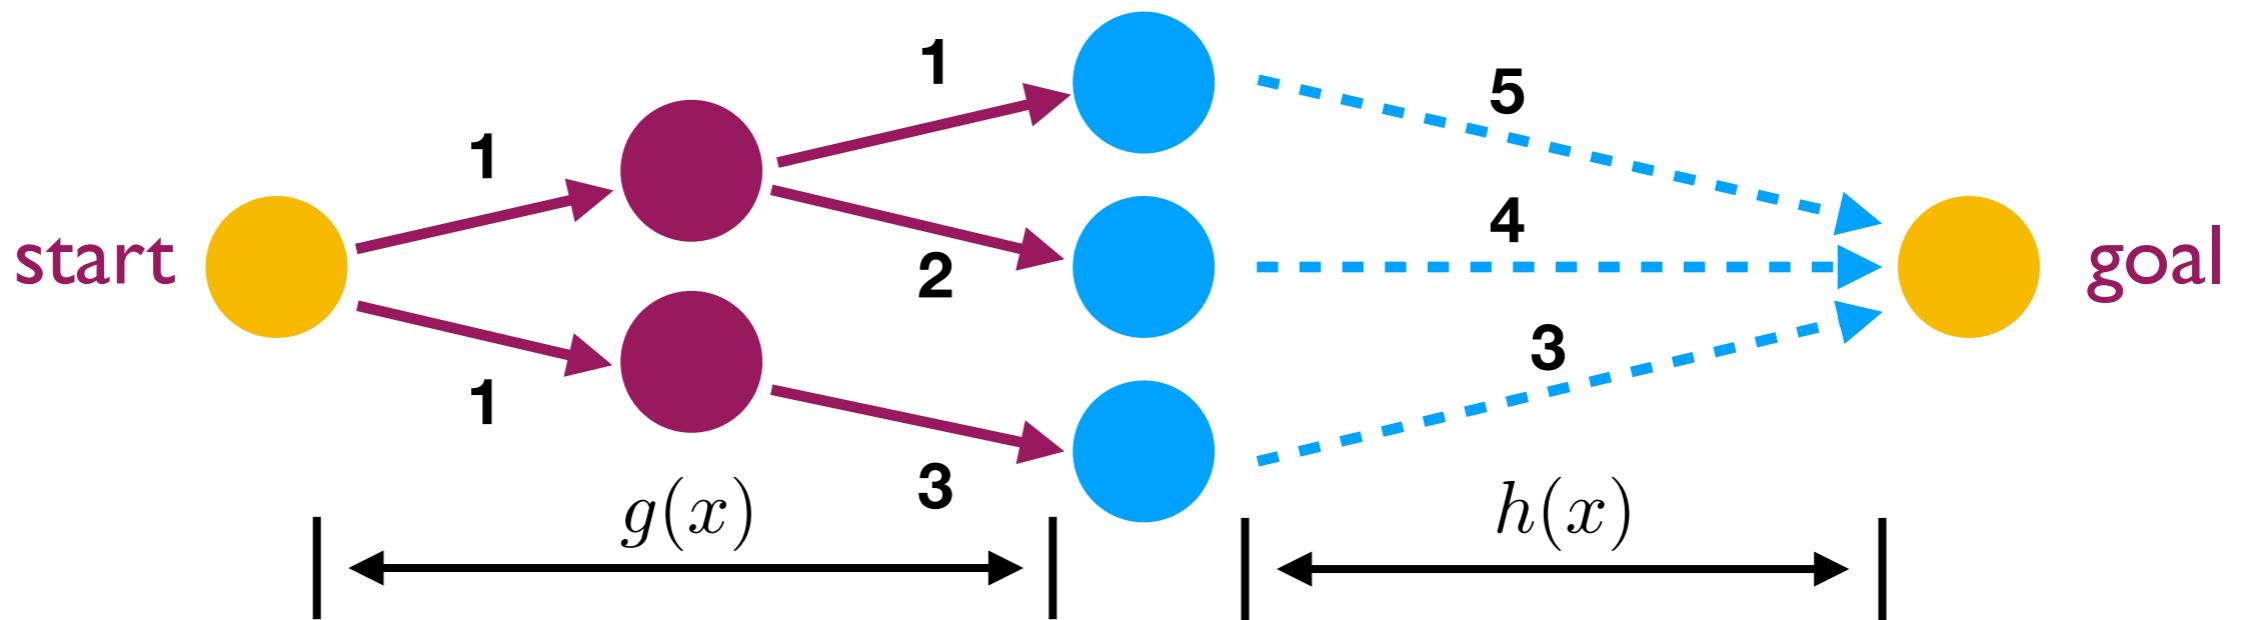


Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Always expand the node with the best heuristic function value

As usual, greedy strategy can lead to sub-optimal results.

A* Search



- A* chooses **one** x to expand with minimum $f(x) = g(x) + h(x)$
- $g(x)$: **known** cost from **the start to nodes**.
- $h(x)$: **heuristic** to **estimate** the cost from **nodes to the goal**.

A* guarantees to find the optimal solution when $h(x)$ never underestimates the true cost.

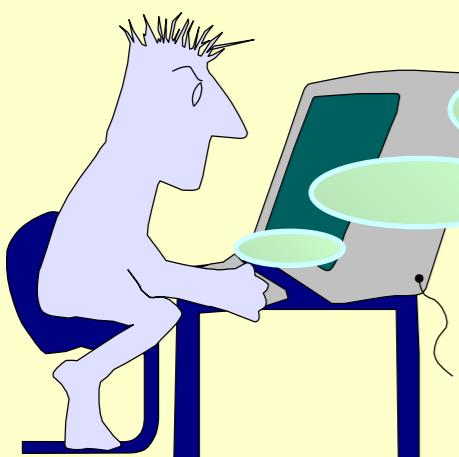
Search & Backtracking

- Fundamental search
- **Search with backtracking**
- Search by simulation
- Take-home messages

Rationale of the Backtracking Algorithms

A sure-fire way to find the answer to a problem is to **make a list of all candidate answers, examine each**, and following the examination of all or some of the candidates, declare the identified answer.

Suuuure — if the list is finite and it is possible to identify the answer following the examinations. **AND**, if there are not **too many** candidates.



Rationale of the Backtracking Algorithms

A sure-fire way to find the answer to a problem is to **make a list of all candidate answers, examine each**, and following the examination of all or some of the candidates, declare the identified answer.

Pruning in the search tree

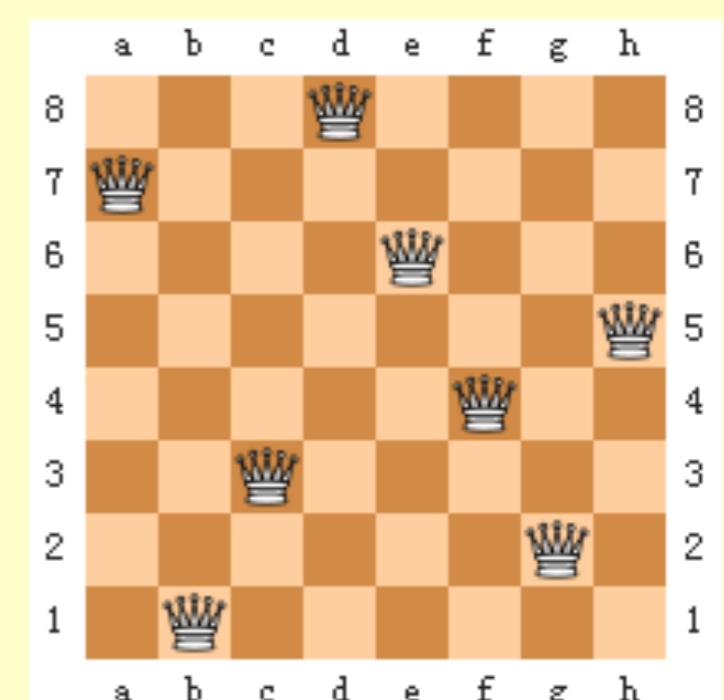
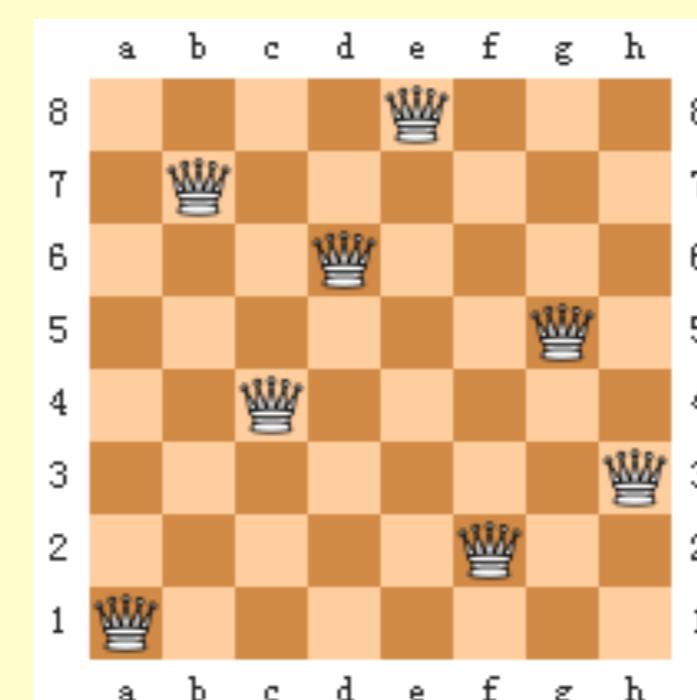
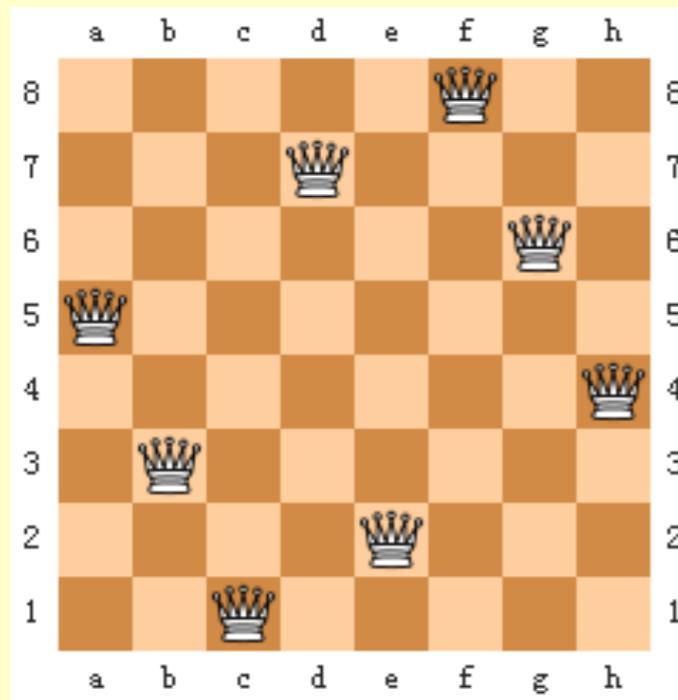
Backtracking enables us to **eliminate** the explicit examination of **a large subset** of the candidates while still guaranteeing that the answer will be found if the algorithm is run to termination.

The **basic idea** is that suppose we have a partial solution (x_1, \dots, x_i) where each $x_k \in S_k$ for $1 \leq k \leq i < n$. First we add $x_{i+1} \in S_{i+1}$ and check if $(x_1, \dots, x_i, x_{i+1})$ satisfies the constraints. If the answer is “yes” we **continue** to add the next x , else we delete x_i and **backtrack** to the previous partial solution (x_1, \dots, x_{i-1}) .

Eight Queens

Find a placement of **8 queens** on an 8×8 chessboard such that no two **queens attack**.

Two queens are said to **attack** iff they are in the same row, column, diagonal, or anti-diagonal of the chessboard.



1	2	3	4	5	6	7	8
1			Q				
					Q		
		Q					
						Q	
			Q				
							Q
8	7	6	5	4	3	2	1

$Q_i ::=$ queen in the i -th row

$x_i ::=$ the column index in which Q_i is

Solution = (x_1, x_2, \dots, x_8)

= (4, 6, 8, 2, 7, 1, 3, 5)

Constraints: ① $S_i = \{ 1, 2, 3, 4, 5, 6, 7, 8 \}$ for $1 \leq i \leq 8$

1	2	3	4	5	6	7	8
1			Q				
2					Q		
3		Q					
4						Q	
5	Q						
6			Q				
7					Q		
8							Q

$Q_i ::=$ queen in the i -th row

$x_i ::=$ the column index in which Q_i is

Solution = (x_1, x_2, \dots, x_8)

= (4, 6, 8, 2, 7, 1, 3, 5)

This implies 8^8 candidates
in the solution space.

Constraints: ① $S_i = \{ 1, 2, 3, 4, 5, 6, 7, 8 \}$ for $1 \leq i \leq 8$

	1	2	3	4	5	6	7	8
1				Q				
2					Q			
3						Q		
4		Q						
5							Q	
6								Q
7			Q					
8					Q			

$Q_i ::=$ queen in the i -th row

$x_i ::=$ the column index in which Q_i is

This implies
that the solution must
be a permutation of 1, 2, ..., 8.
Thus the number of candidates
in the solution space
is reduced to 8!.

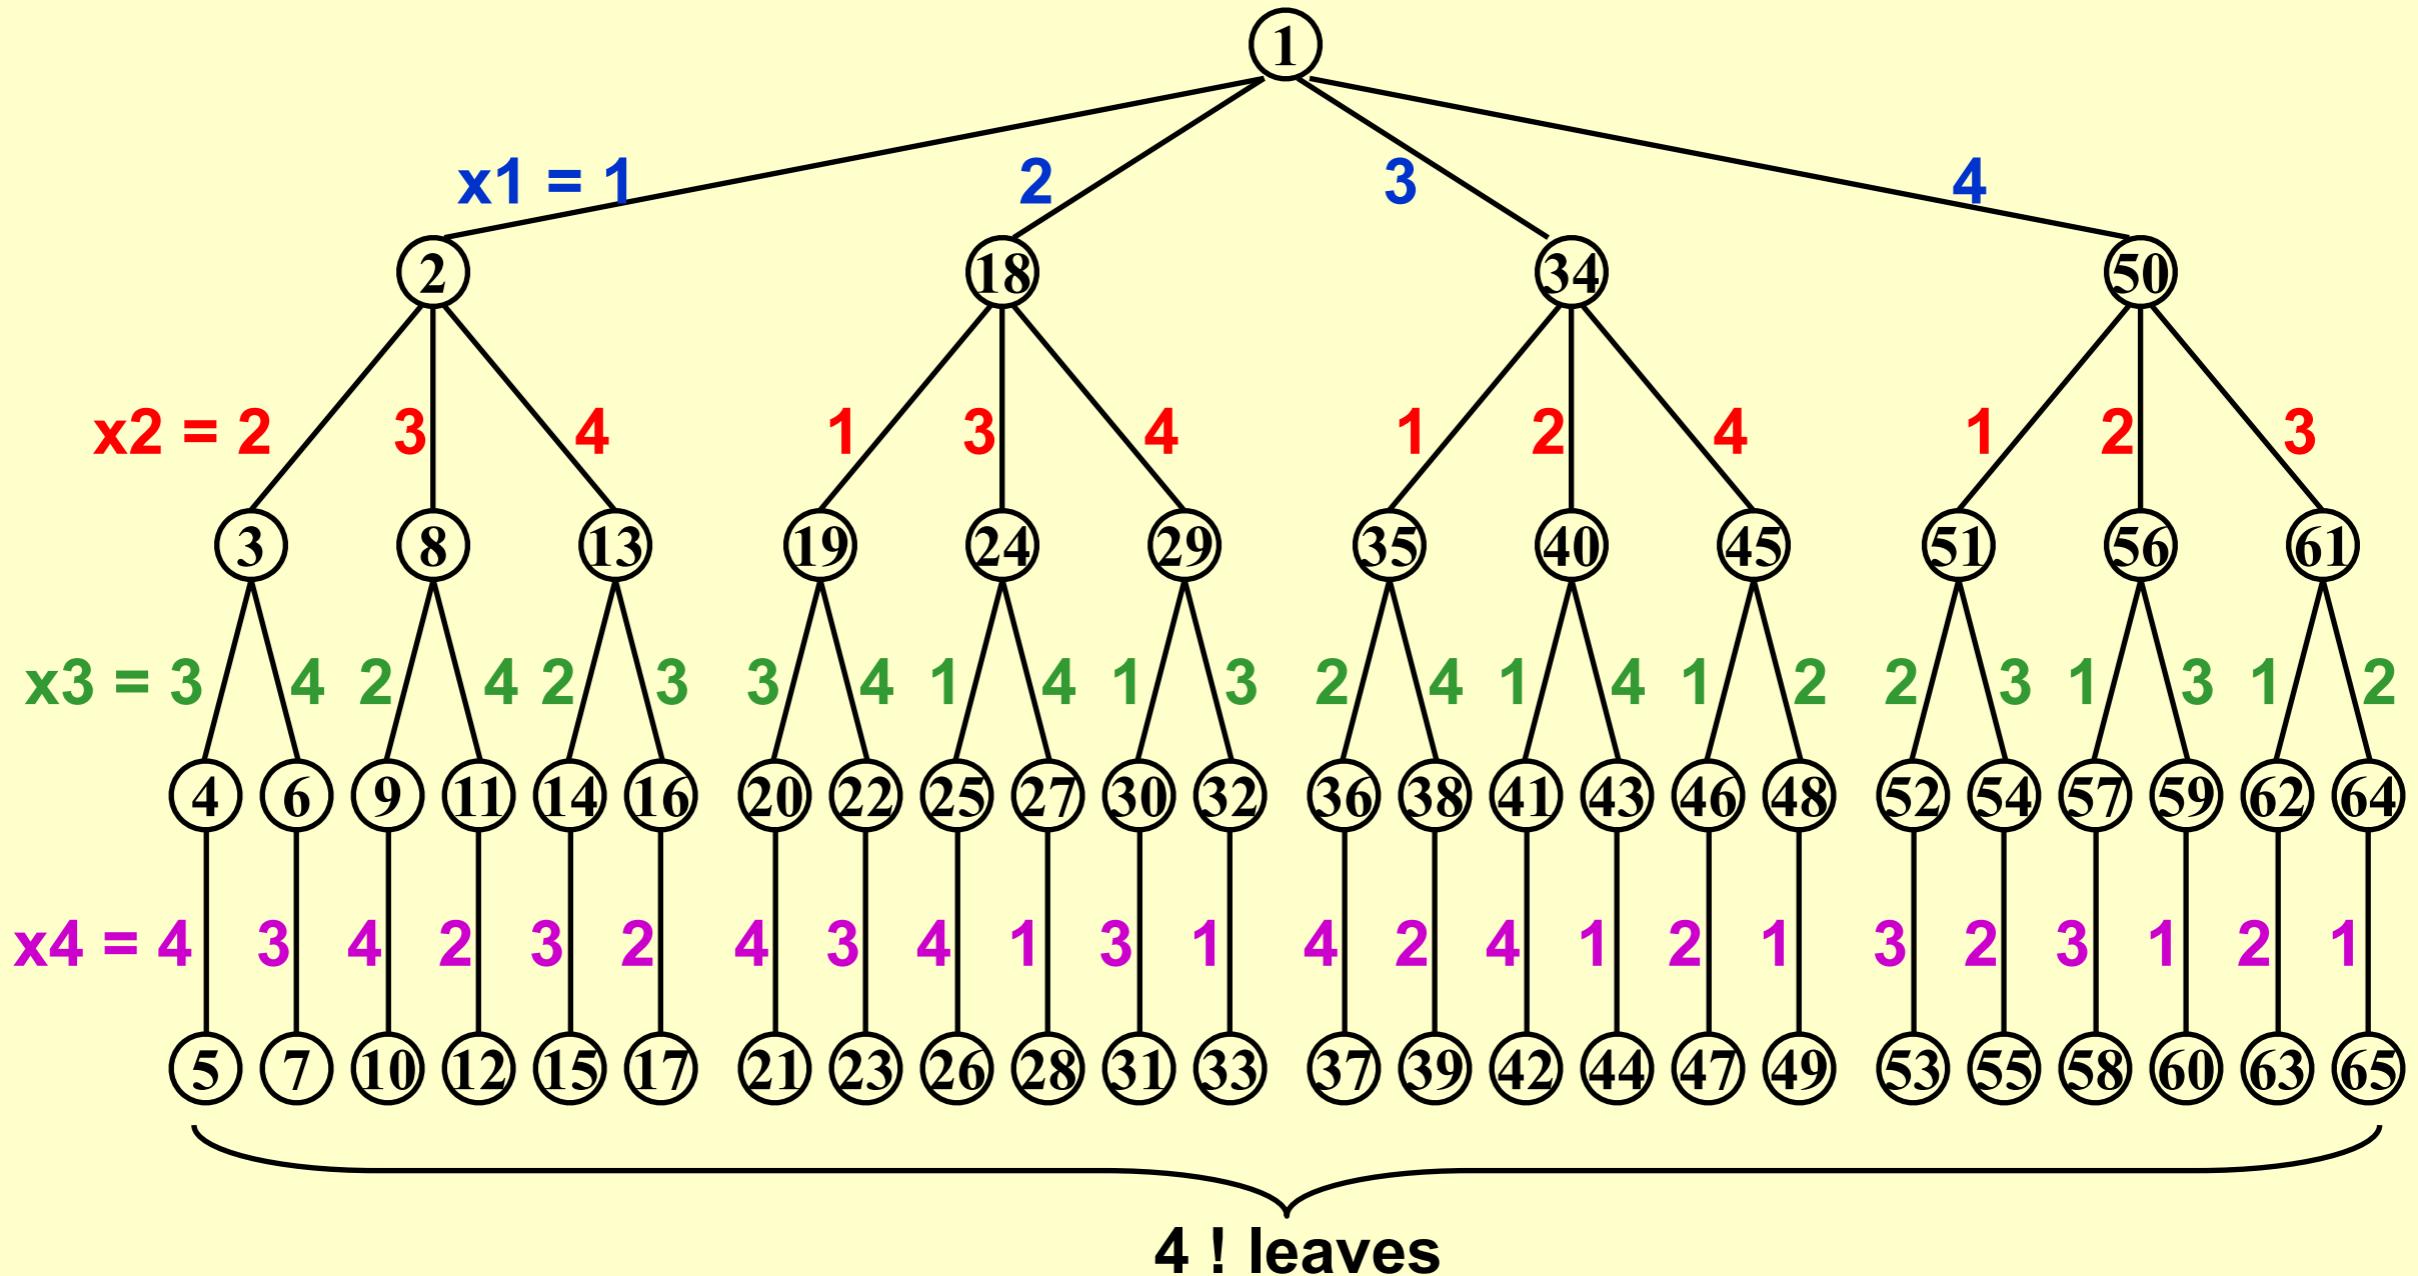
Constraints:

- ① $S_i = \{ 1, 2, 3, 4, 5, 6, 7, 8 \}$ for $1 \leq i \leq 8$
- ② $x_i \neq x_j$ if $i \neq j$
- ③ $(x_i - x_j) / (i - j) \neq \pm 1$

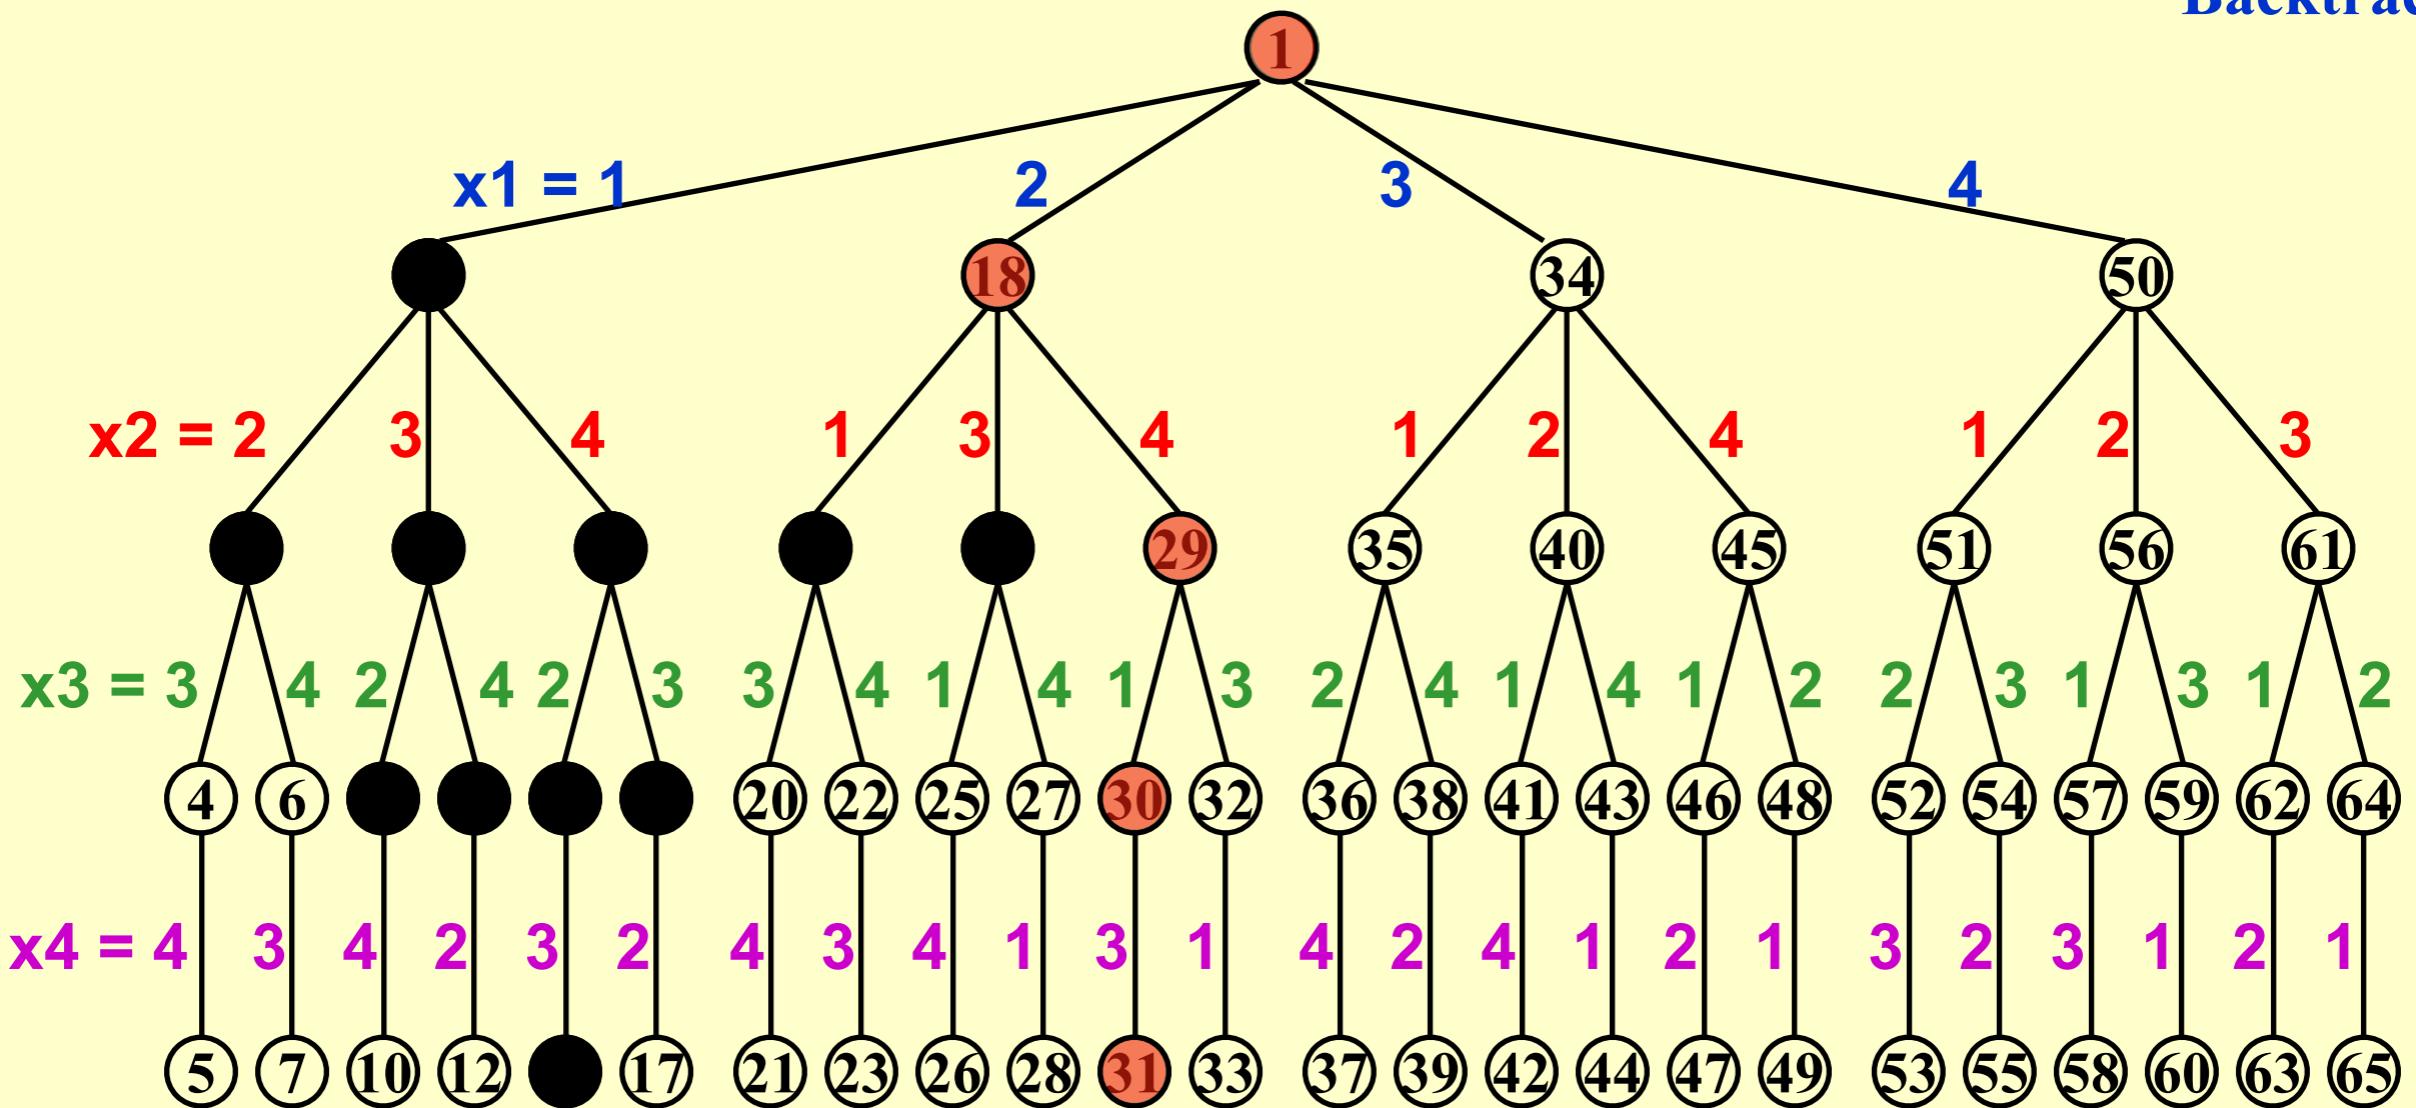
For the problem with n queens,
there are $n!$ candidates
in the solution space.

Method: Take the problem of 4 queens as an example

Step 1: Construct a game tree



Each path from the **root** to a **leaf** defines an element of the solution space.



Step 2: Perform a **depth-first search** (post-order traversal) to examine the paths

(2, 4, 1, 3)

Note: No tree is actually constructed. The game tree is just an abstract concept.

	Q	
		Q
Q		
		Q

The Turnpike Reconstruction Problem

Given N points on the x -axis with coordinates $x_1 < x_2 < \dots < x_N$.

Assume that $x_1 = 0$. There are $N(N-1)/2$ distances between every pair of points.

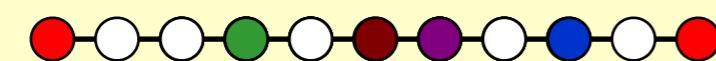
Given $N(N-1)/2$ distances. Reconstruct a point set from the distances.

Given $D = \{ 1, 2, 2, 2, 3, 3, 3, 4, 5, 5, 5, 6, 7, 8, 10 \}$

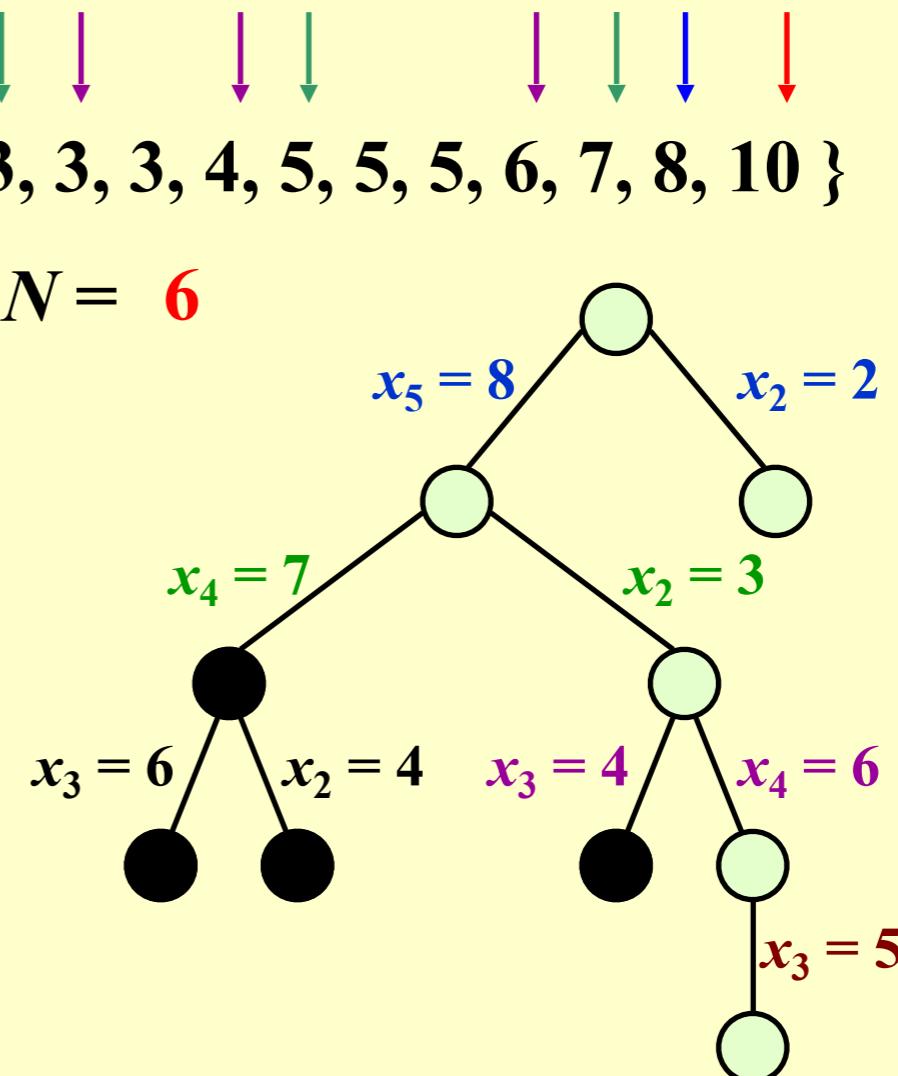
Step 1: $N(N-1)/2 = 15$ implies $N = 6$

Step 2: $x_1 = 0$ and $x_6 = 10$

Step 3: find the next largest distance and check



(0, 3, 5, 6, 8, 10)



```

bool Reconstruct ( DistType X[ ], DistSet D, int N, int left, int right )
{ /* X[1]...X[left-1] and X[right+1]...X[N] are solved */
  bool Found = false;
  if ( Is_Empty( D ) )
    return true; /* solved */
  D_max = Find_Max( D );
  /* option 1: X[right] = D_max */
  /* check if |D_max-X[i]| ∈ D is true for all X[i]'s that have been solved */
  OK = Check( D_max, N, left, right ); /* pruning */
  if ( OK ) { /* add X[right] and update D */
    X[right] = D_max;
    for ( i=1; i<left; i++ ) Delete( |X[right]-X[i]|, D );
    for ( i=right+1; i<=N; i++ ) Delete( |X[right]-X[i]|, D );
    Found = Reconstruct ( X, D, N, left, right-1 );
    if ( !Found ) { /* if does not work, undo */
      for ( i=1; i<left; i++ ) Insert( |X[right]-X[i]|, D );
      for ( i=right+1; i<=N; i++ ) Insert( |X[right]-X[i]|, D );
    }
  }
  /* finish checking option 1 */
}

```

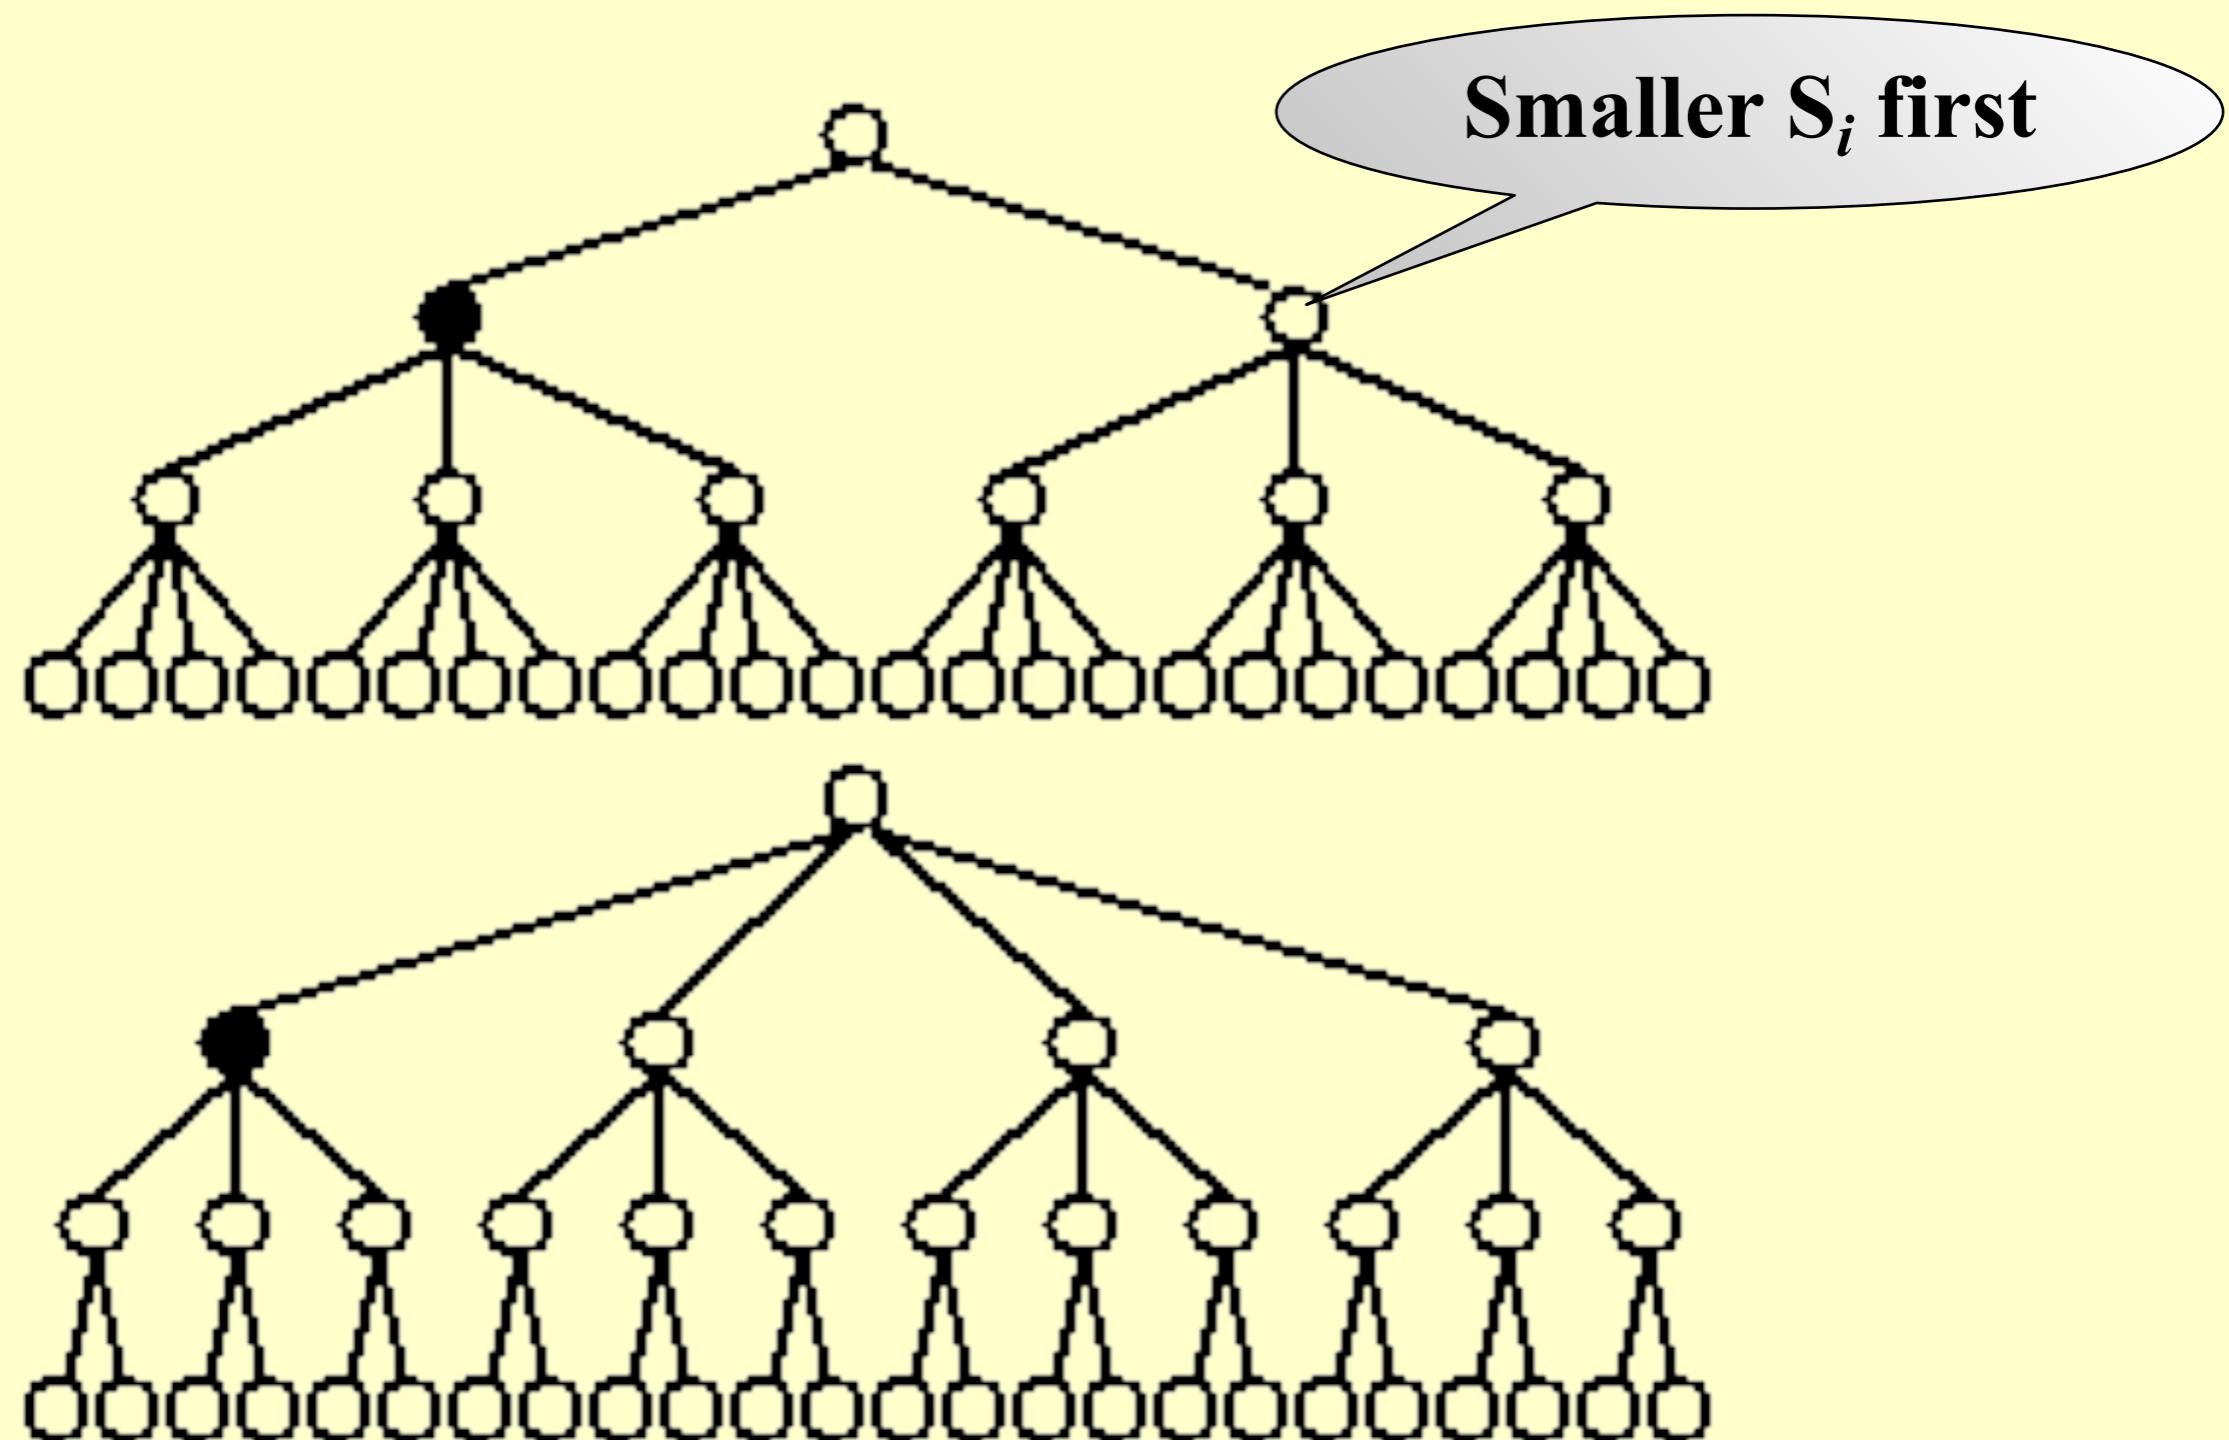
```
if ( !Found ) { /* if option 1 does not work */
    /* option 2: X[left] = X[N]-D_max */
    OK = Check( X[N]-D_max, N, left, right );
    if ( OK ) {
        X[left] = X[N] - D_max;
        for ( i=1; i<left; i++ ) Delete( |X[left]-X[i]|, D );
        for ( i=right+1; i<=N; i++ ) Delete( |X[left]-X[i]|, D );
        Found = Reconstruct (X, D, N, left+1, right );
        if ( !Found ) {
            for ( i=1; i<left; i++ ) Insert( |X[left]-X[i]|, D );
            for ( i=right+1; i<=N; i++ ) Insert( |X[left]-X[i]|, D );
        }
    }
    /* finish checking option 2 */
} /* finish checking all the options */

return Found;
}
```

A Template

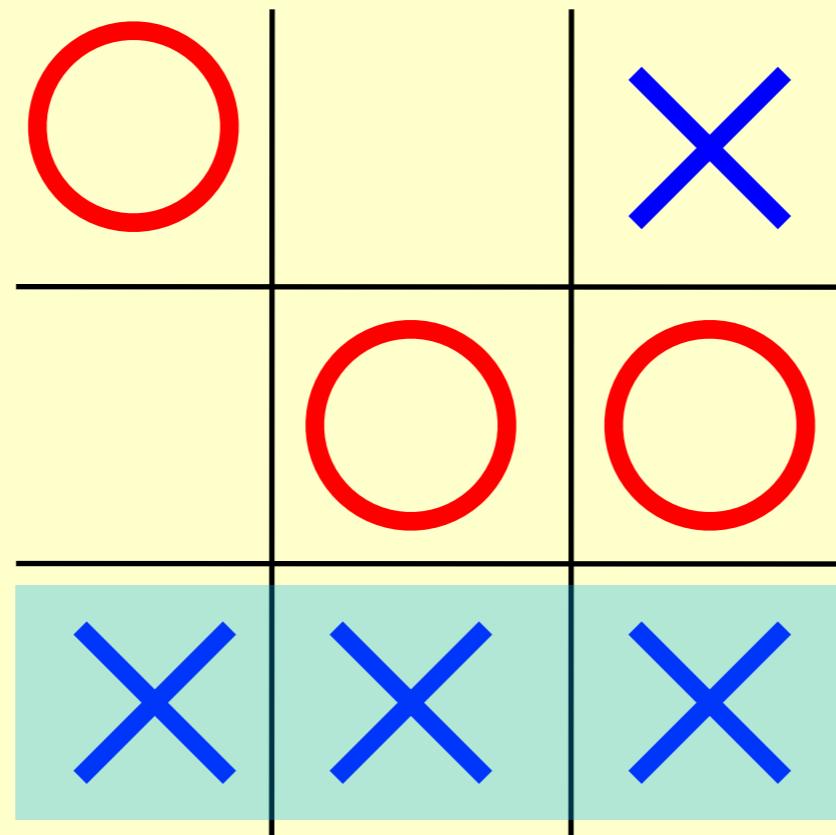
```
bool Backtracking ( int i )
{ Found = false;
  if ( i > N )
    return true; /* solved with (x1, ..., xN) */
  for ( each xi ∈ Si ) {
    /* check if satisfies the restriction R */
    OK = Check((x1, ..., xi) , R ); /* pruning */
    if ( OK ) {
      Count xi in;
      Found = Backtracking( i+1 );
      if ( !Found )
        Undo( i ); /* recover to (x1, ..., xi-1) */
    }
    if ( Found ) break;
  }
  return Found;
}
```

When different S_i 's have different sizes



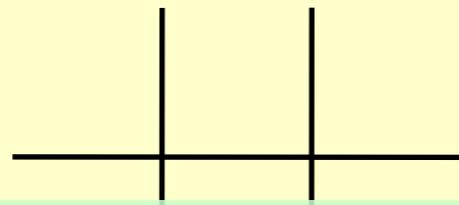
Games – *how did AlphaGo win*

Tic-tac-toe

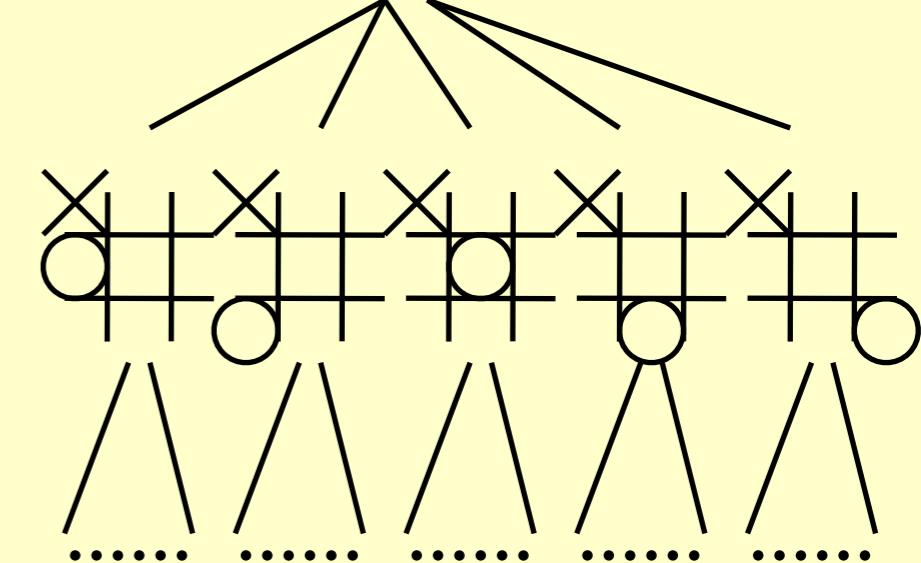
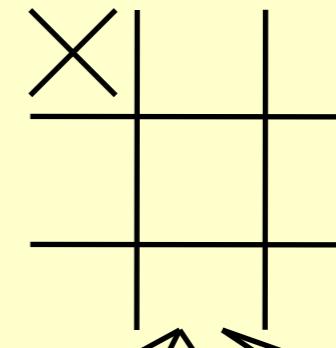
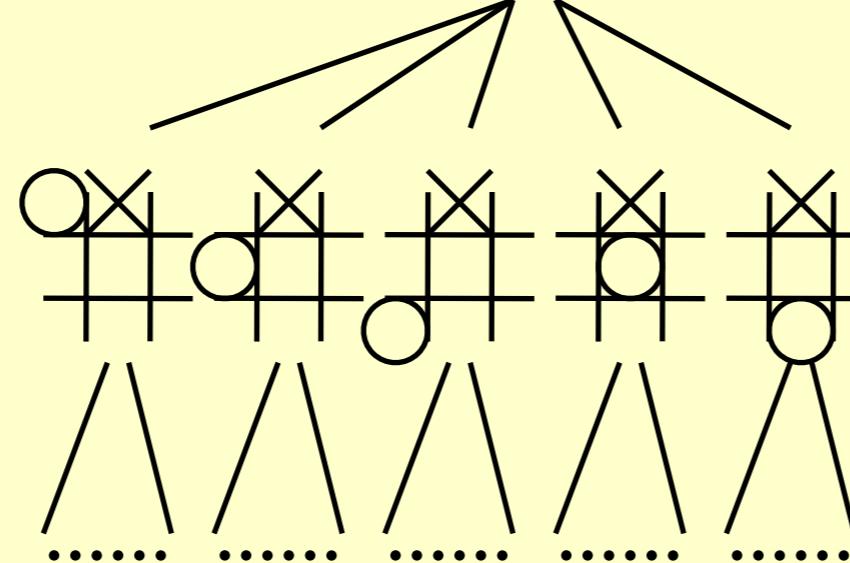
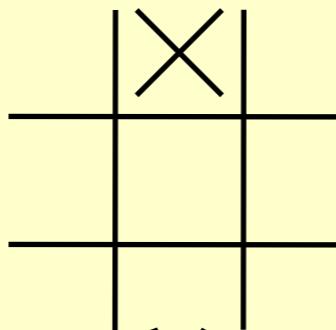
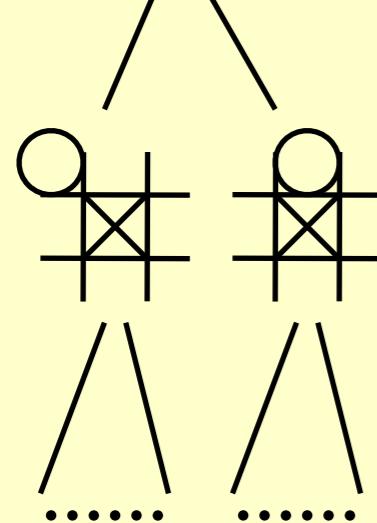
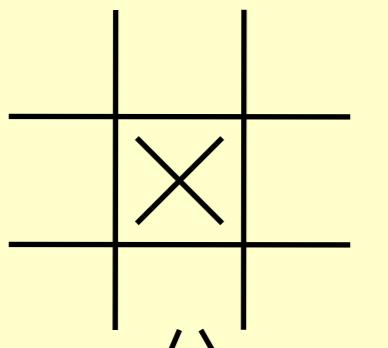


The player who succeeds in placing three of their marks in a horizontal, vertical, or diagonal row wins the game.

Tic-tac-toe



- **19,683** possible board layouts (3^9 since each of the nine spaces can be X, O or blank), and
- **362,880** (i.e., $9!$) possible games (different sequences for placing the Xs and Os on the board)

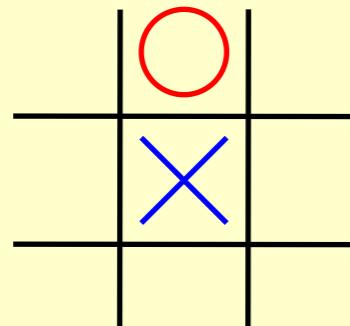


Tic-tac-toe: Minimax Strategy

Use an evaluation function to quantify the "**goodness**" of a position. For example:

$$f(P) = W_{Computer} - W_{Human}$$

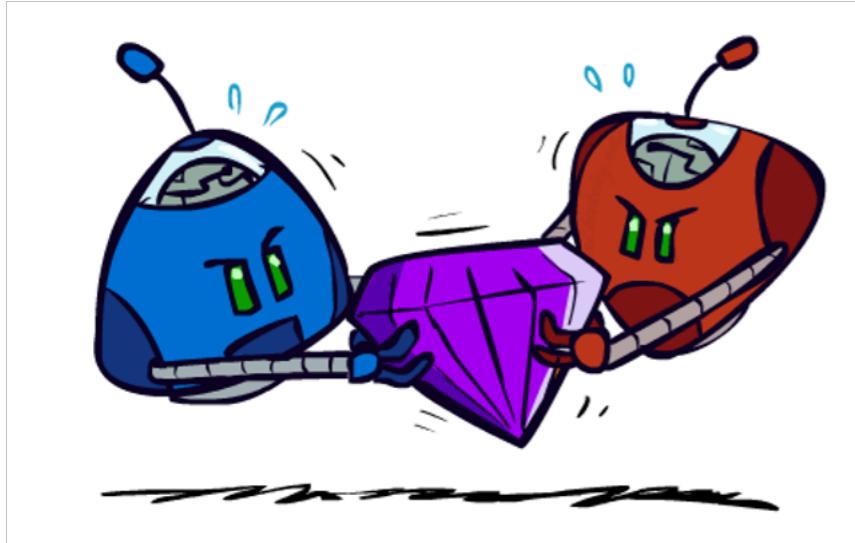
where W is the number of potential wins at position P .



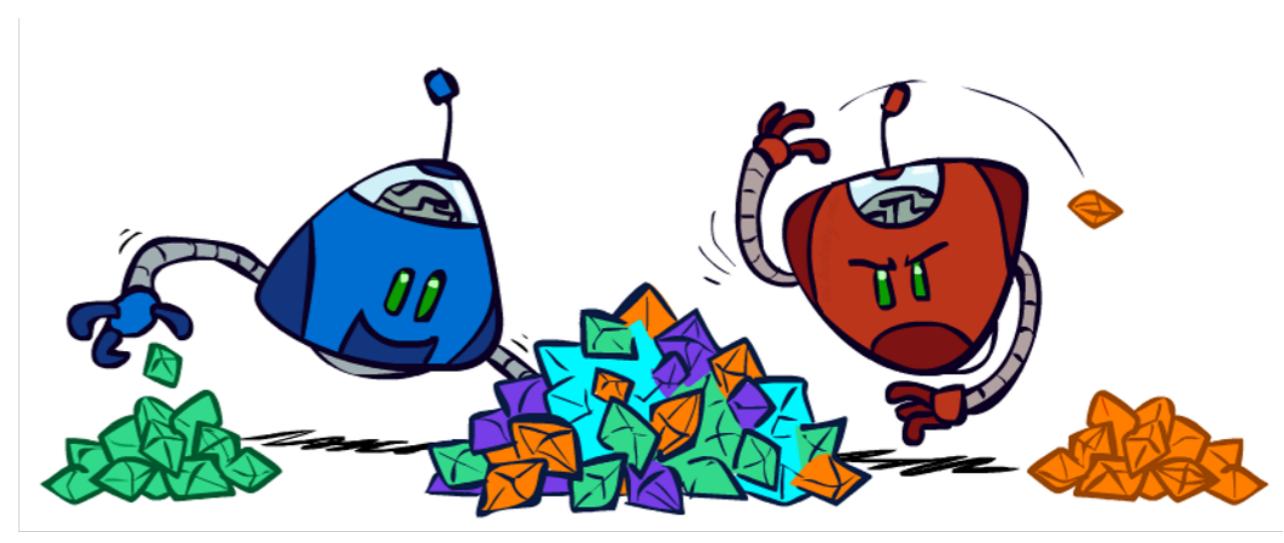
$$f(P) = 6 - 4 = 2$$

The **human** is trying to *minimize* the value of the position P , while the **computer** is trying to *maximize* it.

Two-Player Zero-Sum Game



vs.



In zero-sum games, the utility functions of the two players are coupled:
How much one wins equals to how much the other loses: competitive.
Can the players be cooperative?

The Simplest Formulation: Two-Step Game

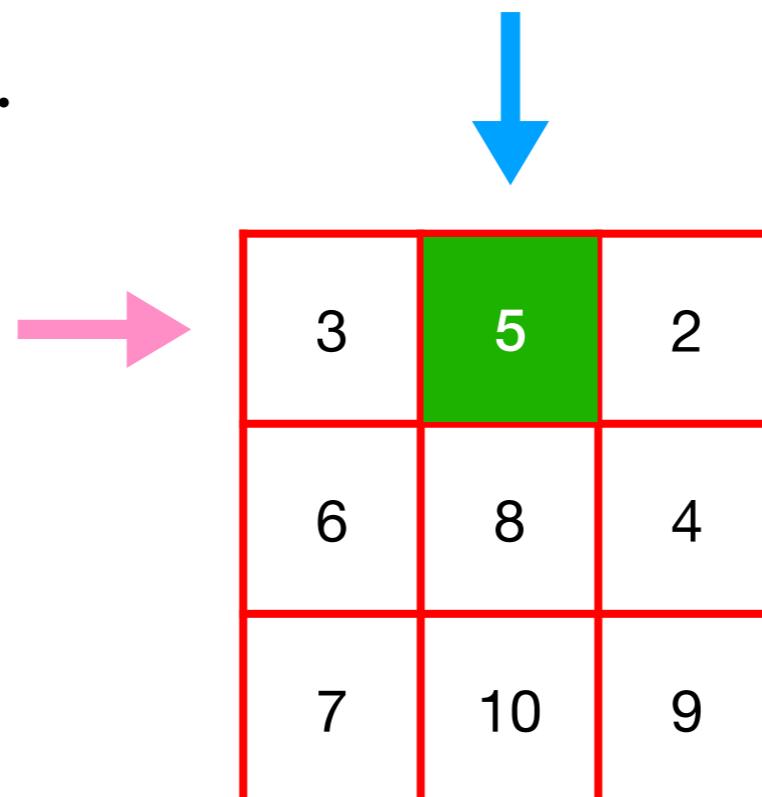
- Two step game: Alice chooses row i , then Bob chooses column j
- Outcome: Alice loses (Bob wins) the utility in entry i, j zero-sum
- A MinMax Game.

3	5	2
6	8	4
7	10	9

$$U(i^*, j^*) = \min_i \max_j U(i, j)$$

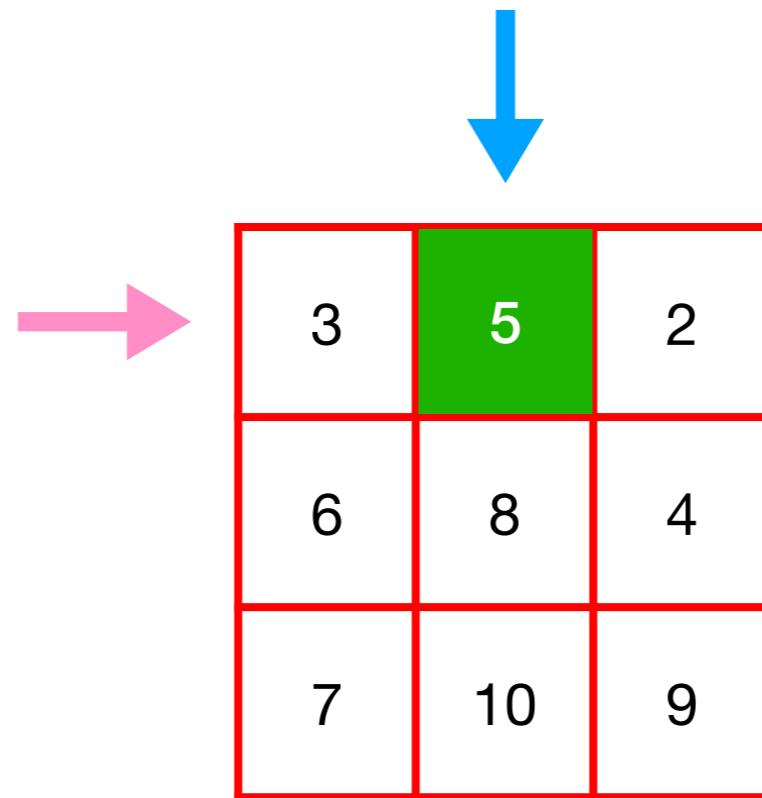
The Simplest Formulation: Two-Step Game

- Two step game: Bob chooses row j , then Alice chooses column i
- Outcome: Alice loses (Bob wins) the utility in entry i, j
- A MaxMin Game.



$$U(i^*, j^*) = \max_j \min_i U(i, j)$$

Is MinMax equivalent to MaxMin?



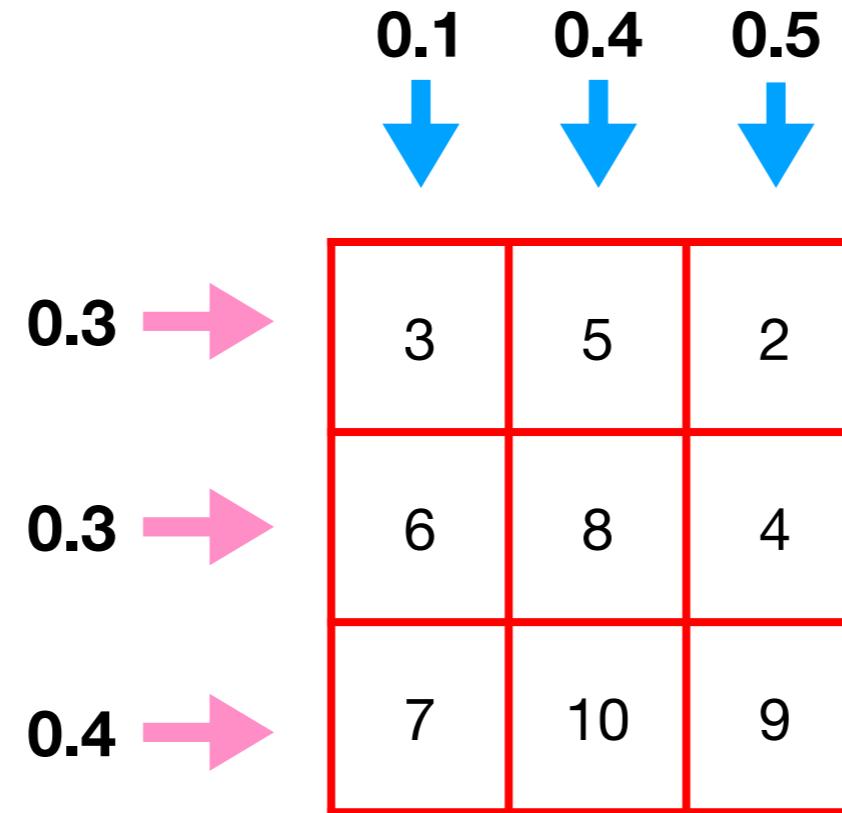
Does Alice (Bob) lose (win) the same utility in MinMax and MaxMin games?

Theorem:

$$\max_j \min_i U(i, j) \leq \min_i \max_j U(i, j)$$

Moving first is always worse!

Deterministic vs. Stochastic Strategy



Mixed (Stochastic) strategy: choosing a distribution over actions instead of a single action (pure or deterministic strategy)

Von Neumann's Minimax Theorem:

$$\min_{P_i} \max_{P_j} \mathbb{E}_{P_i, P_j} [U(i, j)] = \max_{P_j} \min_{P_i} \mathbb{E}_{P_i, P_j} [U(i, j)]$$

MinMax and MaxMin Games are equivalent for mixed strategy!

Nash Equilibrium

$$\min_{P_i} \max_{P_j} \mathbb{E}_{P_i, P_j} [U(i, j)] = \max_{P_j} \min_{P_i} \mathbb{E}_{P_i, P_j} [U(i, j)]$$

- Strong duality: the solution pair P_i^*, P_j^* achieving this equation is called the **saddle point** of the two-step zero-sum game.
- This is also the **Nash equilibrium** of the game.

Under Nash equilibrium, changing the strategy for any player herself would not be a good idea.

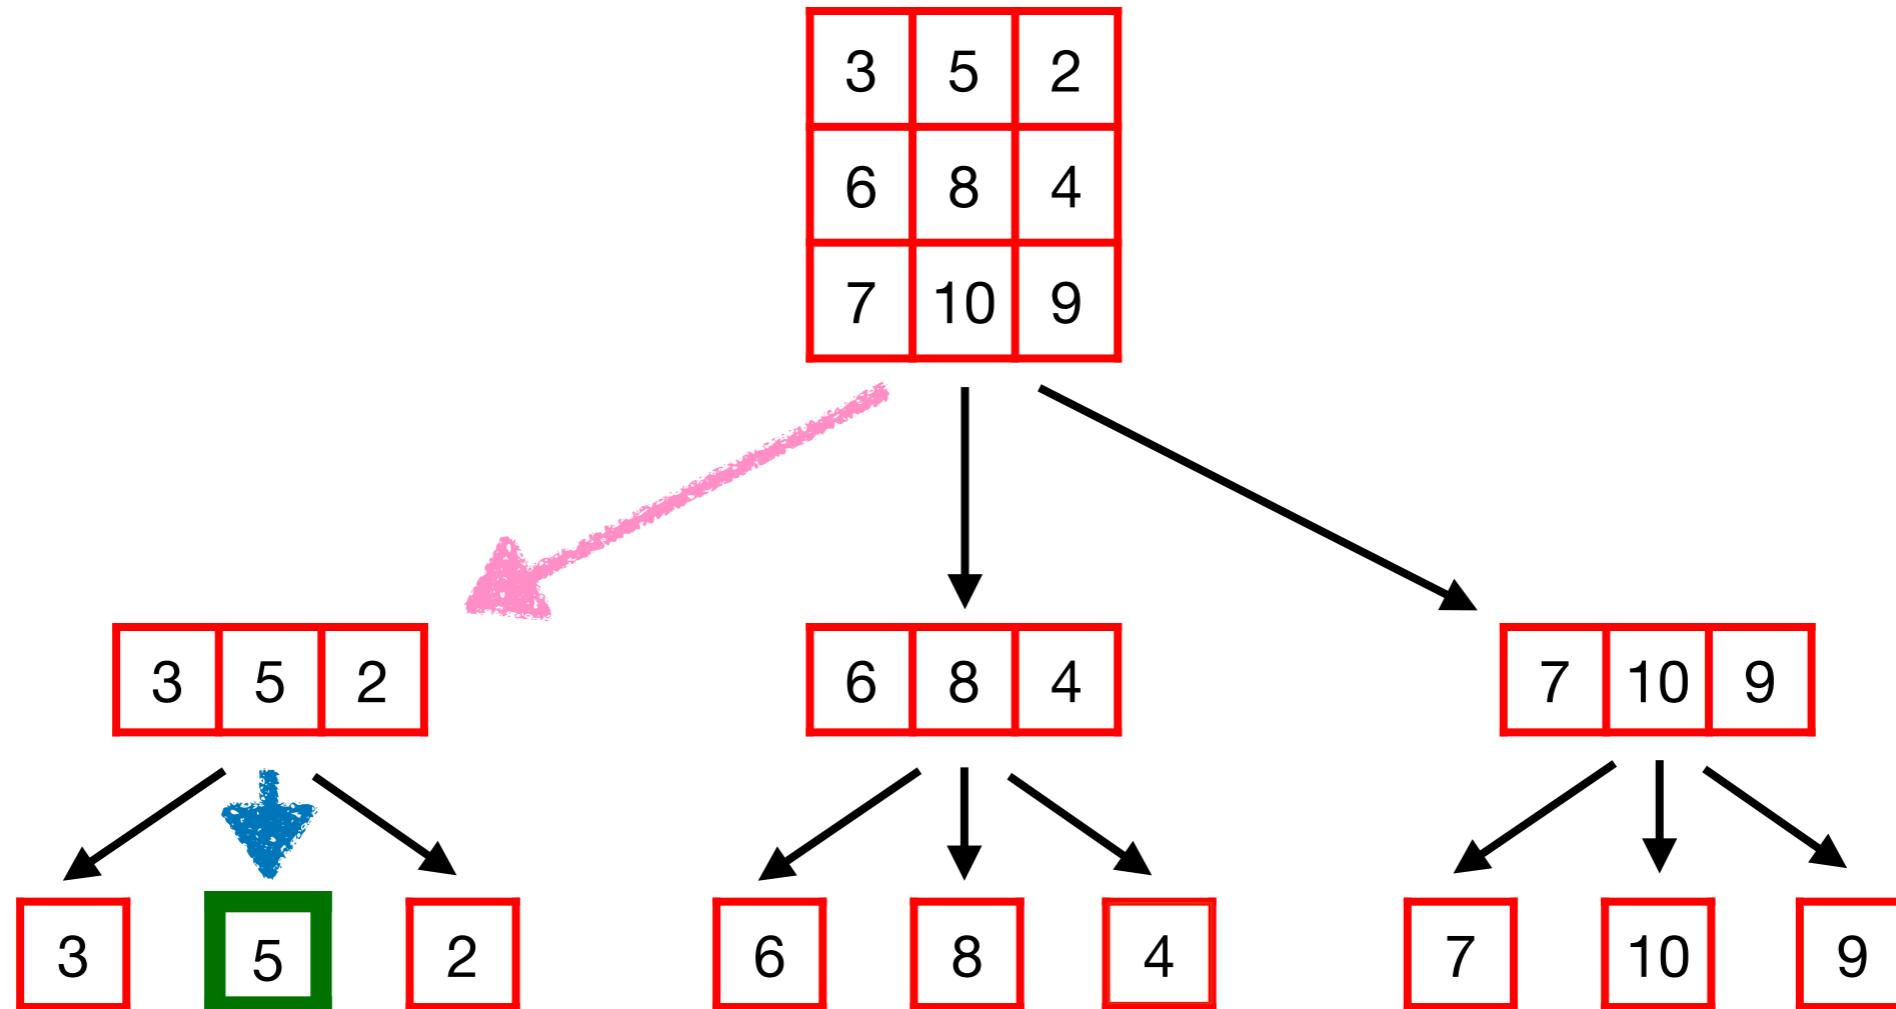
But Nash equilibrium does not mean optimal strategy!
You know the prisoner's dilemma.

Multi-Step Zero-Sum Game

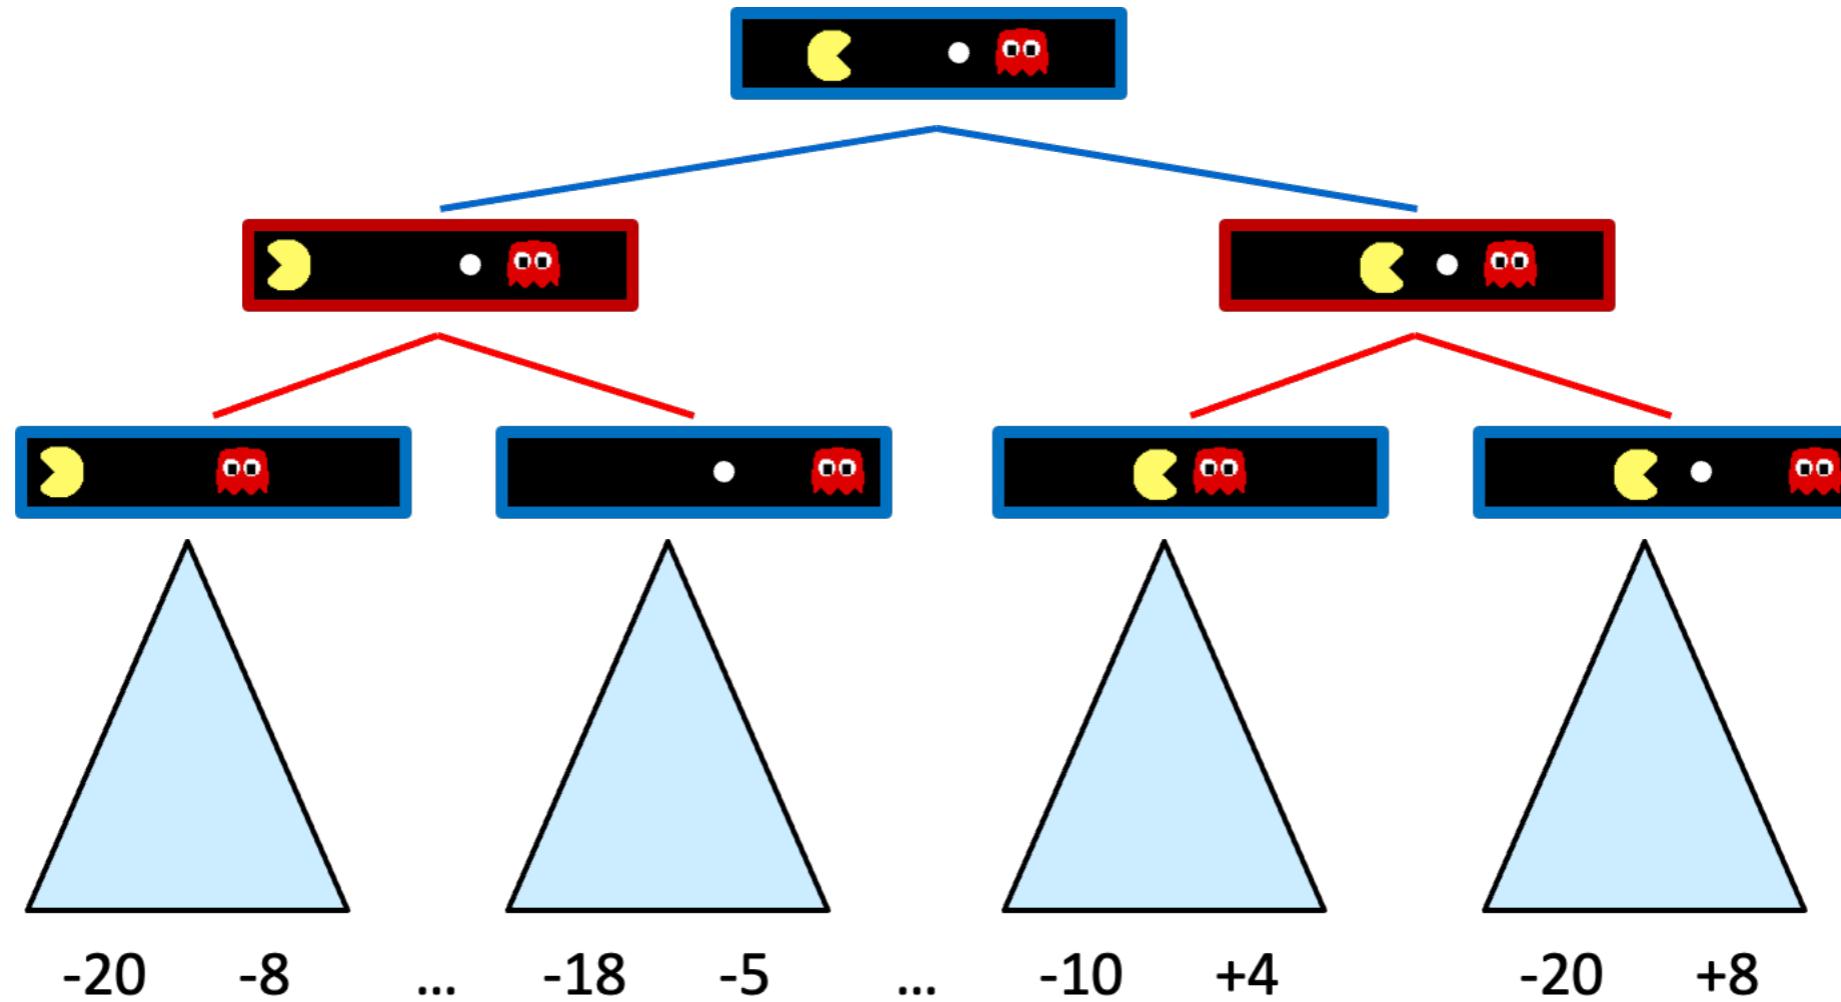
3	5	2
6	8	4
7	10	9



Game Tree of Two-Step Games



Multi-Step Zero-Sum Game

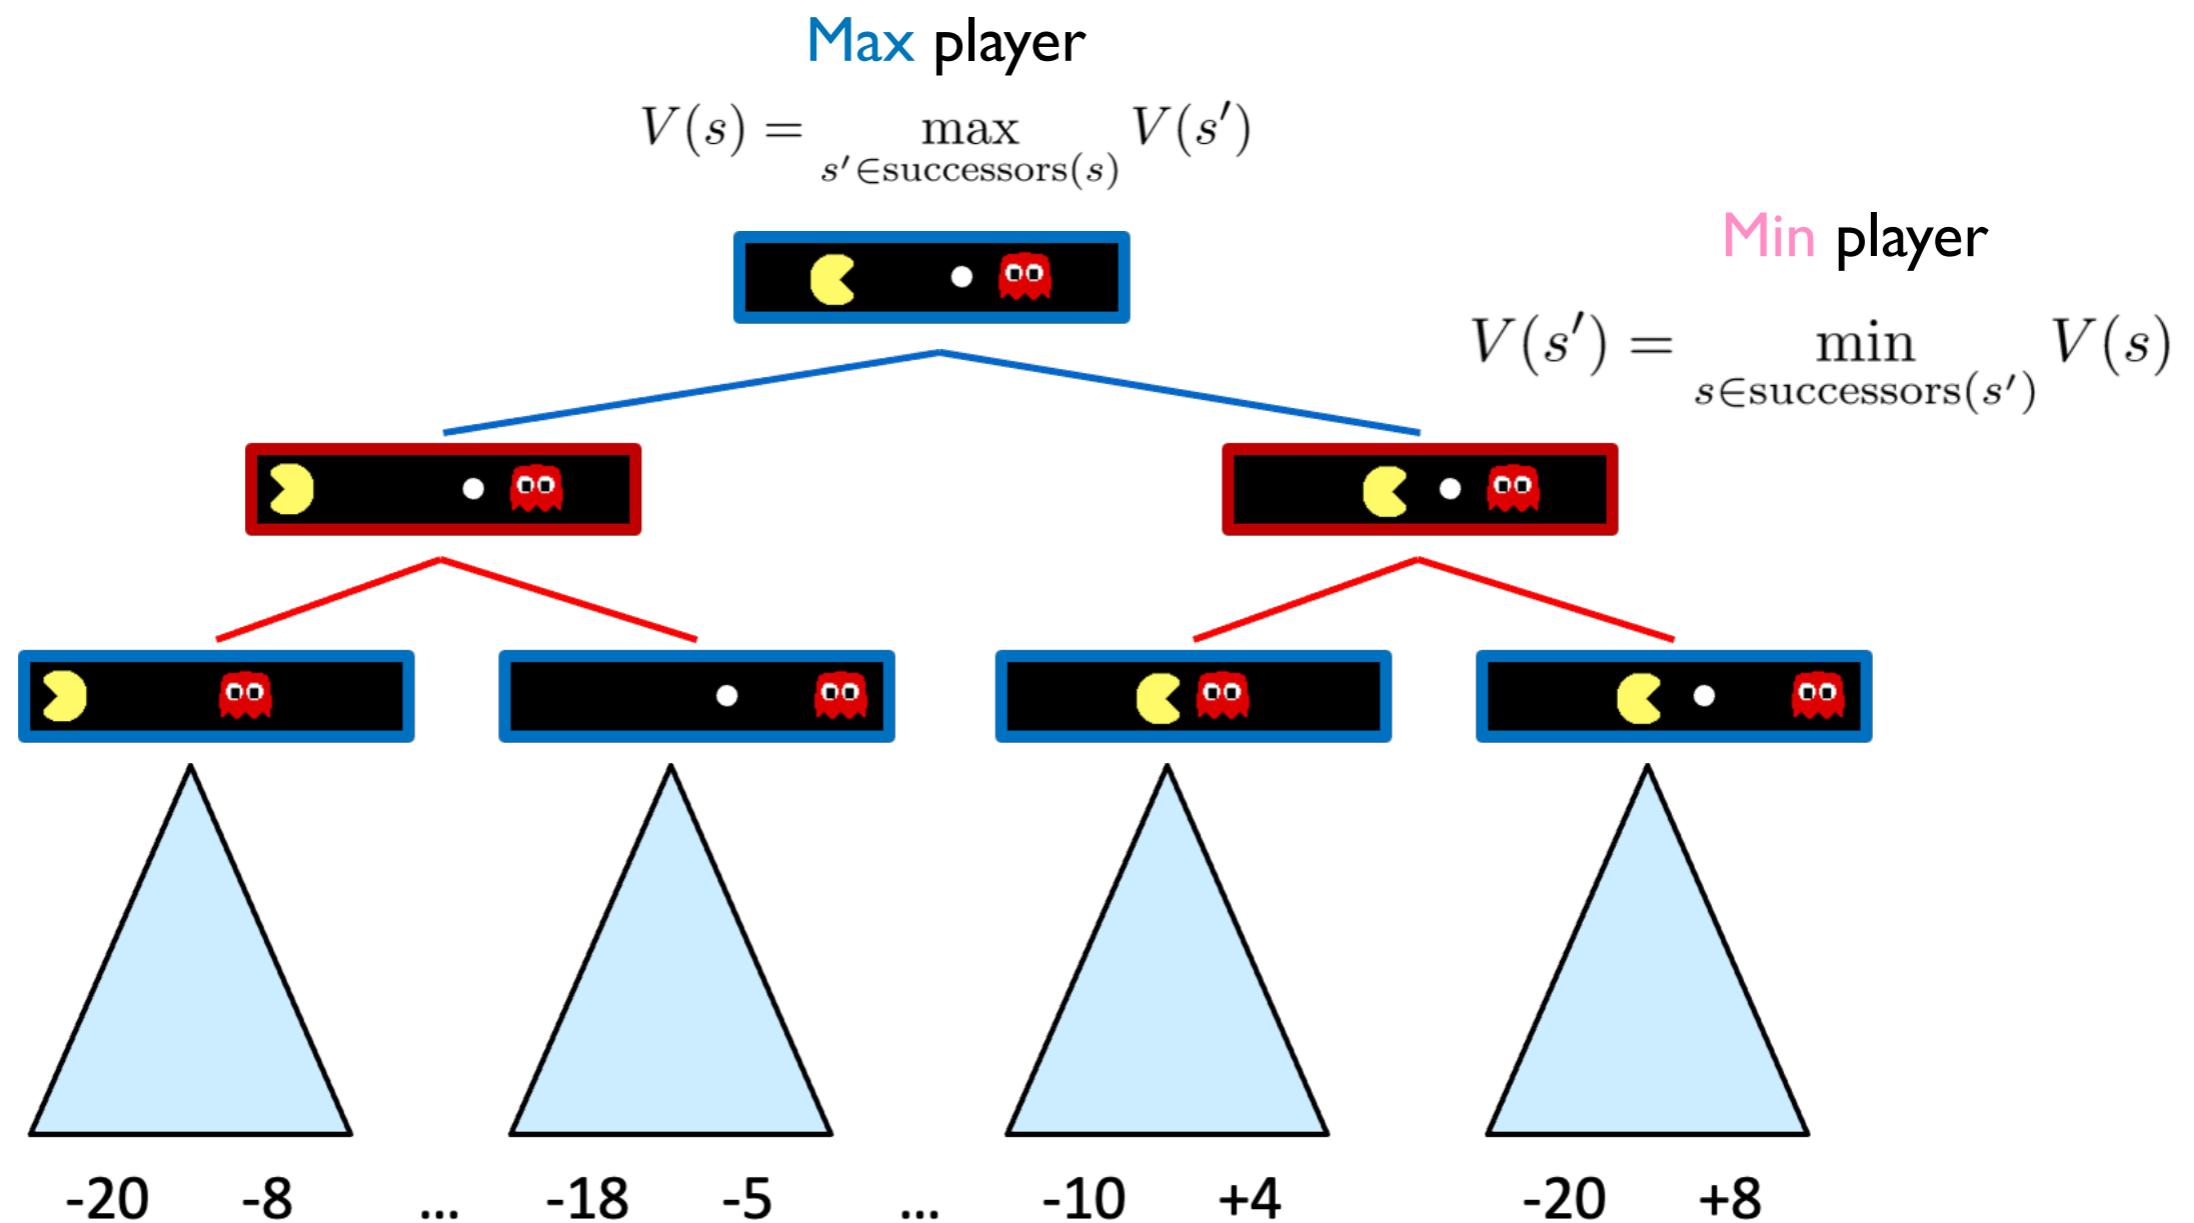


Can also be modeled as a search tree.

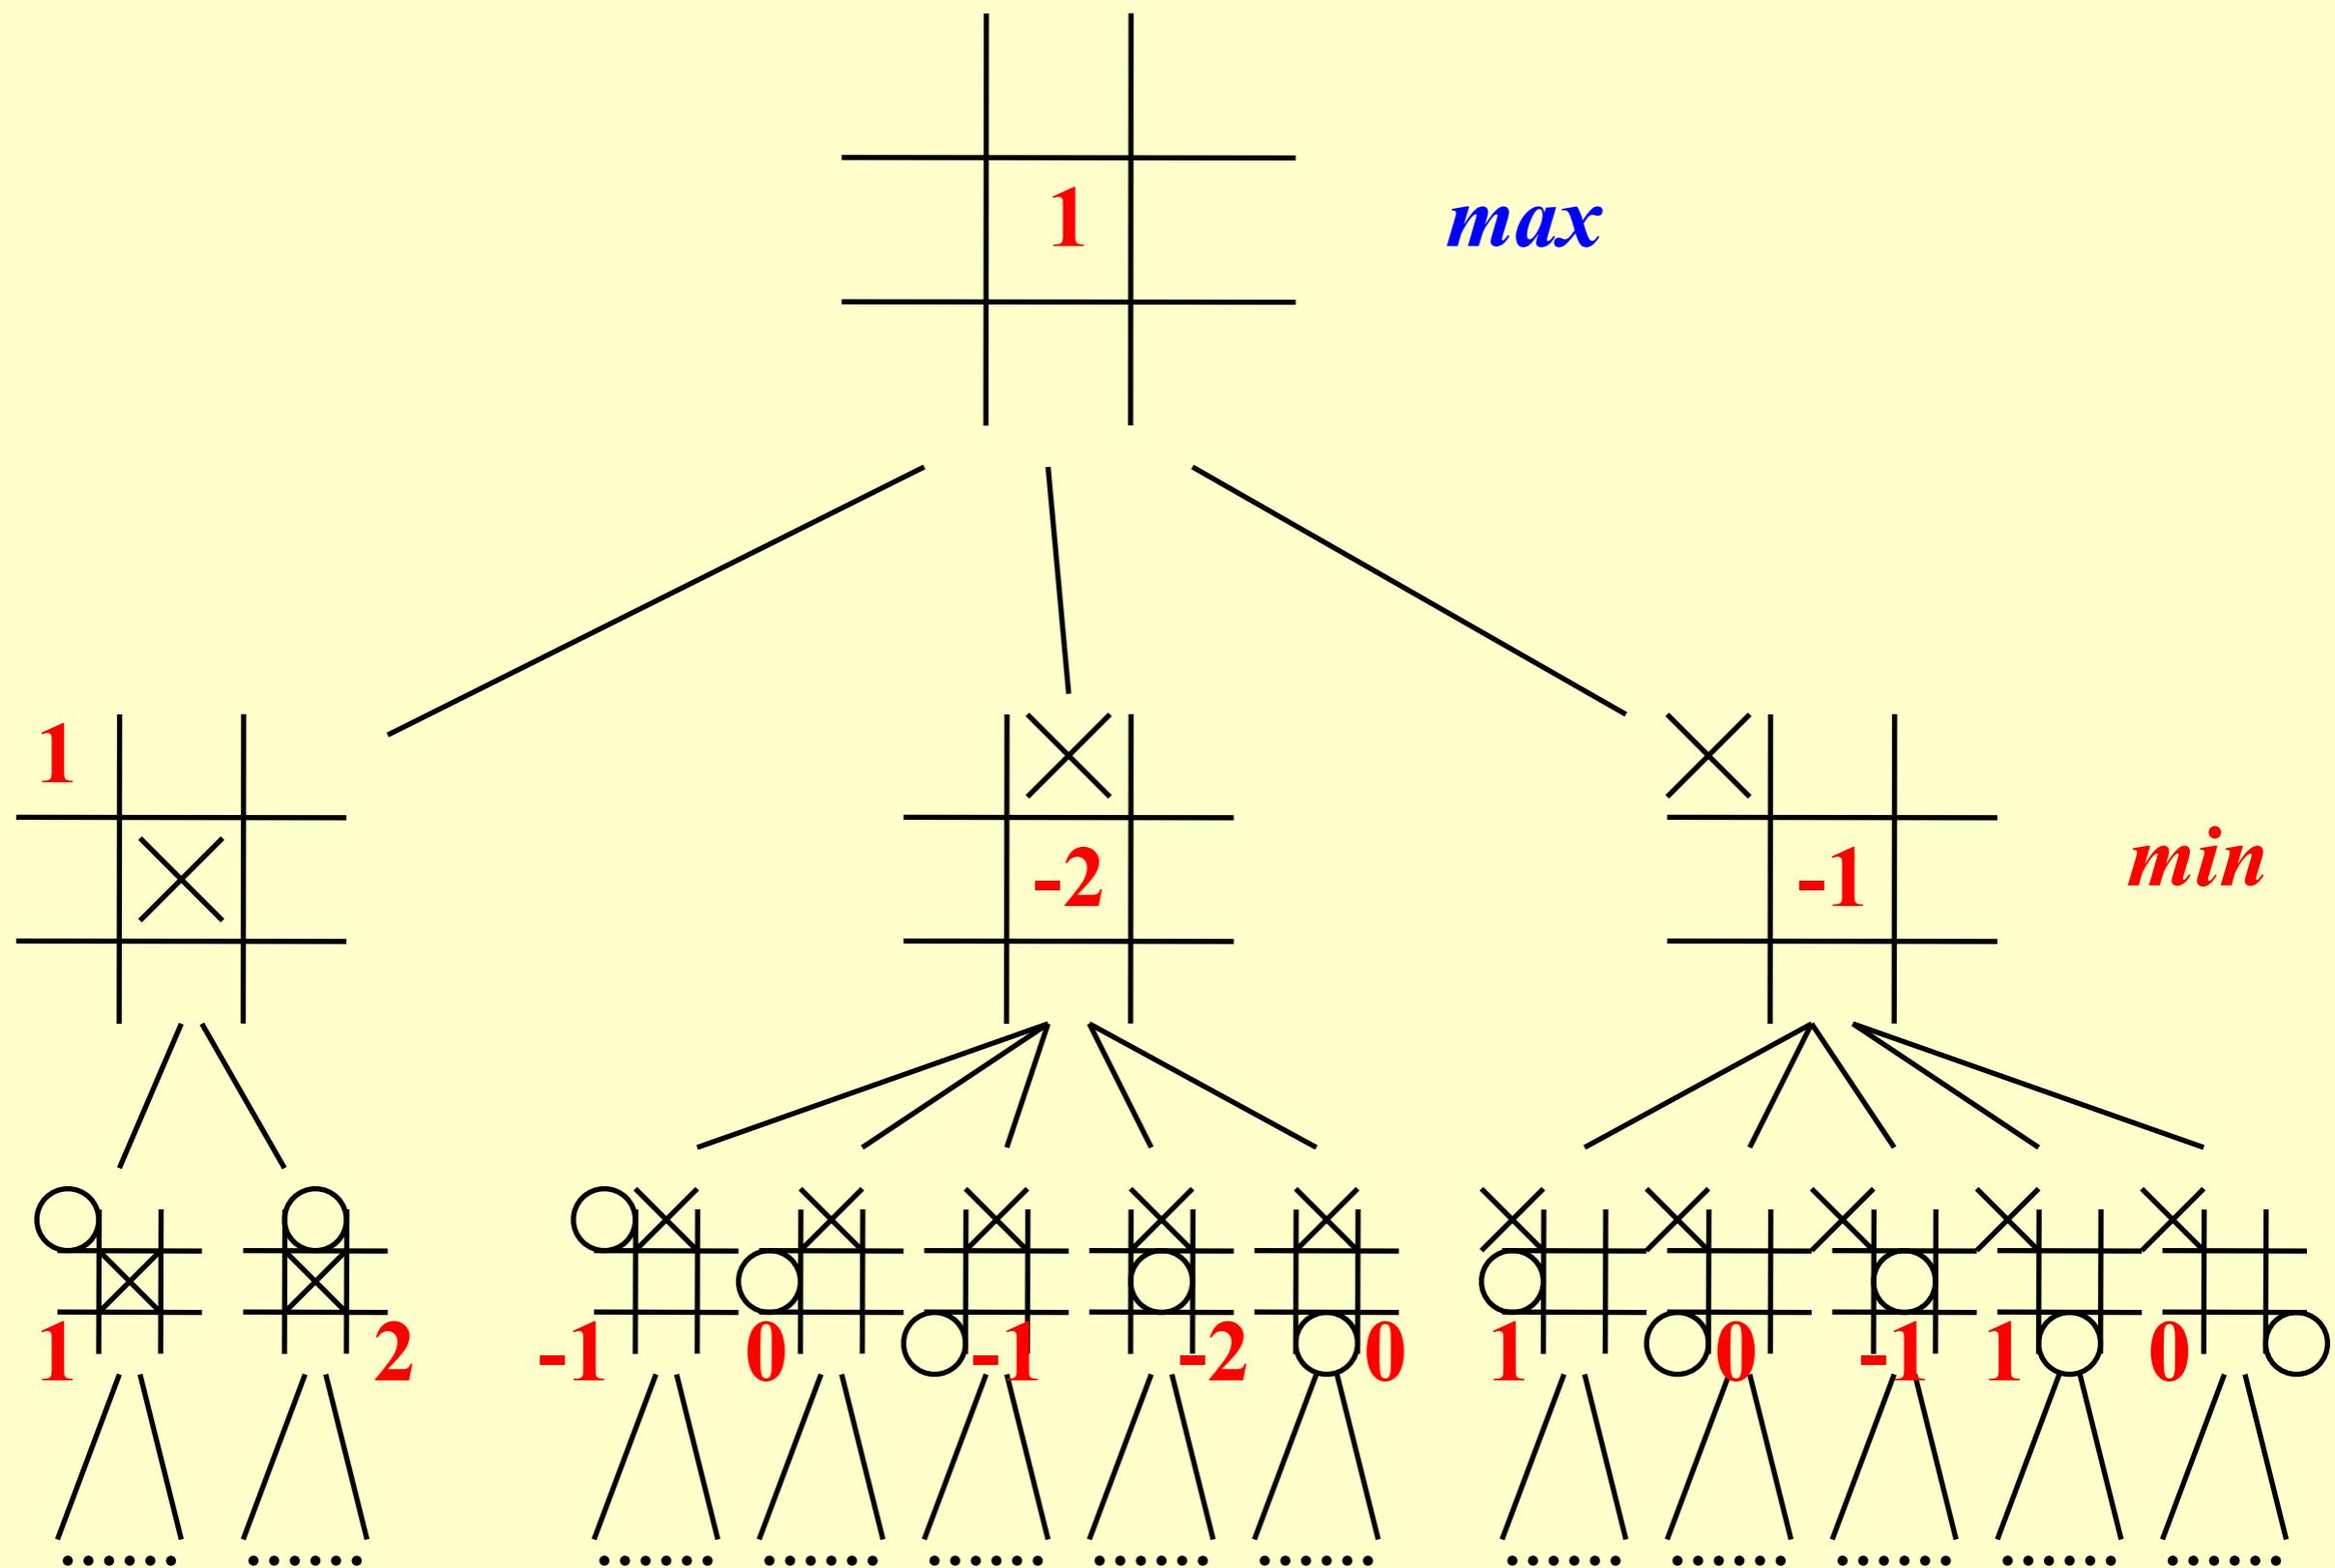
Lesson learned from two-step game:
The players should play **rationally**: use min and max strategies.
The players should **look ahead** when making decisions.

Min and Max Strategy

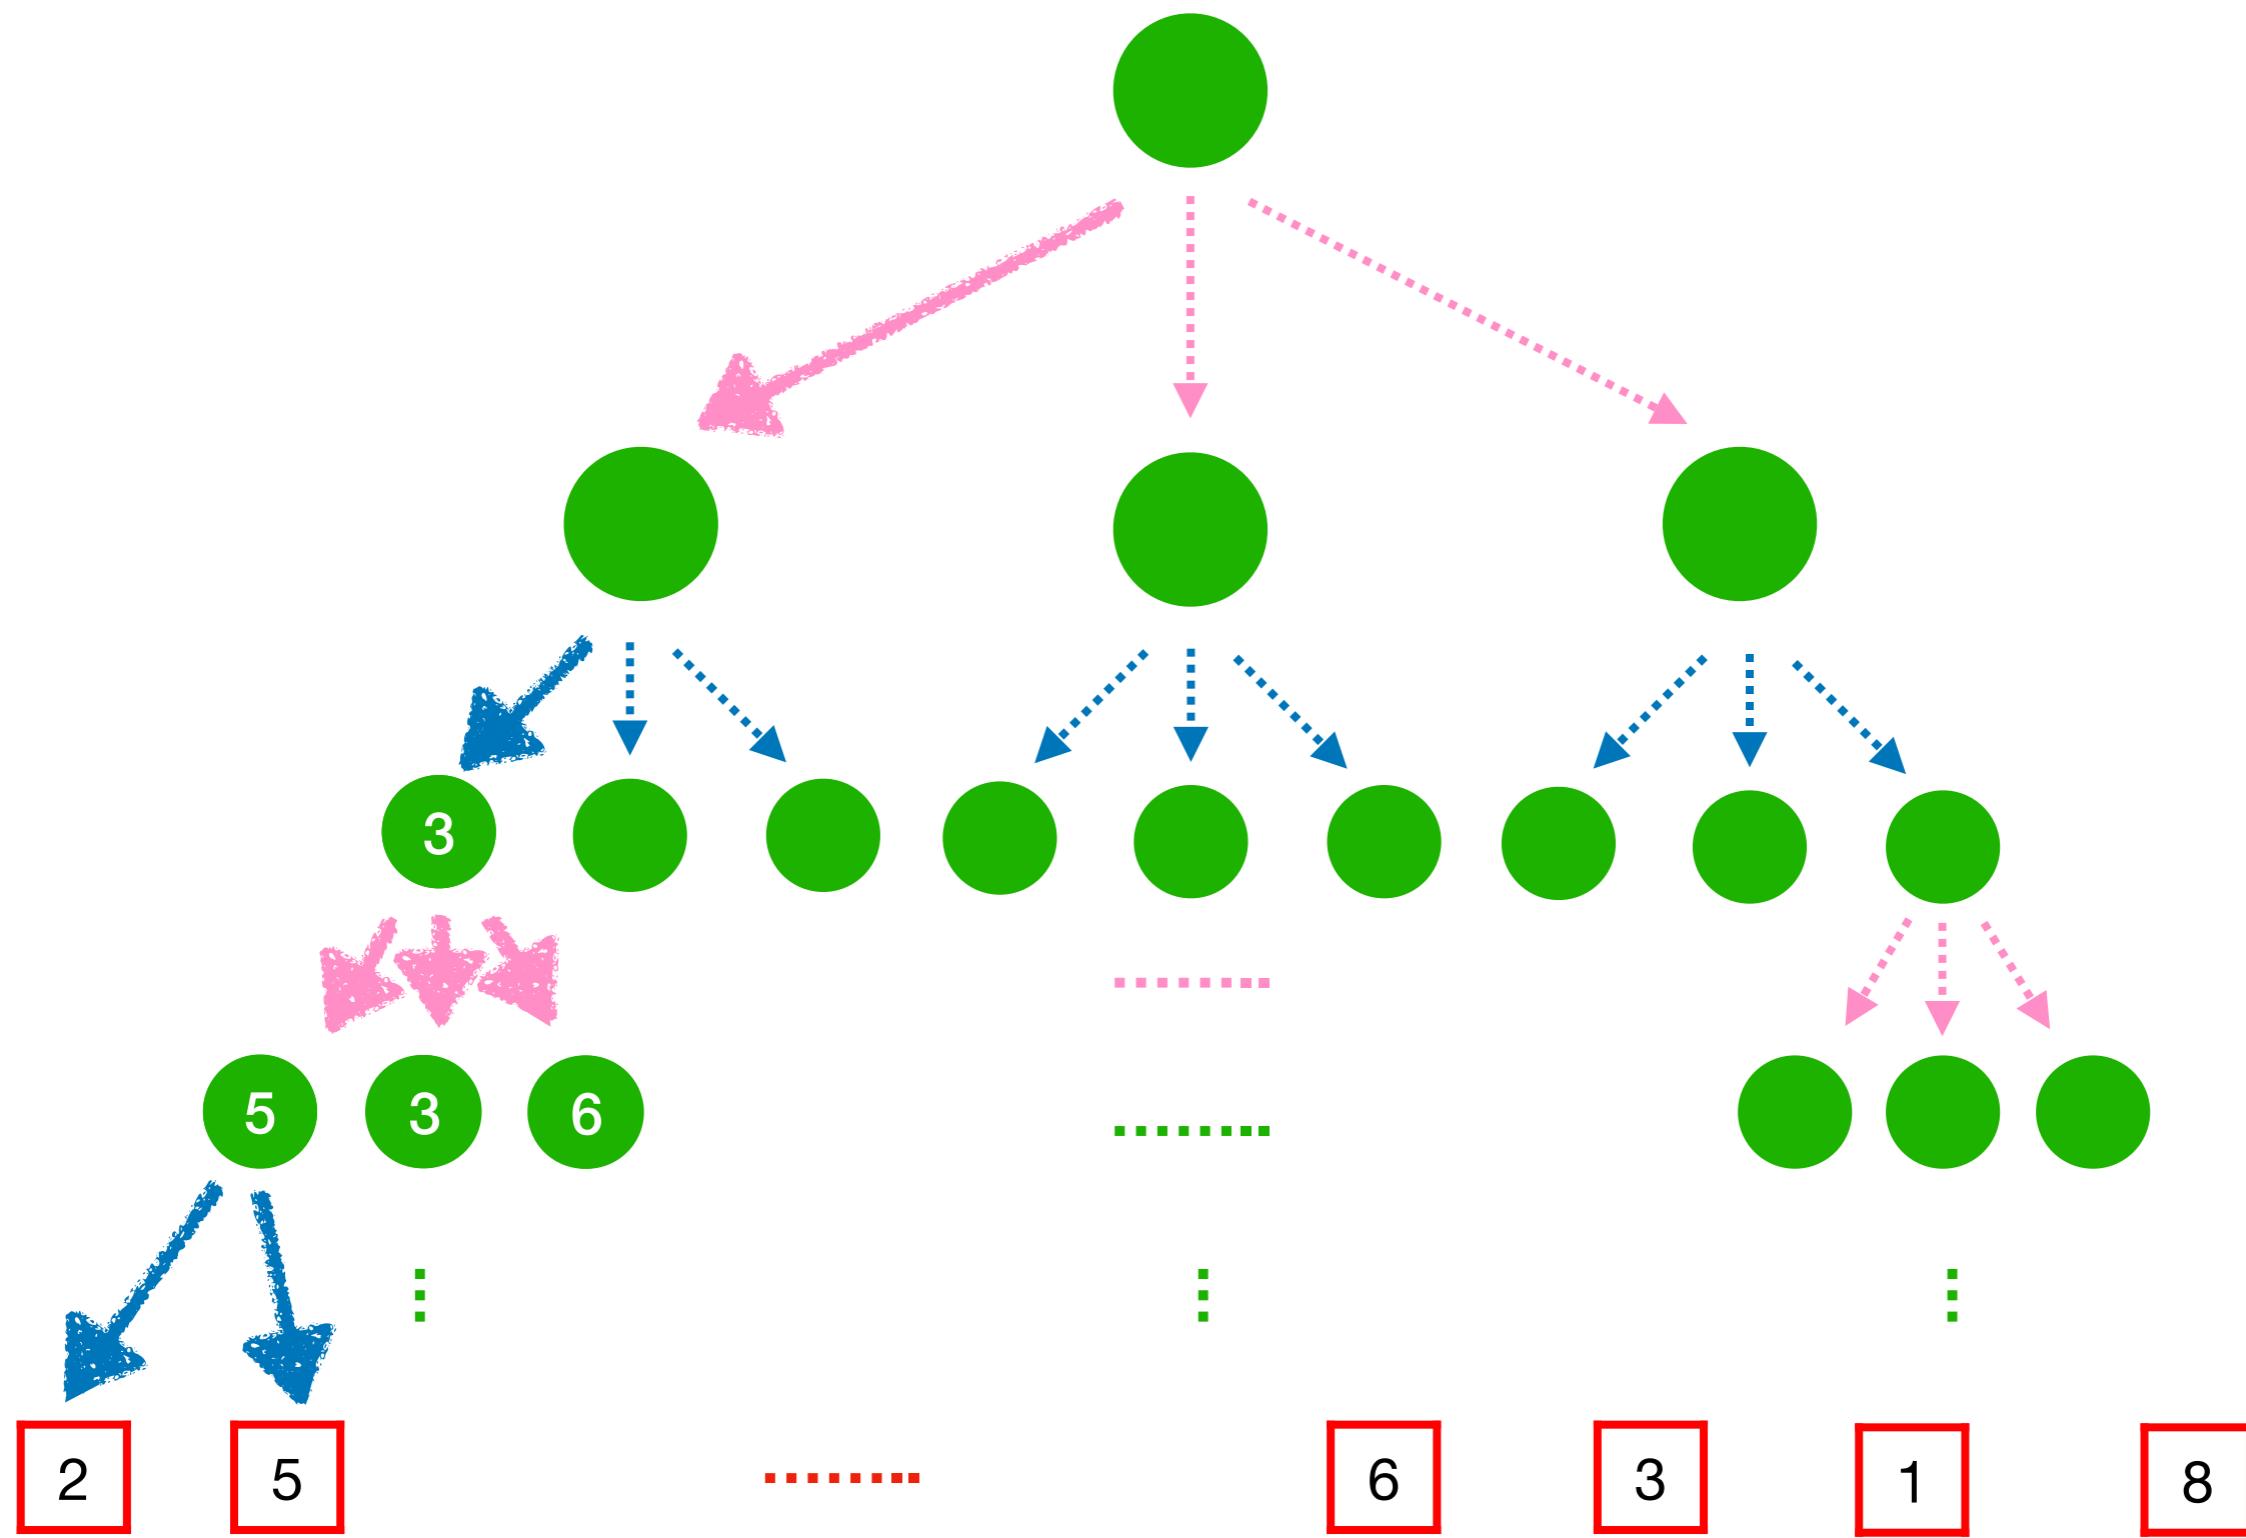
The game is deterministic and complete:
The search tree is known to the players.



Tic-tac-toe



Minimax Search

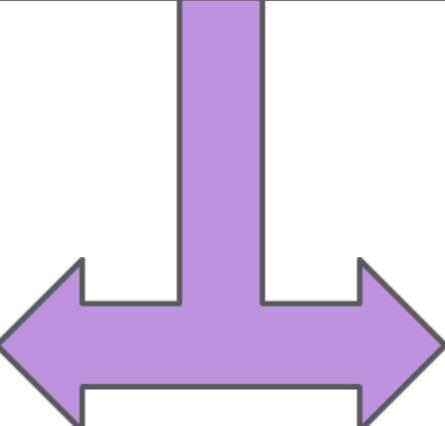


Minimax Search

```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is MIN: return min-value(state)
```

```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, value(successor))
    return v
```

```
def min-value(state):
    initialize v = +∞
    for each successor of state:
        v = min(v, value(successor))
    return v
```



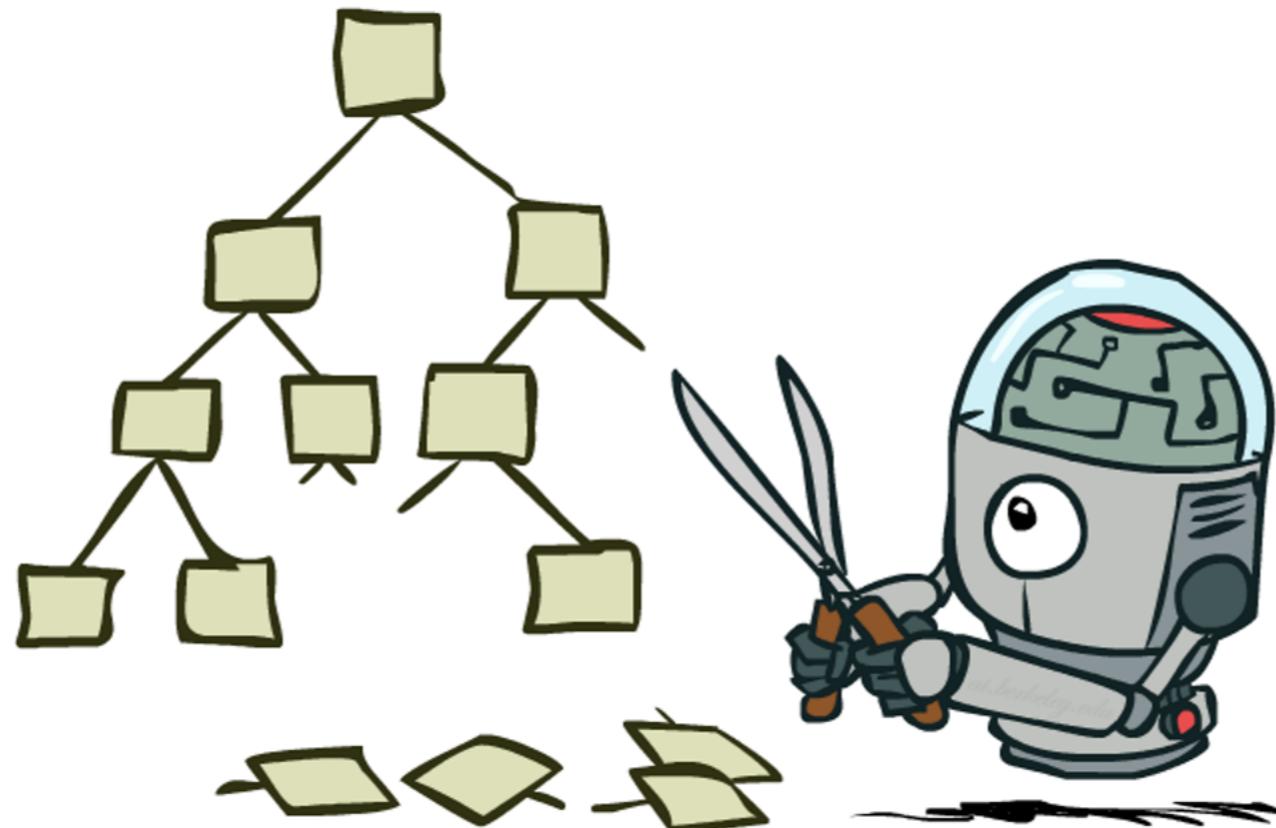
Works quite similar to depth-first search:

Search down then trace back
from left to right

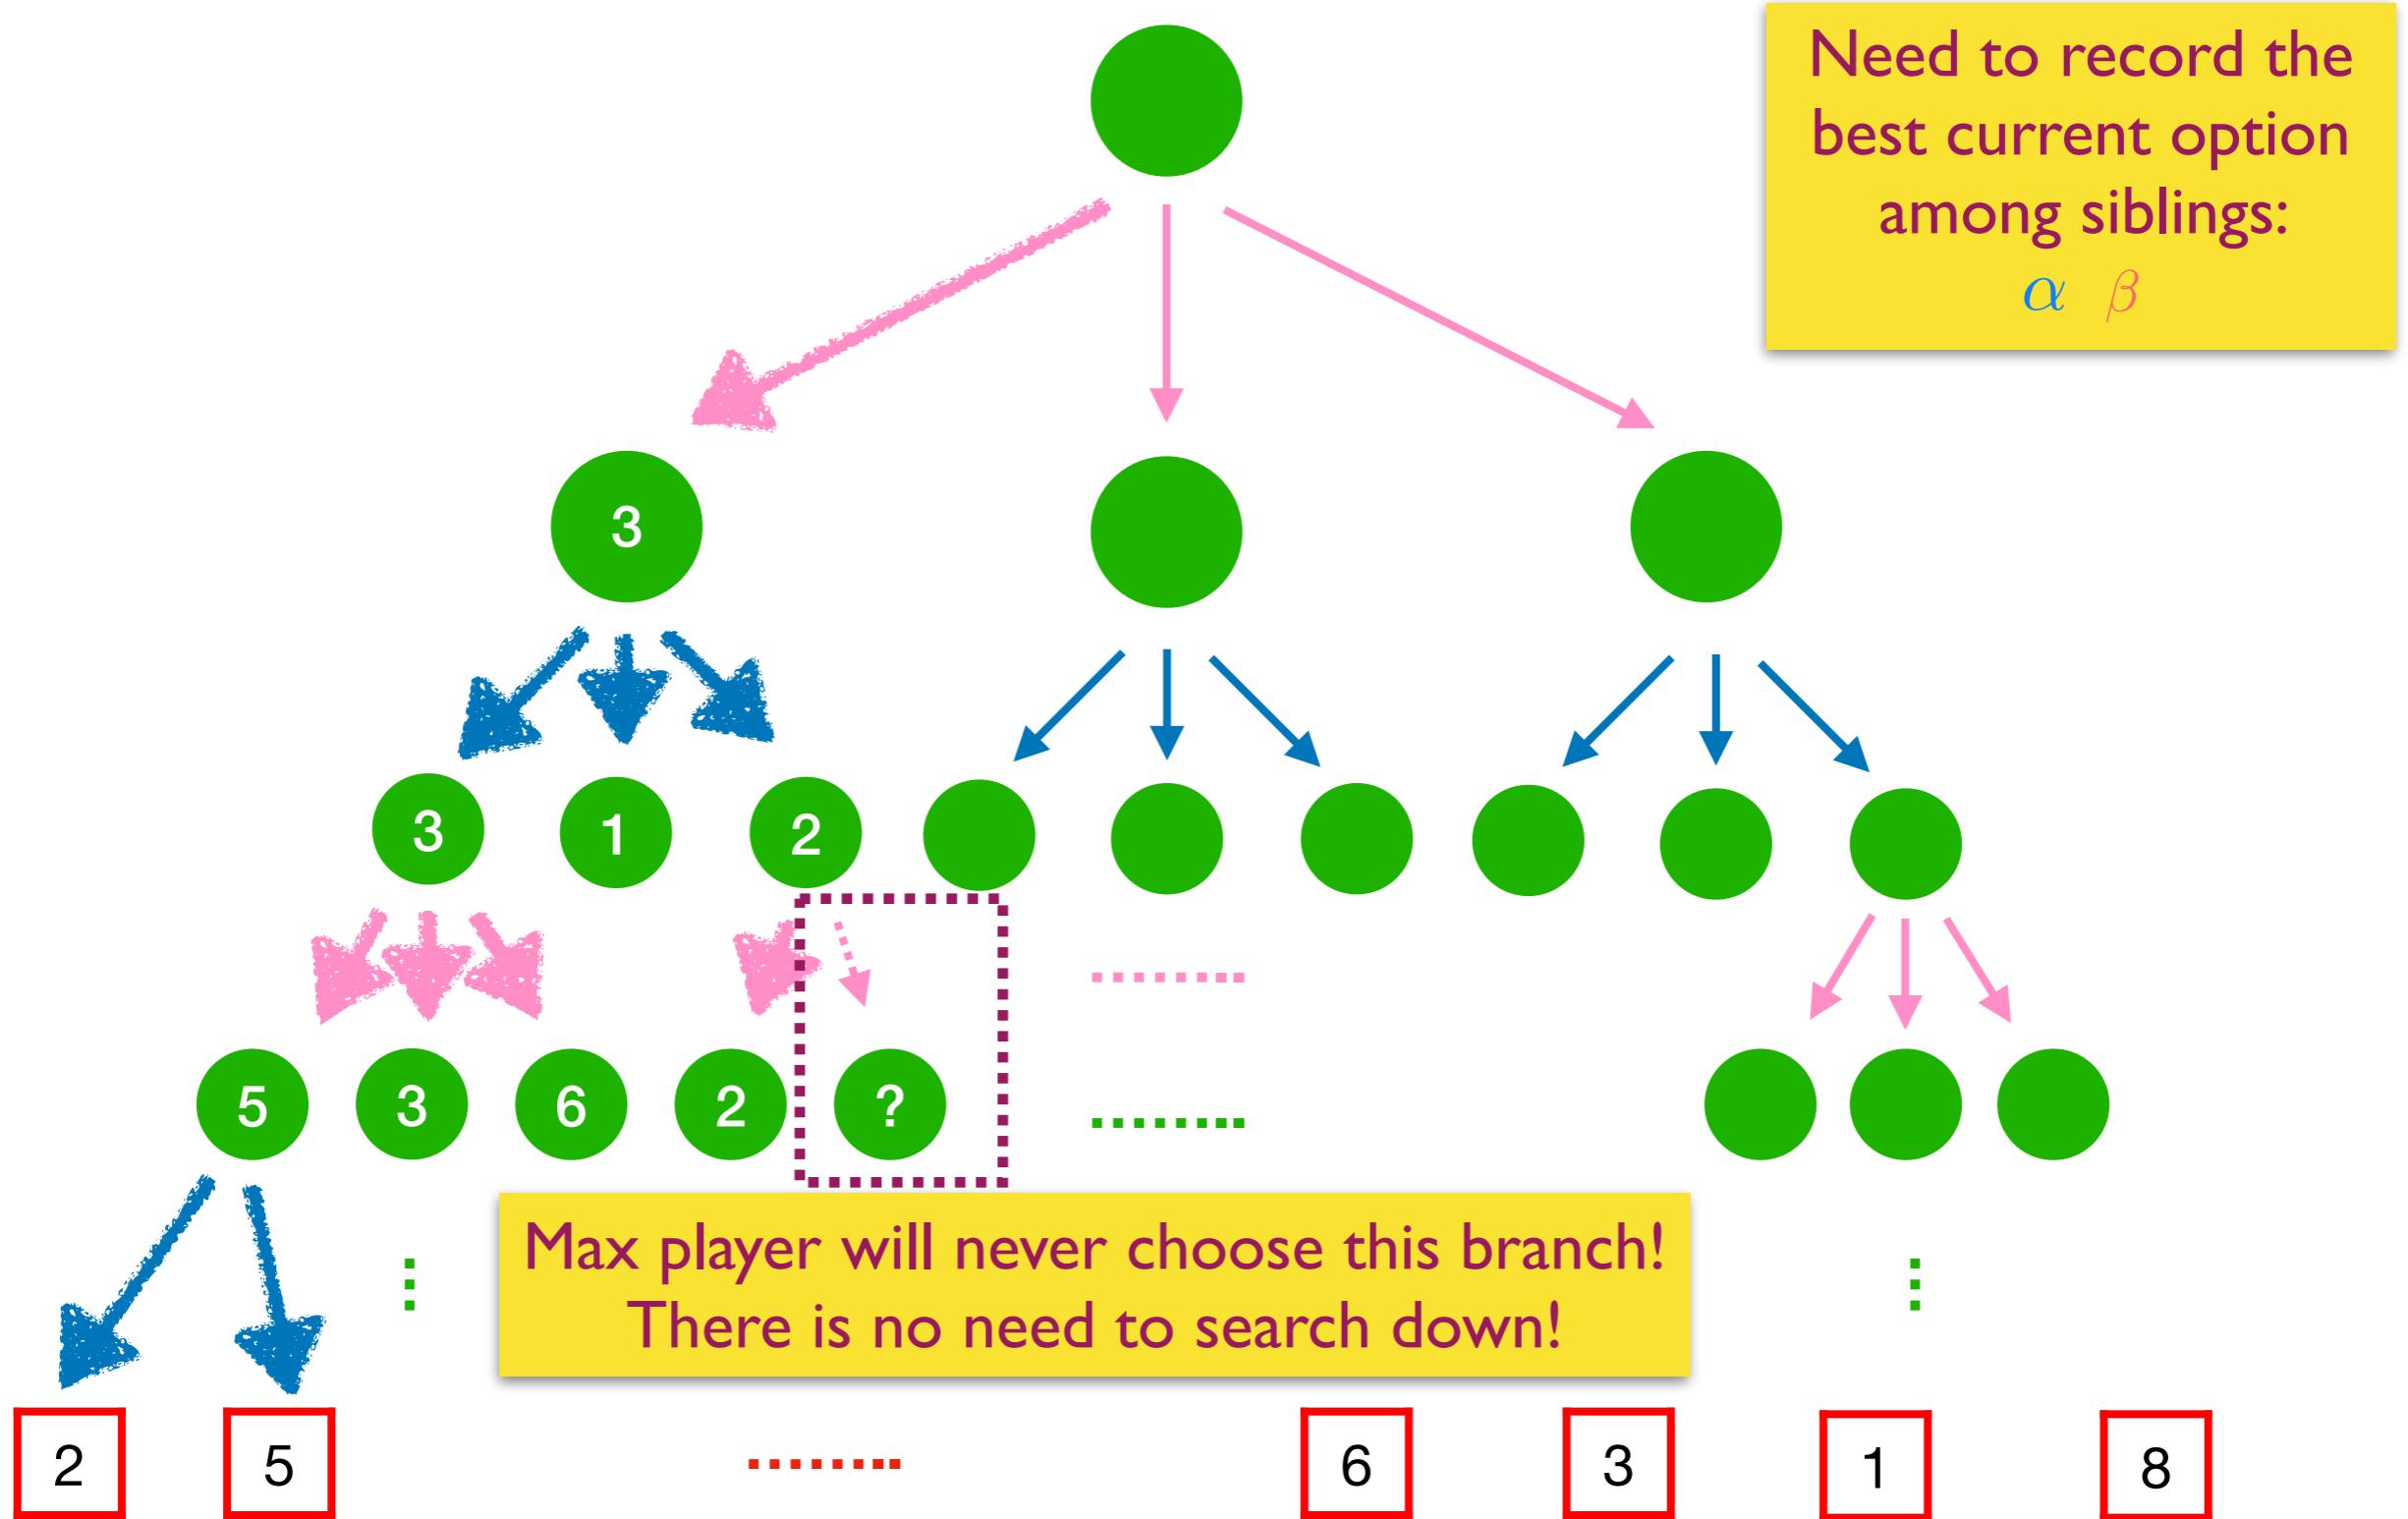
Similar time and space complexity to DFS

Can We Search More Efficiently?

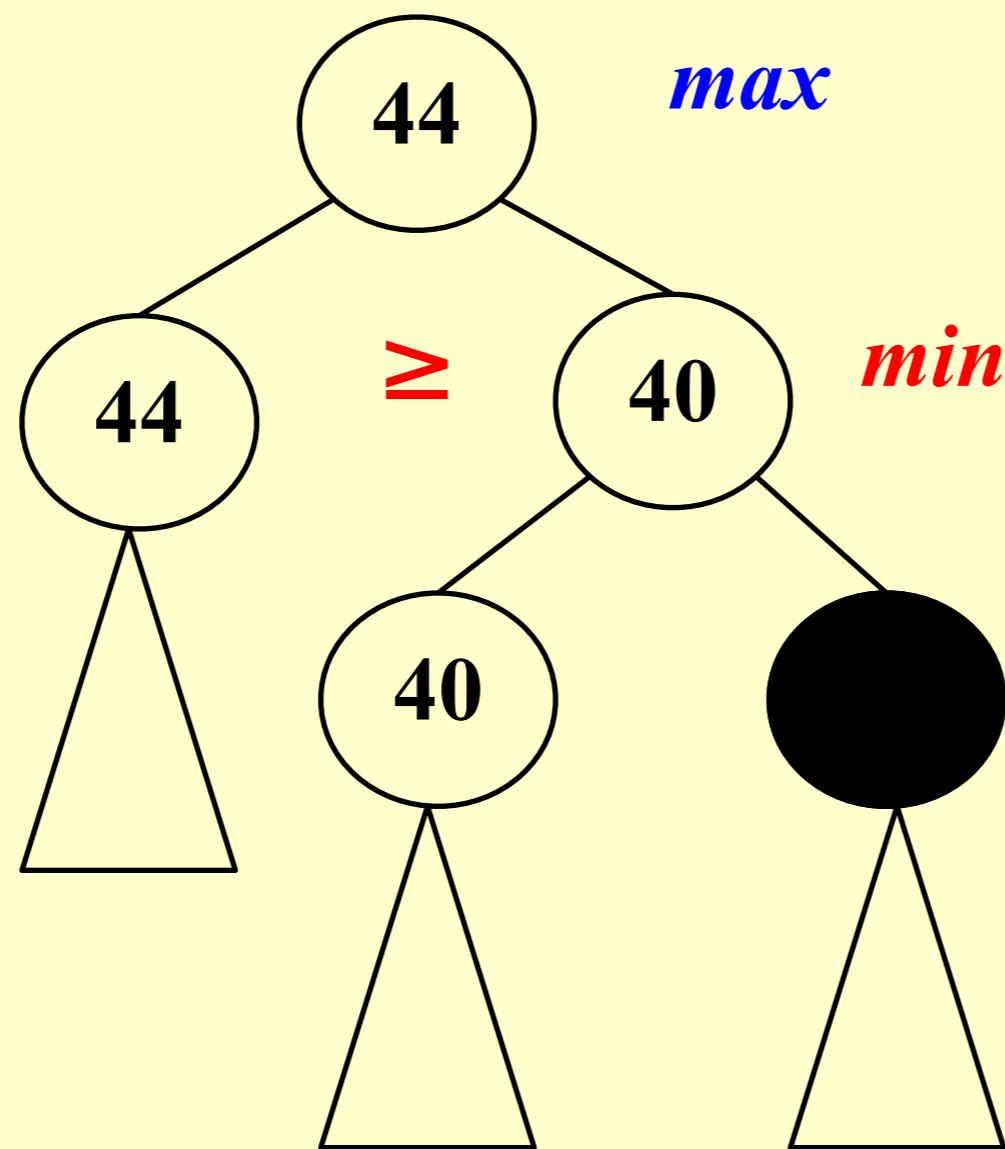
- Limit the search depth and use heuristic functions to estimate the final utility: e.g. how far away is the ghost?
- Prune the search tree.



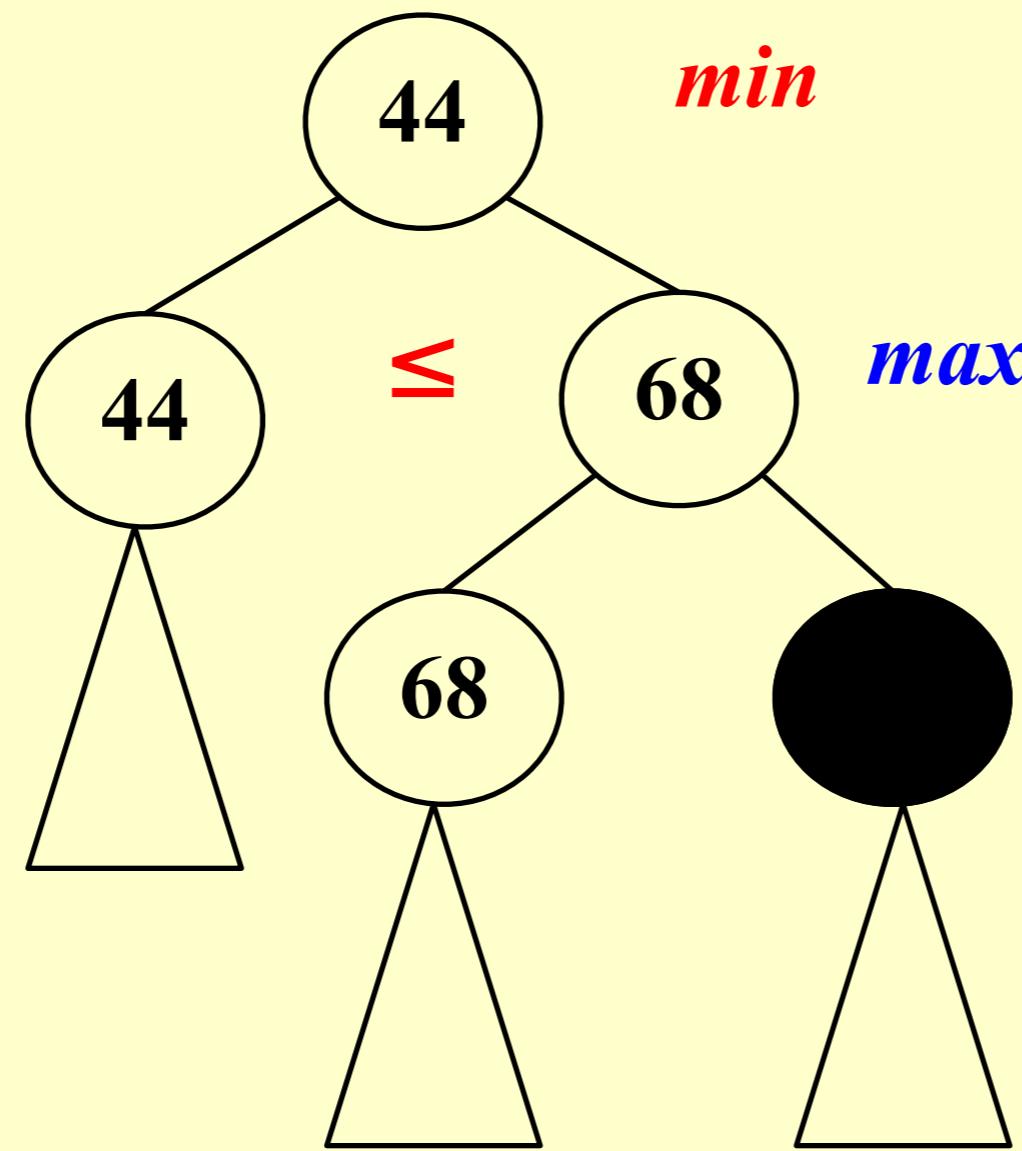
Alpha-Beta Pruning



α pruning



β pruning



α - β pruning: when both techniques are combined. In practice, it limits the searching to only $O(\sqrt{N})$ nodes, where N is the size of the full game tree.

Alpha-Beta Pruning

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):
```

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$ return v

$\alpha = \max(\alpha, v)$

 return v

 return since the min
 father will not choose
 this branch already

```
def min-value(state,  $\alpha$ ,  $\beta$ ):
```

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \leq \alpha$ return v

$\beta = \min(\beta, v)$

 return v

 return since the max
 father will not choose
 this branch already

Search order matters!

In the worst case, no pruning will be done with bad order!

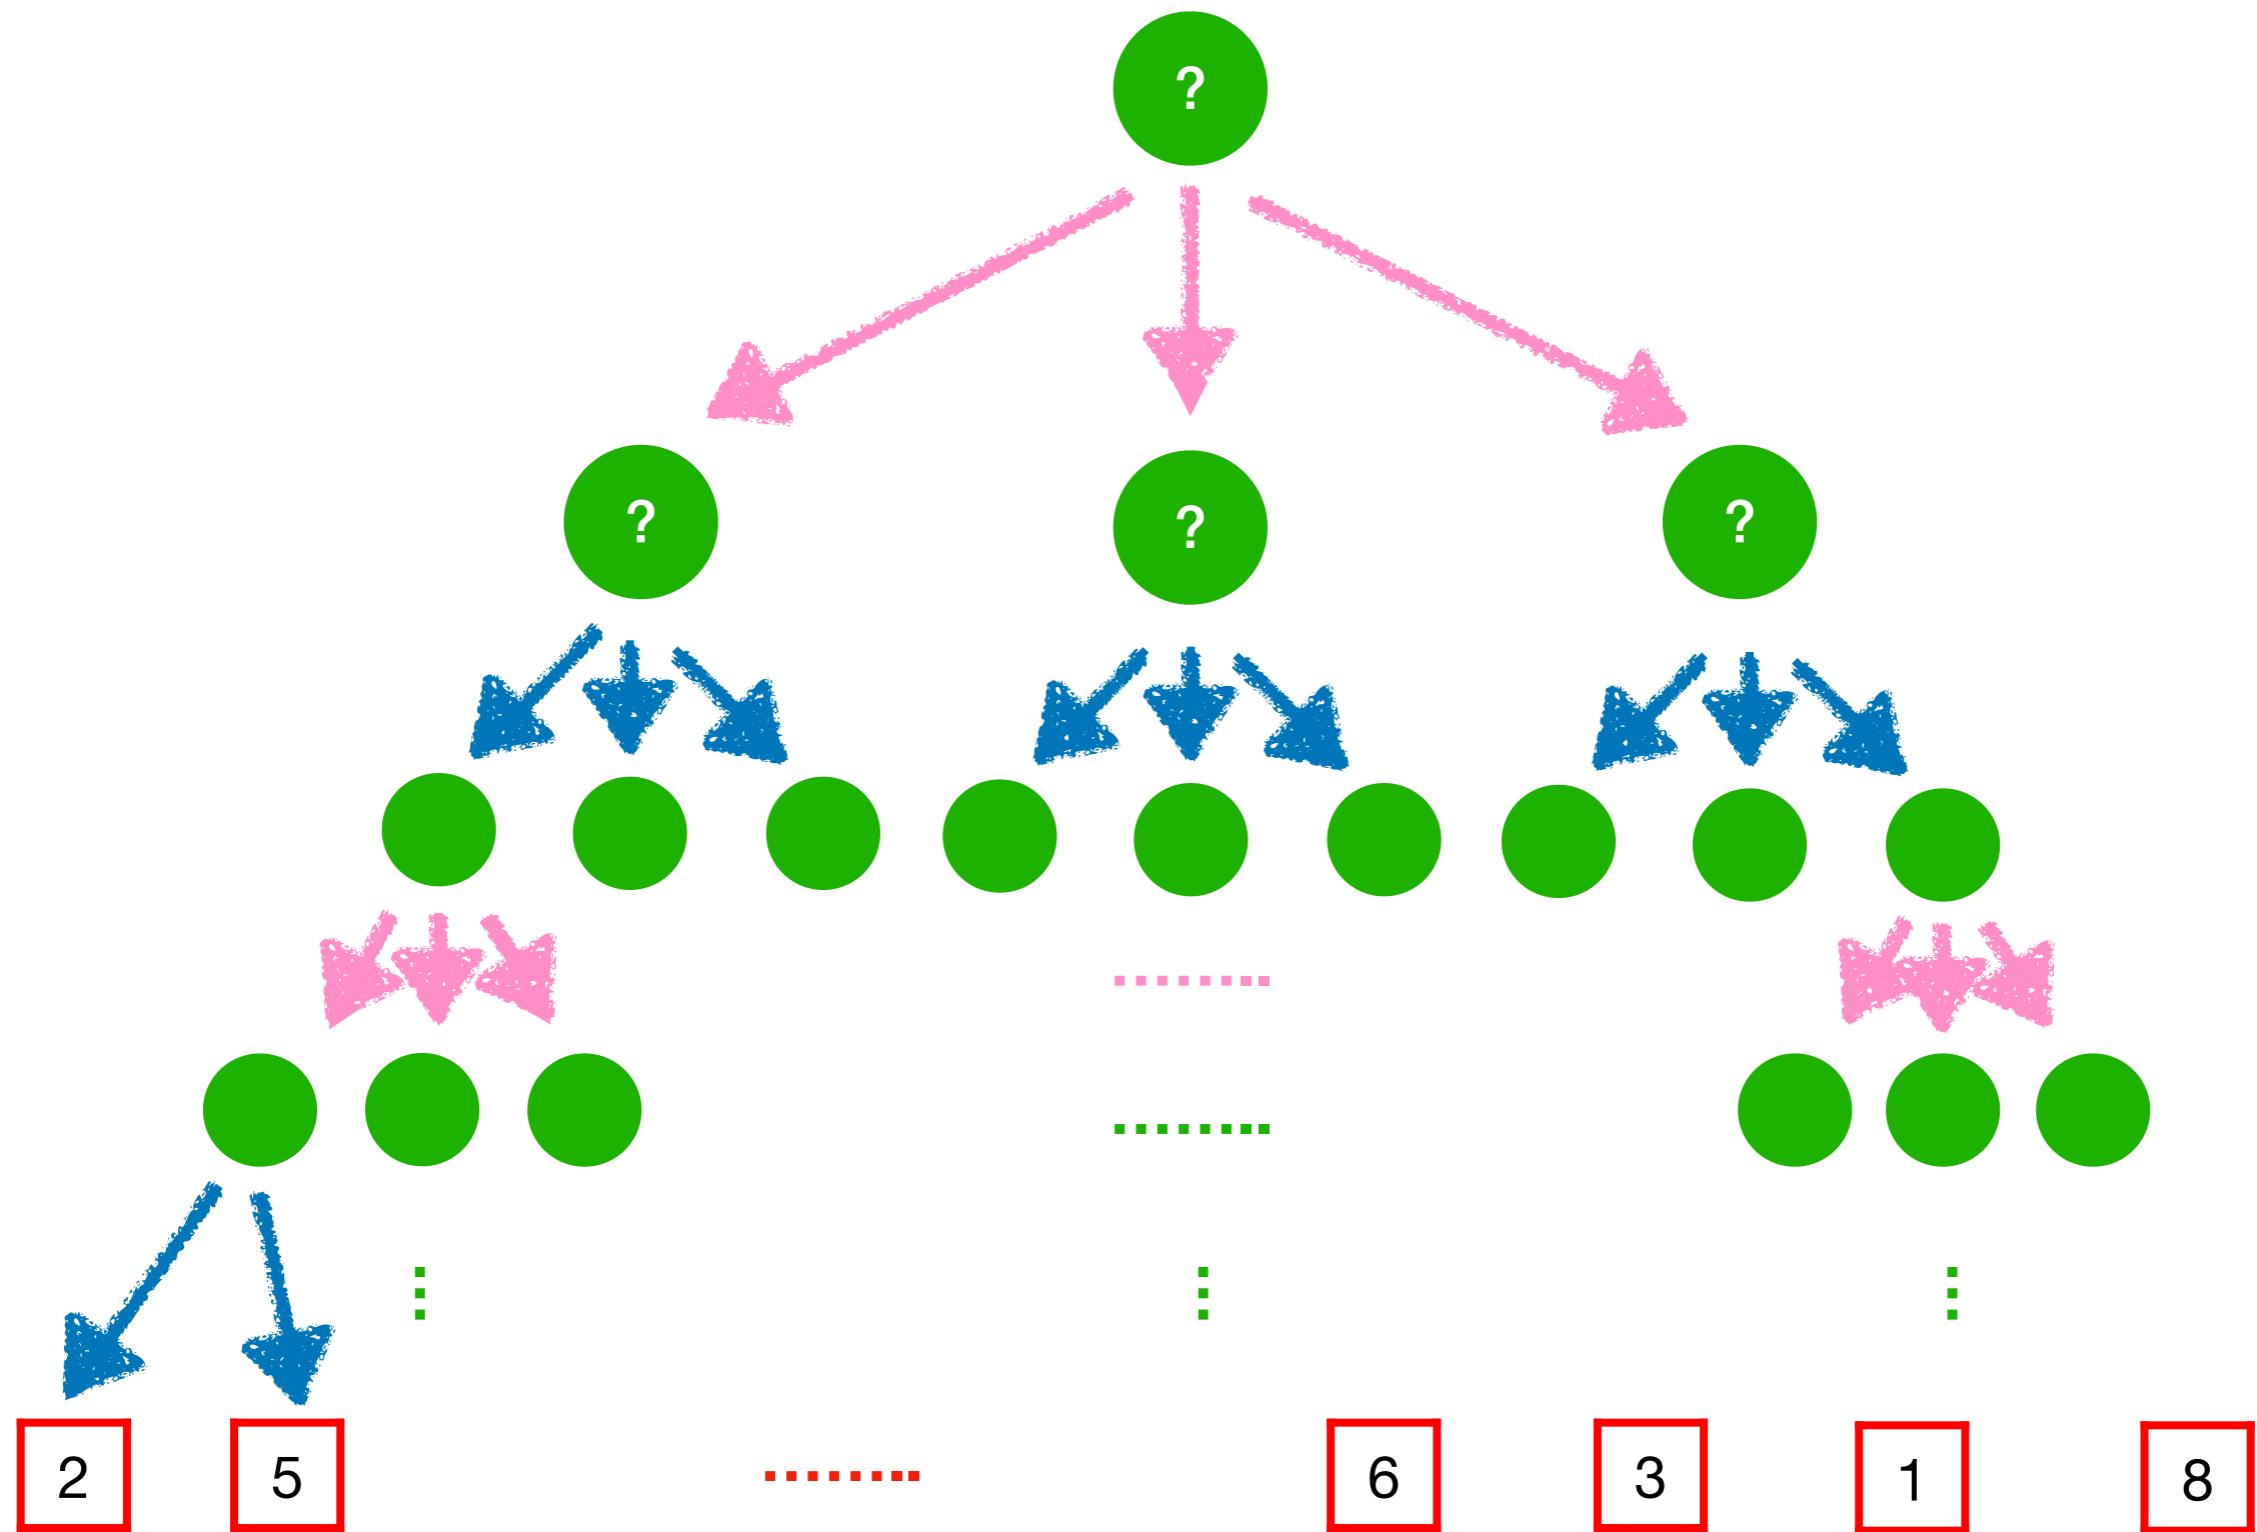
Deterministic search is not aware of tree structure.

Can we search more smartly?

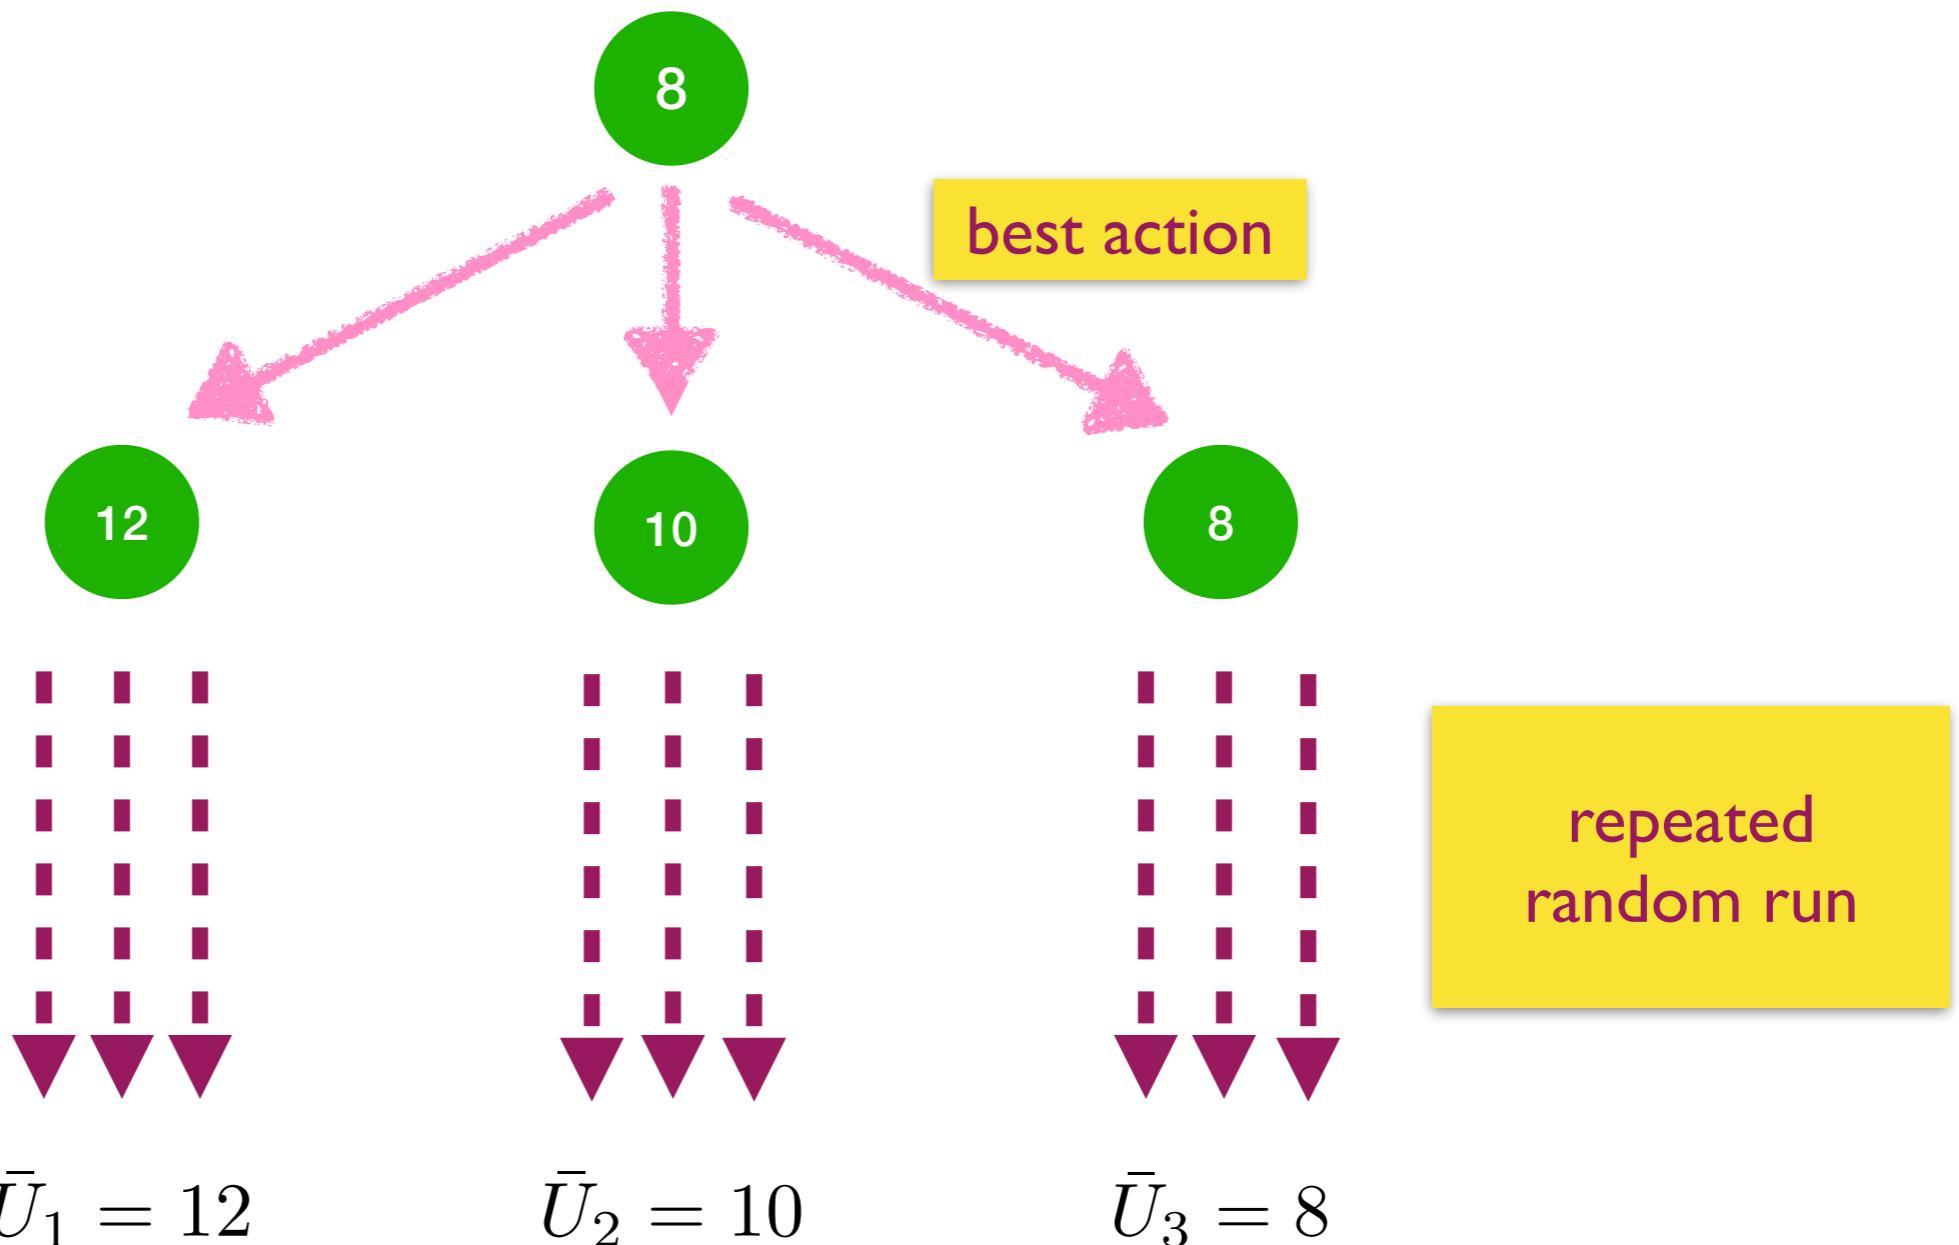
Search & Backtracking

- Fundamental search
- Search with backtracking
- **Search by simulation**
- Take-home messages

Suppose that the search tree is fully expanded,
how to evaluate each mode?



Monte-Carlo Simulation



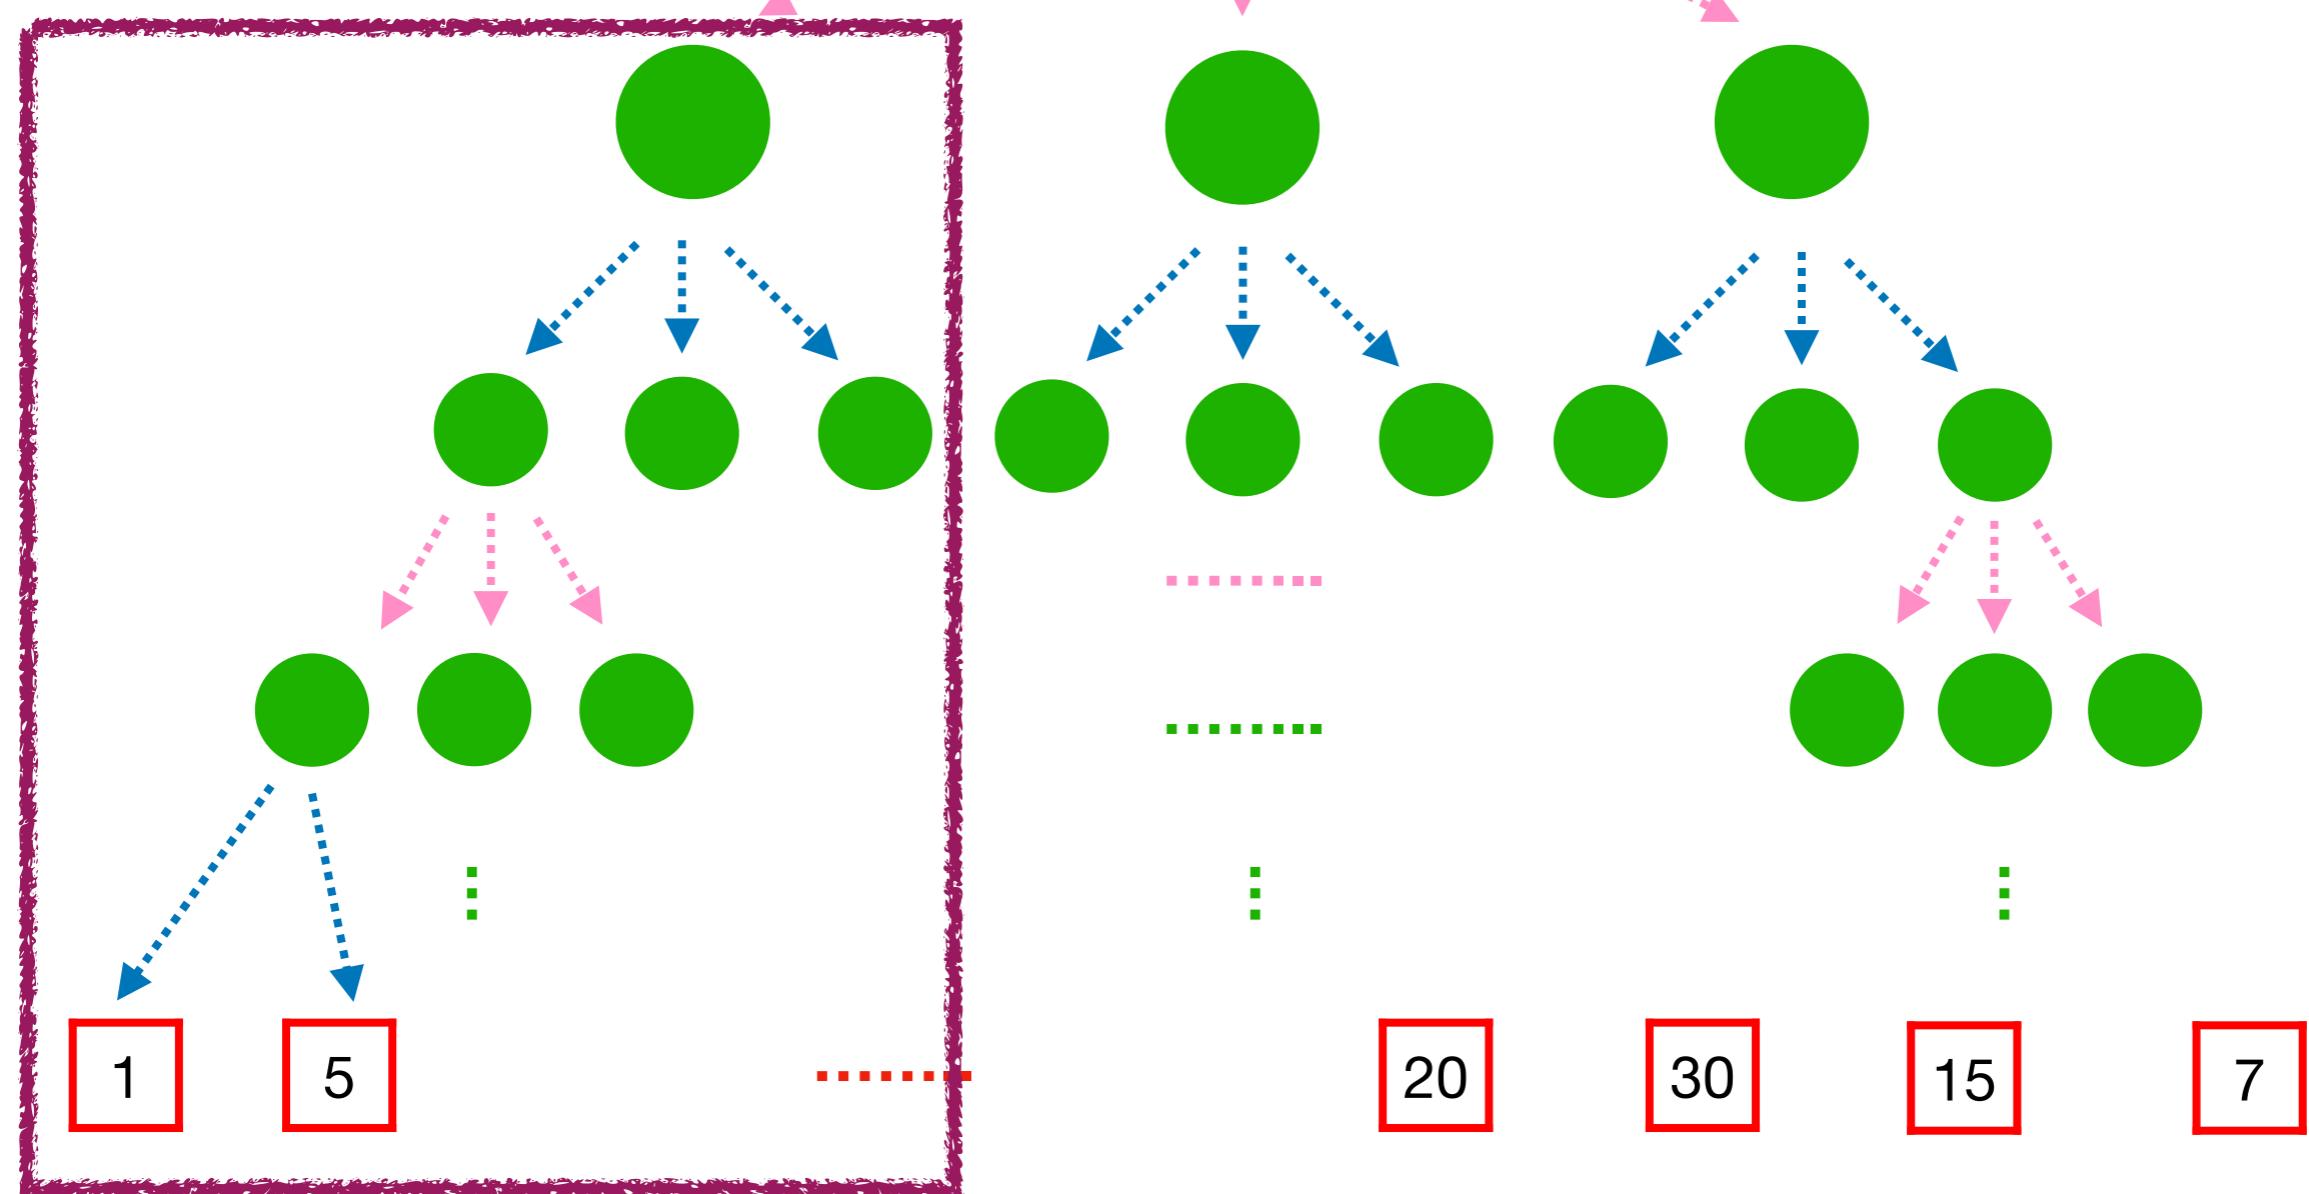
Converge to true utility when the #run is sufficient!

But in real game playing, the time and space for simulation is limited.
We need a smart strategy to decide the order of simulation.

The Order of Simulation Matters

Chicken & Egg problem:
How to know which is promising
when not simulated?

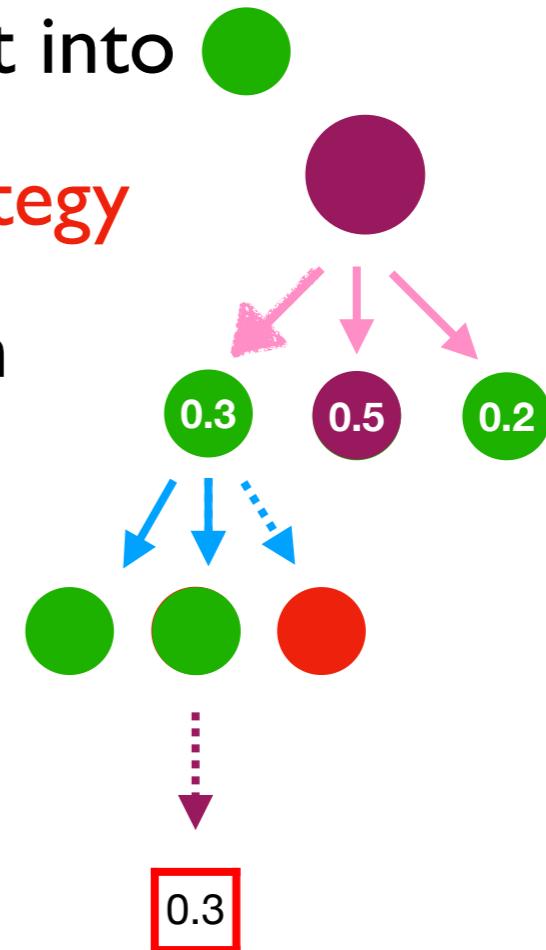
promising branch to simulate more



Monte-Carlo Tree Search (MCTS)

- Nodes of two kinds:  visited node with unexpanded child, and  unvisited node.
- During MCTS, we use two strategies: **tree** and **default**
- Algorithm: repeat until time or space limit:
 - Selection: choose one node among  using **tree strategy**
 - Expansion: expand an unvisited child and put into 
 - Simulation: simulate down using **default strategy**
 - Update: update MC estimation through path
- Output the best action to play

The **default strategy** is usually random play
The **tree strategy** is essential:
Deciding the order of search



AlphaGo: Integration of Machine Learning & Tree Search



Why is Go hard for computers to play?

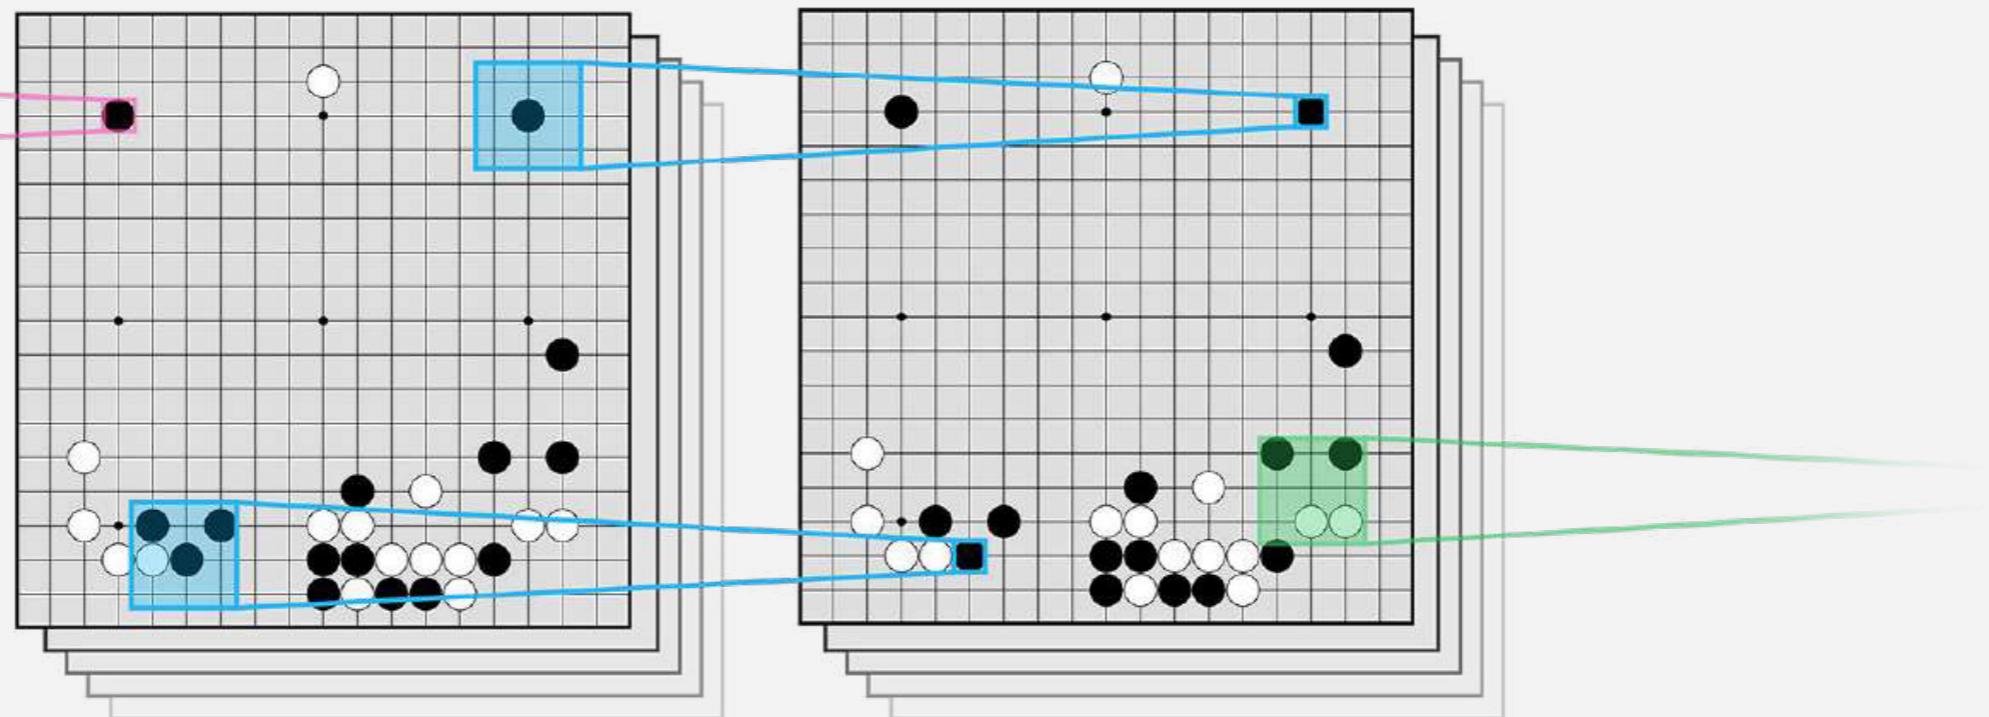
Game tree complexity = b^d

Brute force search intractable:

1. Search space is huge
2. “Impossible” for computers to evaluate who is winning



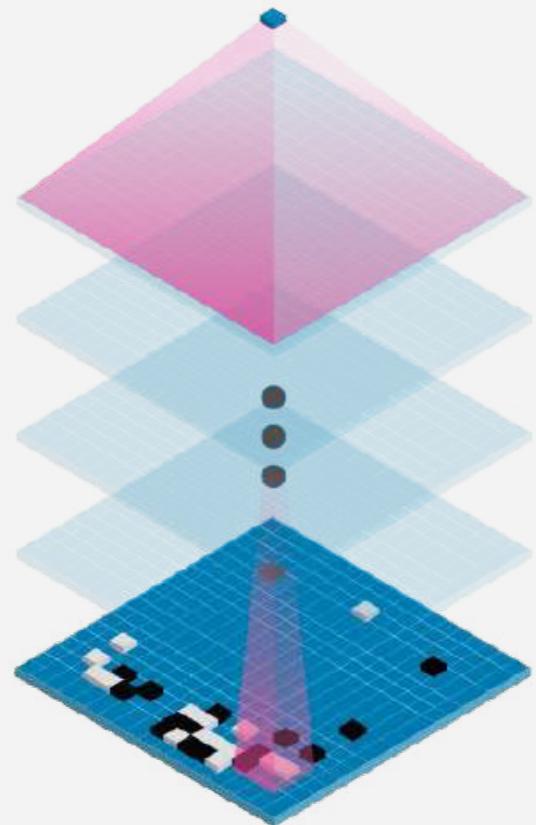
Convolutional neural network



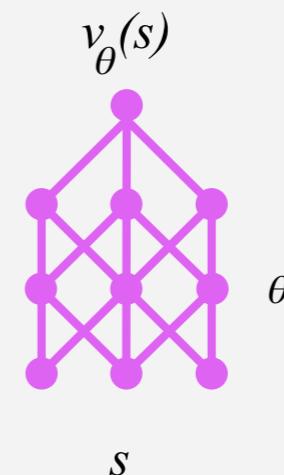
AlphaGo introduces three convolutional nets for learning.
Use images of board as state input.

Value network

Evaluation



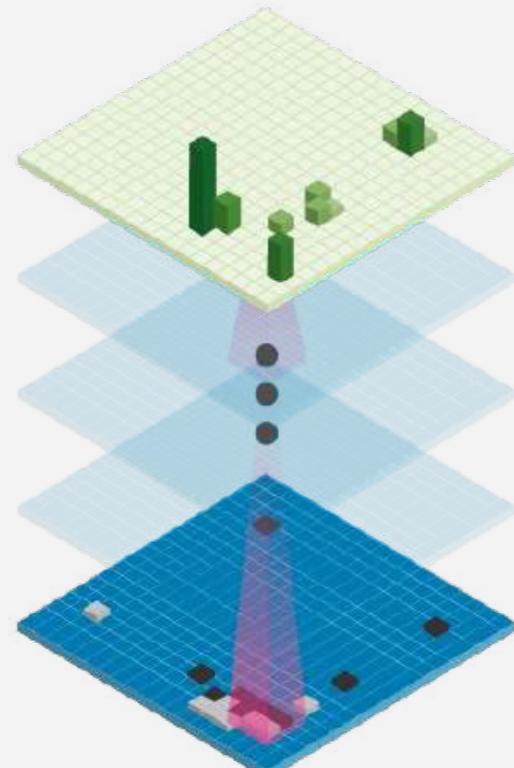
Position



The value function evaluates a state of the game.

Policy network

Move probabilities



Position

$$p_{\sigma}(a|s)$$

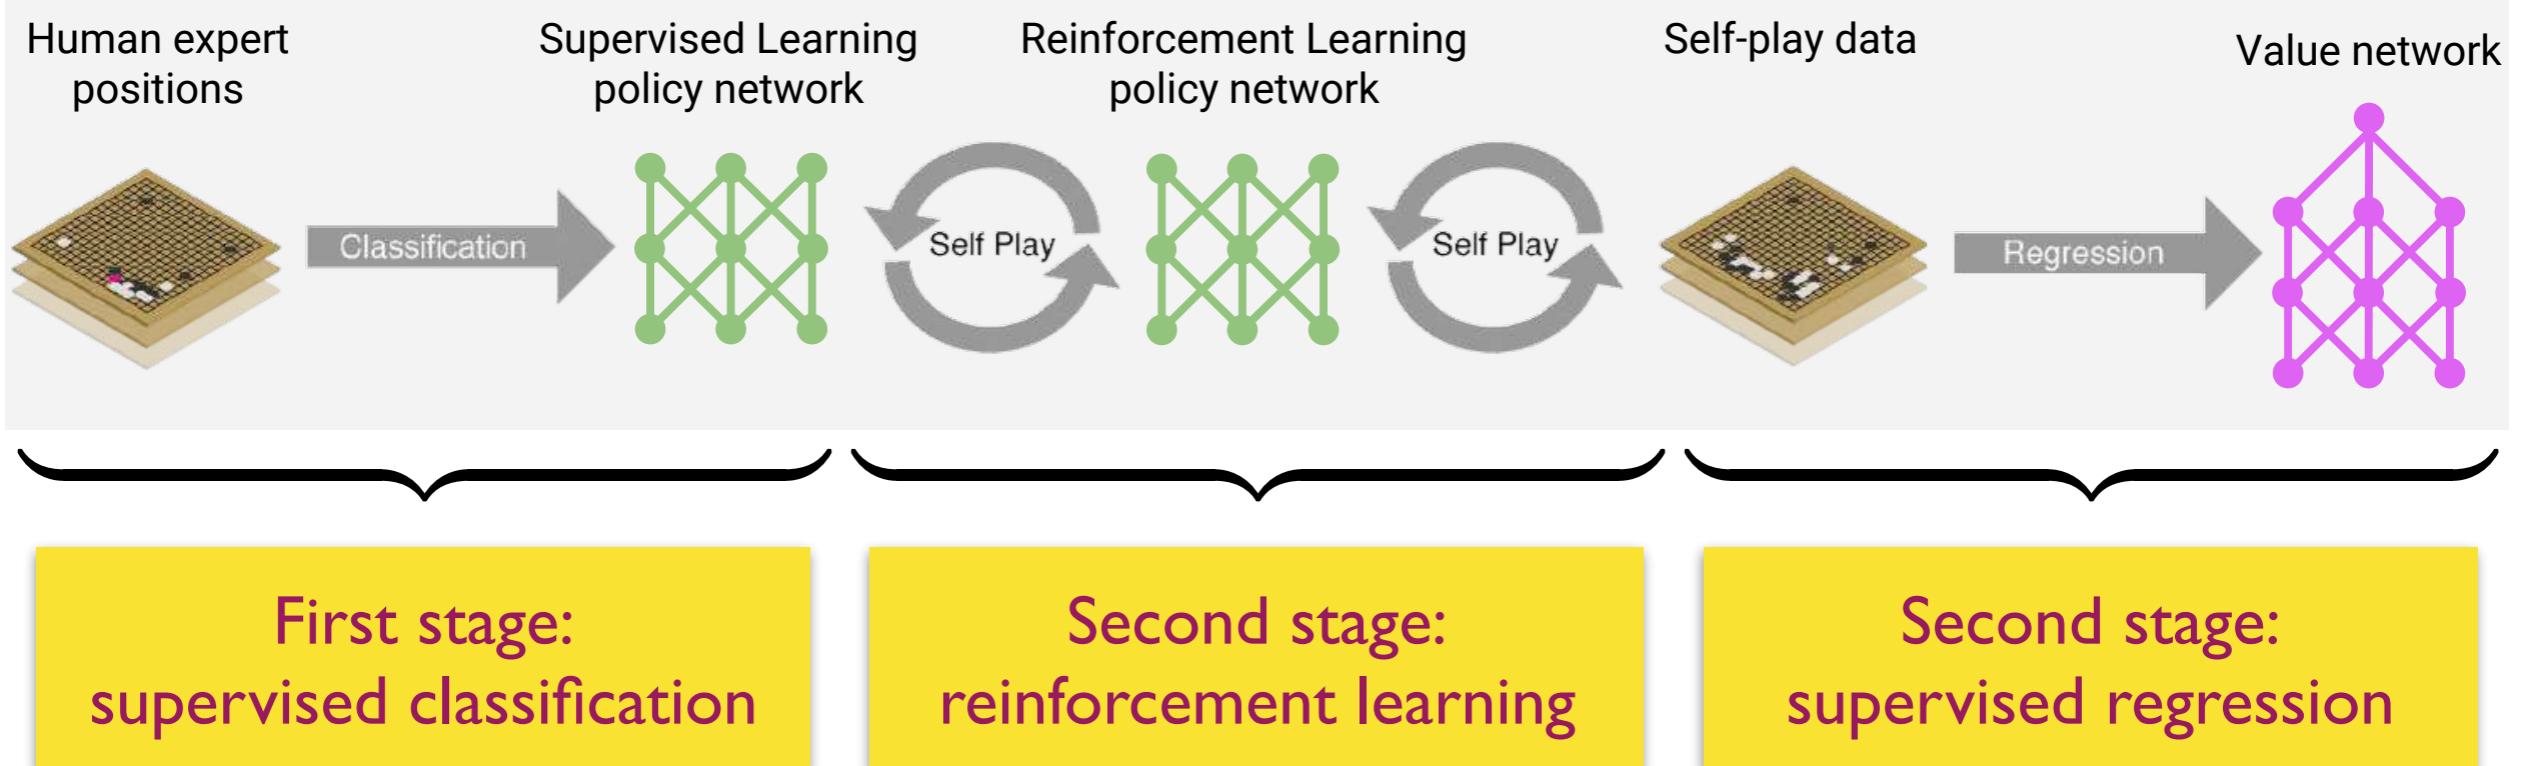
s

σ



The policy predicts which is better for the next step

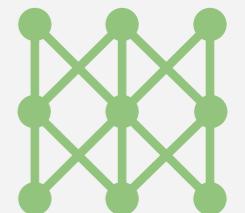
Neural network training pipeline



Learning: First Stage

Supervised learning of policy networks

Policy network: 12 layer convolutional neural network



Training data: 30M positions from human expert games (KGS 5+ dan)

Training algorithm: maximise likelihood by stochastic gradient descent

$$\Delta\sigma \propto \frac{\partial \log p_\sigma(a|s)}{\partial \sigma}$$

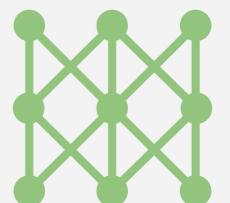
Training time: 4 weeks on 50 GPUs using Google Cloud

Results: 57% accuracy on held out test data (state-of-the art was 44%)

Learning: Second Stage

Reinforcement learning of policy networks

Policy network: 12 layer convolutional neural network



Training data: games of self-play between policy network

Training algorithm: maximise wins z by policy gradient reinforcement learning

$$\Delta\sigma \propto \frac{\partial \log p_\sigma(a|s)}{\partial \sigma} z$$

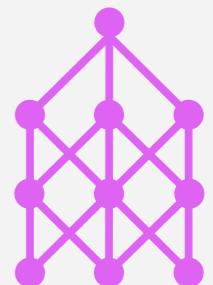
Training time: 1 week on 50 GPUs using Google Cloud

Results: 80% vs supervised learning. Raw network ~3 amateur dan.

Learning: Third Stage

Reinforcement learning of value networks

Value network: 12 layer convolutional neural network



Training data: 30 million games of self-play

Training algorithm: minimise MSE by stochastic gradient descent

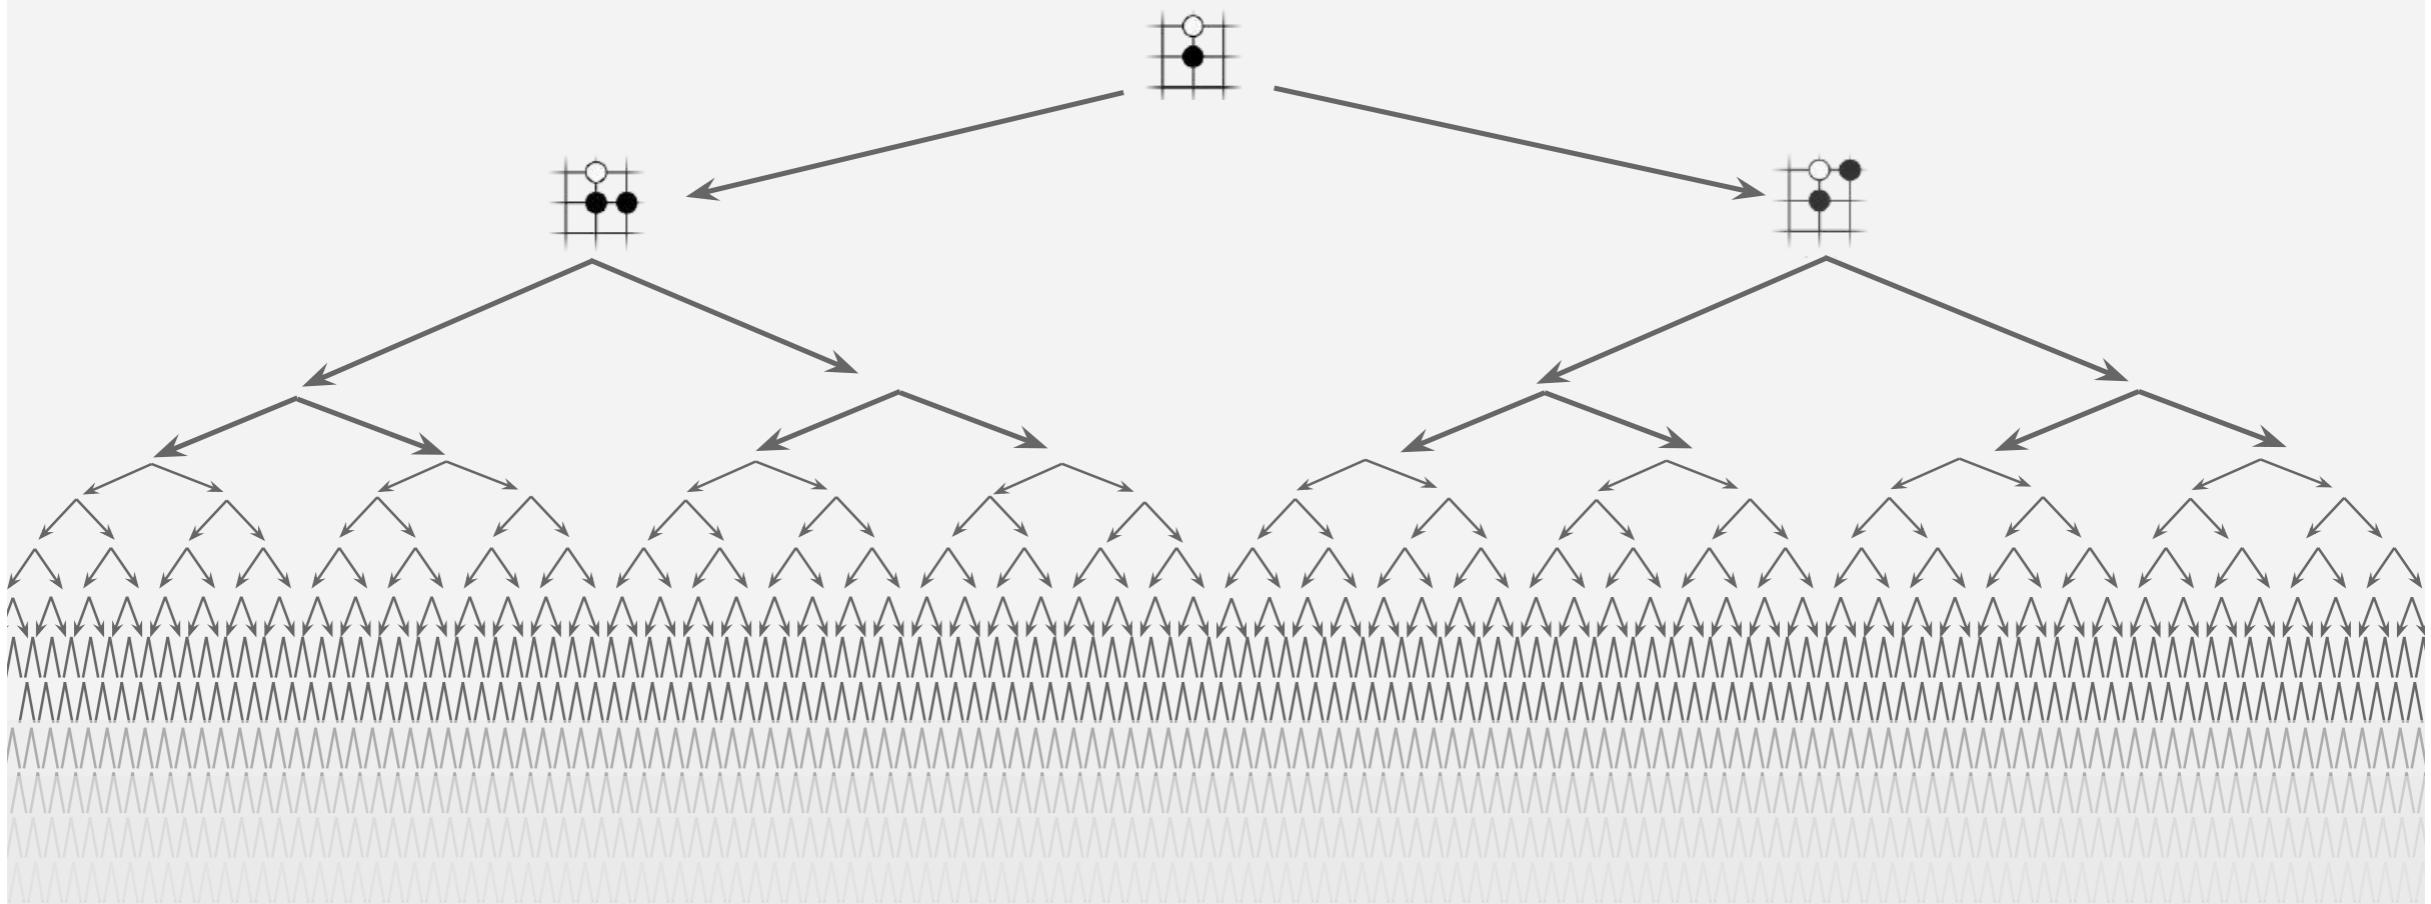
$$\Delta\theta \propto \frac{\partial v_\theta(s)}{\partial \theta} (z - v_\theta(s))$$

Training time: 1 week on 50 GPUs using Google Cloud

Results: First strong position evaluation function - previously thought impossible

Real Play: MCTS

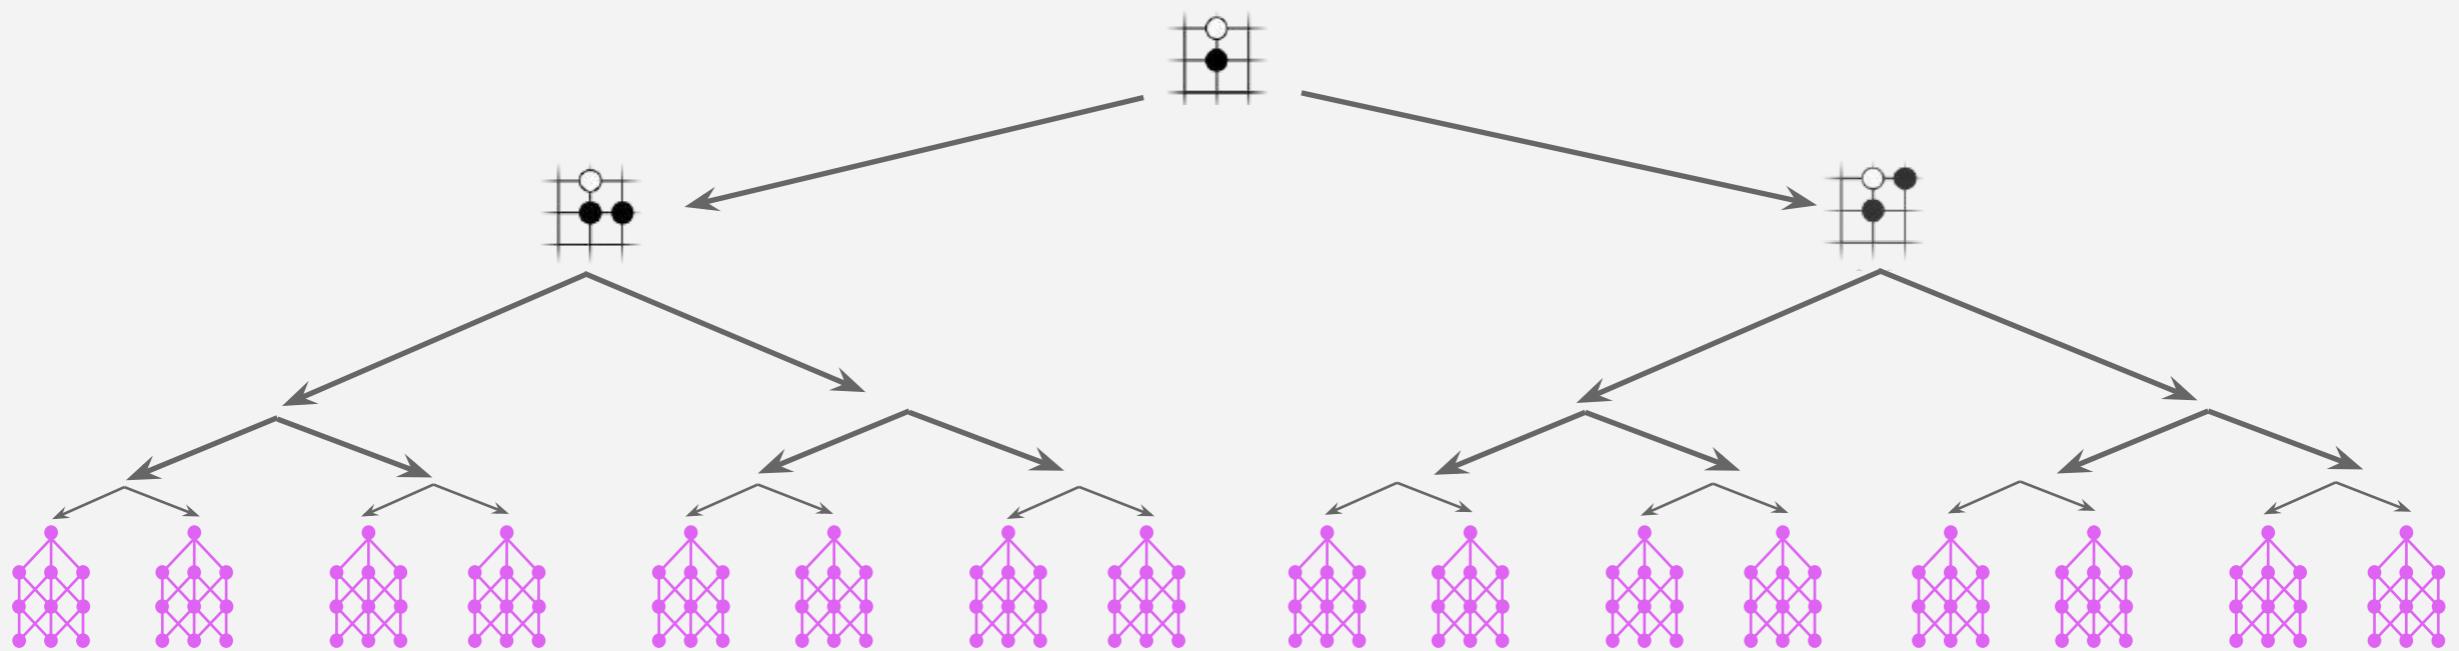
Exhaustive search



Two key steps:
node expansion and repeated random simulation

Real Play: MCTS

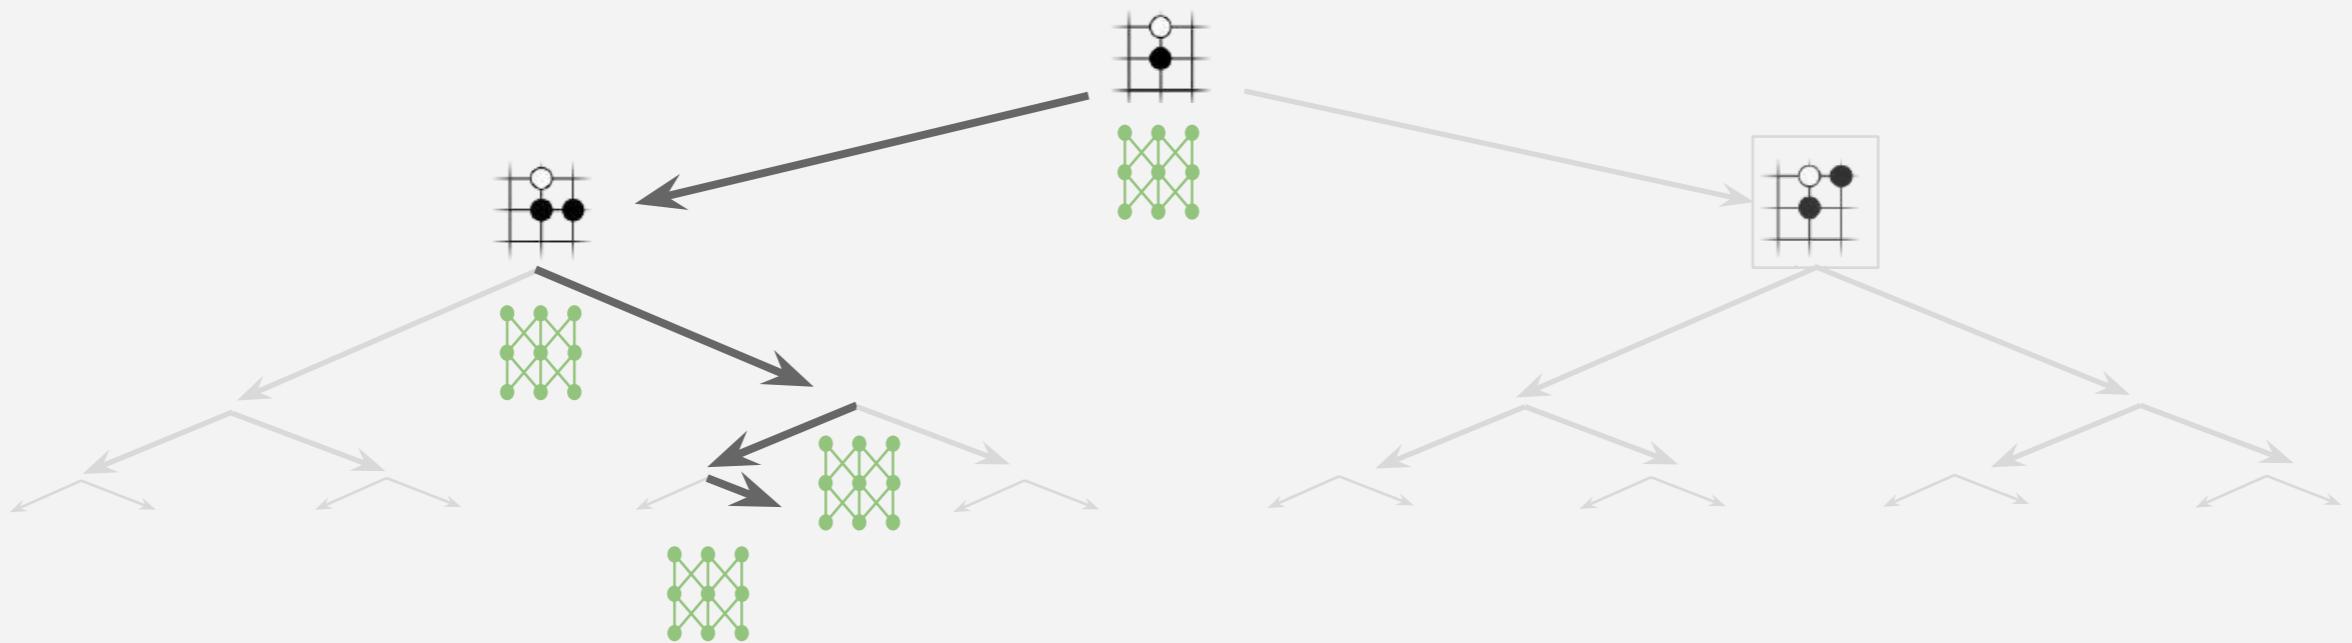
Reducing depth with value network



With value network, we can expand fewer depth
since the value of nodes can also be obtained from the value network.

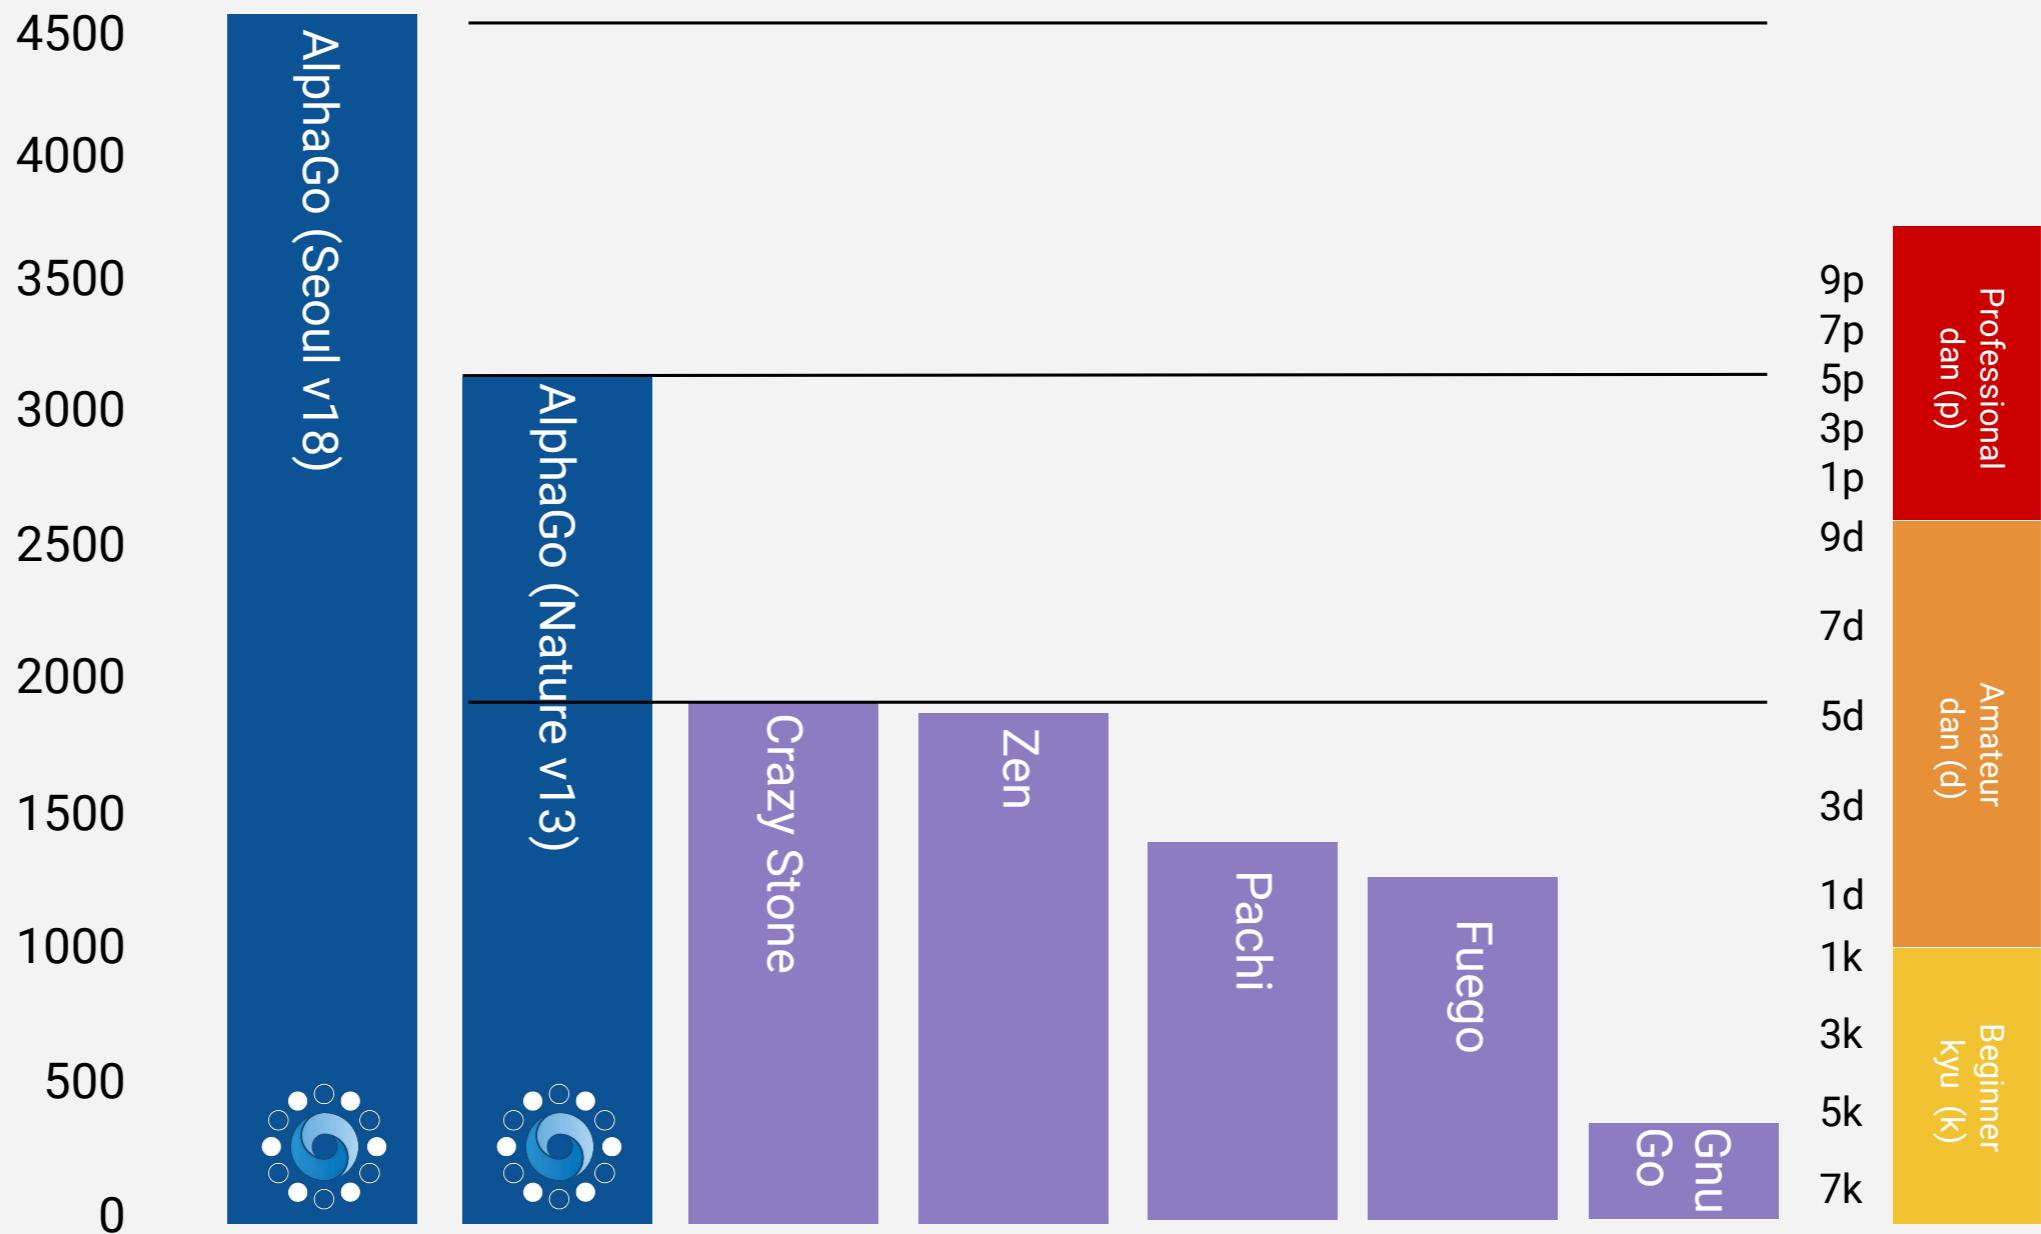
Real Play: MCTS

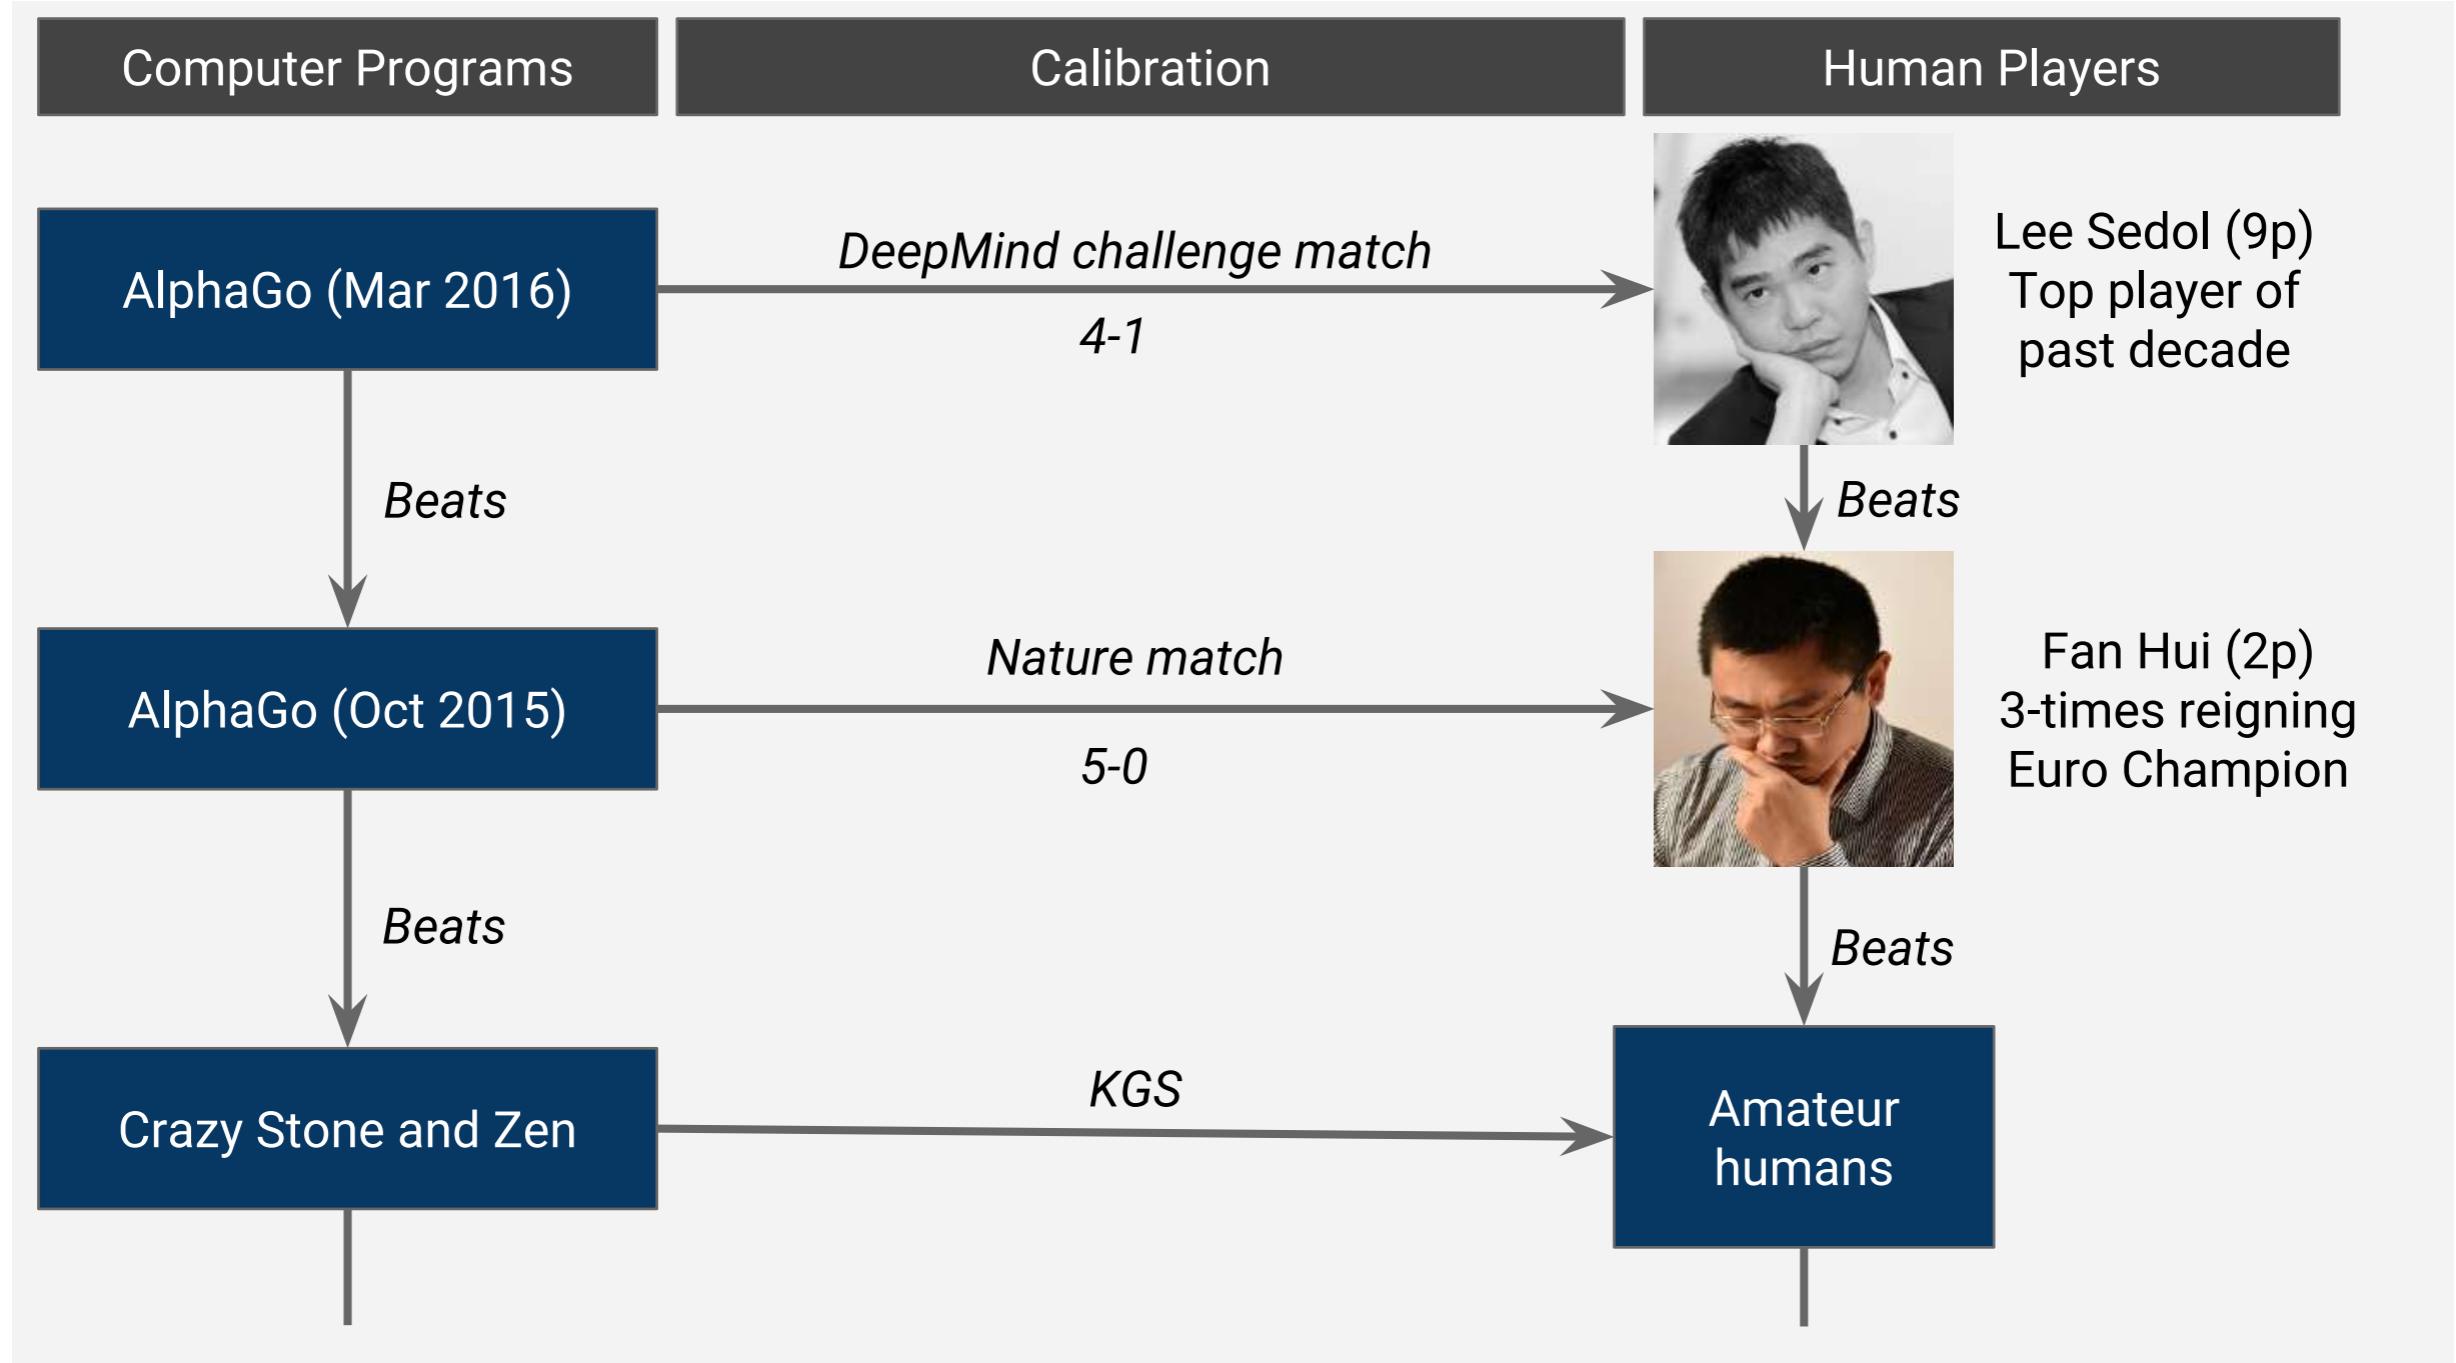
Reducing breadth with policy network



With policy network, we can simulate with fewer times
since the simulation can be guided by policy network over random play.

Evaluating AlphaGo against computers





Search & Backtracking

- Fundamental search
- Search with backtracking
- Search by simulation
- Take-home messages

Search & Backtracking

- Backtracking: search + pruning w.r.t. problem constraints.
- Minimax search and alpha-beta pruning: deterministic search without benefiting from the structure of the search tree.
- Search by simulation: (*optional to learn)
 - Monte-Carlo tree search: benefit from exploration of tree structure and random simulation.

Thanks for your attention!
Discussions?

Reference

Data Structure and Algorithm Analysis in C (2nd Edition): Sec. 10.5.

Artificial Intelligence:A Modern Approach (4th edition): Chap. 3, Sec.5.1-5.4.

http://ai.berkeley.edu/lecture_slides.html (Lecture 2, 3, 6)