

Advanced Data Structures and Algorithm Analysis

丁尧相
浙江大学

Fall and Winter 2025
Lecture 8

Inverted File Index

- Inverted File Index
- Take-home messages

Inverted File Index

- Inverted File Index
- Take-home messages

AVL Trees

AVL Trees

Splay Trees

AVL Trees

Splay Trees

Red-Black Trees

AVL Trees

Splay Trees

Red-Black Trees

B+ Trees

AVL Trees

Splay Trees

Red-Black Trees

B+ Trees



1 | Google

1 - eBizMBA Rank | **1,800,000,000** - Estimated Unique Monthly Visitors | 1 - Quantcast Rank | 1 - Alexa Rank | 1 - SimilarWeb Rank | *Last Updated: February 1, 2020. The Best Search Engines* | eBizMBA

2 | Bing

33 - eBizMBA Rank | **500,000,000** - Estimated Unique Monthly Visitors | 8 - Quantcast Rank | 40 - Alexa Rank | 43 - SimilarWeb Rank | *Last Updated: February 1, 2020. The Best Search Engines* | eBizMBA

3 | Yahoo! Search

43 - eBizMBA Rank | **490,000,000** - Estimated Unique Monthly Visitors | 8 - Quantcast Rank | *56* - Alexa Rank | *67* - SimilarWeb Rank | *Last Updated: February 1, 2020. The Best Search Engines* | eBizMBA

4 | Baidu

54 - eBizMBA Rank | **480,000,000** - Estimated Unique Monthly Visitors | *150* - Quantcast Rank | 4 - Alexa Rank | 9 - SimilarWeb Rank | *Last Updated: February 1, 2020. The Best Search Engines* | eBizMBA


5 | Ask

205 - eBizMBA Rank | **300,000,000** - Estimated Unique Monthly Visitors | 329 - Quantcast Rank | 110 - Alexa Rank | 177 - SimilarWeb Rank | *Last Updated: February 1, 2020. The Best Search Engines* | eBizMBA

6 | AOL Search

273 - eBizMBA Rank | **200,000,000** - Estimated Unique Monthly Visitors | *350* - Quantcast Rank | 276 - Alexa Rank | *194* - SimilarWeb Rank | *Last Updated: February 1, 2020. The Best Search Engines* | eBizMBA







How can I find in which
retrieved web pages that include
"Computer Science"?

Microsoft Internet Explorer window showing a Google search for "Computer Science". The search results page displays links to various universities and their computer science departments, including Tsinghua University, Nanjing University, Peking University, and Shanghai Jiaotong University.

Right-click context menu is open on the "Computer Science" search bar in the browser.

Background image shows a computer desktop with a blue and white wavy pattern, a Microsoft Word document window, and a "畅言" (Changyan) communication software window.

Changyan software window showing a list of messages and a video preview.

Information Retrieval

Information Retrieval

- Information Retrieval (IR) is **finding material** (usually documents) of an **unstructured** nature (usually text) that satisfies an **information need** from within **large collections** (usually stored on computers).

Information Retrieval

- Information Retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections (usually stored on computers).
- These days we frequently think first of web search, but there are many other cases:
 - E-mail search
 - Searching your laptop
 - Corporate knowledge bases
 - Legal information retrieval

Basic assumptions of Information Retrieval

Basic assumptions of Information Retrieval


- **Collection:** A set of documents
 - Assume it is a static collection for the moment

Basic assumptions of Information Retrieval


- **Collection:** A set of documents
 - Assume it is a static collection for the moment
- **Goal:** Retrieve documents with information that is **relevant** to the user's **information need** and helps the user complete a **task**

This is different from searching for a fixed key in a data structure or database


The classic search model




The classic search model




The classic search model




The classic search model



The classic search model



The classic search model





Solution 1: Scan each page for the string "Computer Science".



Solution 1: Scan each page for the string "Computer Science".



Wait till your next life !




Solution 1: Scan each page for the string "Computer Science".



Wait till your next life !

How did Google do?





Solution 1: Scan each page for the string "Computer Science".



Have more than 1 billion web pages Indexed

Google 搜索

手气不错

[高级搜索](#)
[使用偏好](#)
[语言工具](#)



Solution 1: Scan each page for the string "Computer Science".



Wait till your next life !



Have more than **1 trillion** web pages Indexed

Google 搜索

手气不错

[高级搜索](#)
[使用偏好](#)
[语言工具](#)



Solution 2: Term-Document Incidence Matrix



Solution 2: Term-Document Incidence Matrix

【Example】 Document sets

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck



Solution 2: Term-Document Incidence Matrix

〔Example〕 Document sets

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck

	1	2	3	4
a	0	1	1	1
arrived	0	0	1	1
damaged	0	1	0	0
delivery	0	0	1	0
fire	0	1	0	0
gold	1	1	0	1
of	0	1	1	1
in	0	1	1	1
shipment	0	1	0	1
silver	1	0	1	0
truck	1	0	1	1



Solution 2: Term-Document Incidence Matrix

〔Example〕 Document sets

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck

	1	2	3	4
a	0	1	1	1
arrived	0	0	1	1
damaged	0	1	0	0
delivery	0	0	1	0
fire	0	1	0	0
gold	1	1	0	1
of	0	1	1	1
in	0	1	1	1
shipment	0	1	0	1
silver	1	0	1	0
truck	1	0	1	1

silver & truck



Solution 2: Term-Document Incidence Matrix

〔Example〕 Document sets

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck

	1	2	3	4
a	0	1	1	1
arrived	0	0	1	1
damaged	0	1	0	0
delivery	0	0	1	0
fire	0	1	0	0
gold	1	1	0	1
of	0	1	1	1
in	0	1	1	1
shipment	0	1	0	1
silver	1	0	1	0
truck	1	0	1	1

silver & truck = 1010 & 1011 = 1010



Solution 2: Term-Document Incidence Matrix

〔Example〕 Document sets

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck

	1	2	3	4
a	0	1	1	1
arrived	0	0	1	1
damaged	0	1	0	0
delivery	0	0	1	0
fire	0	1	0	0
gold	1	1	0	1
of	0	1	1	1
in	0	1	1	1
shipment	0	1	0	1
silver	1	0	1	0
truck	1	0	1	1

silver & truck = 1010 & 1011 = 1010



Solution 3: Compact Version - Inverted File Index



Solution 3: Compact Version - Inverted File Index

[Definition] **Index** is a mechanism for locating a given term in a text.



Solution 3: Compact Version - Inverted File Index

[Definition] **Index** is a mechanism for locating a given term in a text.

[Definition] **Inverted file** contains a list of pointers (e.g. the number of a page) to all occurrences of that term in the text.



Solution 3: Compact Version - Inverted File Index

[Definition] **Index** is a mechanism for locating a given term in a text.

[Definition] **Inverted file** contains a list of pointers (e.g. the number of a page) to all occurrences of that term in the text.

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck



Solution 3: Compact Version - Inverted File Index

[Definition] **Index** is a mechanism for locating a given term in a text.

[Definition] **Inverted file** contains a list of pointers (e.g. the number of a page) to all occurrences of that term in the text.

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck

Inverted
File
Index



No.	Term	Times; Documents
1	a	<3; 2,3,4>
2	arrived	<2; 3,4>
3	damaged	<1; 2>
4	delivery	<1; 3>
5	fire	<1; 2>
6	gold	<3; 1,2,4>
7	of	<3; 2,3,4>
8	in	<3; 2,3,4>
9	shipment	<2; 2,4>
10	silver	<2; 1,3>
11	truck	<3; 1,3,4>



Solution 3: Compact Version - Inverted File Index

[Definition] **Index** is a mechanism for locating a given term in a text.

[Definition] **Inverted file** contains a list of pointers (e.g. the number of a page) to all occurrences of that term in the text.

Doc **Text**

No. **Term** **Times; Documents**

Index

5 **fire** **<1; 2>**
6 **gold** **<3; 1,2,4>**
7 **of** **<3; 2,3,4>**
8 **in** **<3; 2,3,4>**
9 **shipment** **<2; 2,4>**
10 **silver** **<2; 1,3>**
11 **truck** **<3; 1,3,4>**

Inverted because it lists for a *term*,
all documents that contain the term

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck



No.	Term	Times; Documents
1	a	<3; 2,3,4>
2	arrived	<2; 3,4>
3	damaged	<1; 2>
4	delivery	<1; 3>
5	fire	<1; 2>
6	gold	<3; 1,2,4>
7	of	<3; 2,3,4>
8	in	<3; 2,3,4>
9	shipment	<2; 2,4>
10	silver	<2; 1,3>
11	truck	<3; 1,3,4>

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck



No.	Term	Times; Documents
1	a	<3; 2,3,4>
2	arrived	<2; 3,4>
3	damaged	<1; 2>
4	delivery	<1; 3>
5	fire	<1; 2>
6	gold	<3; 1,2,4>
7	of	<3; 2,3,4>
8	in	<3; 2,3,4>
9	shipment	<2; 2,4>
10	silver	<2; 1,3>
11	truck	<3; 1,3,4>



How to easily print the sentences which contain the words and highlight the words?

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck



No.	Term	Times; Documents Words
1	a	<3; (2;6),(3;6),(4;6)>
2	arrived	<2; (3;4),(4;4)>
3	damaged	<1; (2;4)>
4	delivery	<1; (3;1)>
5	fire	<1; (2;7)>
6	gold	<3; (1;1),(2;3),(4;3)>
7	of	<3; (2;2),(3;2),(4;2)>
8	in	<3; (2;5),(3;5),(4;5)>
9	shipment	<2; (2;1),(4;1)>
10	silver	<2; (1;2),(3;3,7)>
11	truck	<3; (1;3),(3;8),(4;7)>



How to easily print the sentences which contain the words and highlight the words?

Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck



No.	Term	Times; Documents Words
1	a	<3; (2;6),(3;6),(4;6)>
2	arrived	<2; (3;4),(4;4)>
3	damaged	<1; (2;4)>
4	delivery	<1; (3;1)>
5	fire	<1; (2;7)>
6	gold	<3; (1;1),(2;3),(4;3)>
7	of	<3; (2;2),(3;2),(4;2)>
8	in	<3; (2;5),(3;5),(4;5)>
9	shipment	<2; (2;1),(4;1)>
10	silver	<2; (1;2),(3;3,7)>
11	truck	<3; (1;3),(3;8),(4;7)>

Term
Dictionary

How to easily print the sentences which contain the words and highlight the words?



Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck



No.	Term	Times; Documents Words
1	a	<3; (2;6),(3;6),(4;6)>
2	arrived	<2; (3;4),(4;4)>
3	damaged	<1; (2;4)>
4	delivery	<1; (3;1)>
5	fire	<1; (2;7)>
6	gold	<3; (1;1),(2;3),(4;3)>
7	of	<3; (2;2),(3;2),(4;2)>
8	in	<3; (2;5),(3;5),(4;5)>
9	shipment	<2; (2;1),(4;1)>
10	silver	<2; (1;2),(3;3,7)>
11	truck	<3; (1;3),(3;8),(4;7)>

Term
Dictionary Posting List

How to easily print the sentences which contain the words and highlight the words?



Doc	Text
1	Gold silver truck
2	Shipment of gold damaged in a fire
3	Delivery of silver arrived in a silver truck
4	Shipment of gold arrived in a truck



No.	Term	Times; Documents Words
1	a	<3; (2;6),(3;6),(4;6)>
2	arrived	<2; (3;4),(4;4)>
3	damaged	<1; (2;4)>
4	delivery	<1; (3;1)>
5	fire	<1; (2;7)>
6	gold	<3; (1;1),(2;3),(4;3)>
7	of	<3; (2;2),(3;2),(4;2)>
8	in	<3; (2;5),(3;5),(4;5)>
9	shipment	<2; (2;1),(4;1)>
10	silver	<2; (1;2),(3;3,7)>
11	truck	<3; (1;3),(3;8),(4;7)>

Term
Dictionary Posting List

How to easily print the sentences which contain the words and highlight the words?



Why do we keep “times” (frequency)?

Index Generator

Index Generator

```
while ( read a document D ) {  
    while ( read a term T in D ) {  
        if ( Find( Dictionary, T ) == false )  
            Insert( Dictionary, T );  
        Get T's posting list;  
        Insert a node to T's posting list;  
    }  
}  
Write the inverted index to disk;
```

Index Generator

```
while ( read a document D ) {
    while ( read a term T in D ) {
        if ( Find( Dictionary, T ) == false )
            Insert( Dictionary, T );
        Get T's posting list;
        Insert a node to T's posting list;
    }
}
Write the inverted index to disk;
```

Index Generator

```
while ( read a document D ) {
    while ( read a term T in D ) {
        if ( Find( Dictionary, T ) == false )
            Insert( Dictionary, T );
        Get T's posting list;
        Insert a node to T's posting list;
    }
}
Write the inverted index to disk;
```

Token Analyzer
Stop Filter

Index Generator

```
while ( read a document D ) {
    while ( read a term T in D ) {
        if ( Find( Dictionary, T ) == false )
            Insert( Dictionary, T );
        Get T's posting list;
        Insert a node to T's posting list;
    }
}
Write the inverted index to disk;
```

Token Analyzer
Stop Filter

Index Generator

```
while ( read a document D ) {
    while ( read a term T in D ) {
        if ( Find( Dictionary, T ) == false )
            Insert( Dictionary, T );
        Get T's posting list;
        Insert a node to T's posting list;
    }
}
Write the inverted index to disk;
```

Token Analyzer
Stop Filter

Vocabulary
Scanner

Index Generator

```
while ( read a document D ) {
    while ( read a term T in D ) {
        if ( Find( Dictionary, T ) == false )
            Insert( Dictionary, T );
        Get T's posting list;
        Insert a node to T's posting list;
    }
}
Write the inverted index to disk;
```

Token Analyzer
Stop Filter

Vocabulary
Scanner

Index Generator

```
while ( read a document D ) {  
    while ( read a term T in D ) {  
        if ( Find( Dictionary, T ) == false )  
            Insert( Dictionary, T );  
        Get T's posting list;  
        Insert a node to T's posting list;  
    }  
}  
Write the inverted index to disk;
```

Token Analyzer
Stop Filter

Vocabulary
Scanner

Vocabulary
Insertor

Index Generator

```
while ( read a document D ) {  
    while ( read a term T in D ) {  
        if ( Find( Dictionary, T ) == false )  
            Insert( Dictionary, T );  
        Get T's posting list;  
        Insert a node to T's posting list;  
    }  
}  
Write the inverted index to disk;
```

Token Analyzer
Stop Filter

Vocabulary
Scanner

Vocabulary
Insertor

Index Generator

```
while ( read a document D ) {  
    while ( read a term T in D ) {  
        if ( Find( Dictionary, T ) == false )  
            Insert( Dictionary, T );  
        Get T's posting list;  
        Insert a node to T's posting list;  
    }  
}  
Write the inverted index to disk;
```

Token Analyzer
Stop Filter

Vocabulary
Scanner

Vocabulary
Insertor

Memory management

While reading a term

While reading a term

❖ *Word Stemming*

While reading a term

❖ *Word Stemming*

Process a word so that only its stem or root form is left.

While reading a term

❖ *Word Stemming*

Process a word so that only its stem or root form is left.

【Example】



While reading a term

❖ *Word Stemming*

Process a word so that only its stem or root form is left.

【Example】



❖ *Stop Words*

While reading a term

❖ *Word Stemming*

Process a word so that only its stem or root form is left.

【Example】

Process	}	process	}	say
processing				
processes				
processed				

❖ *Stop Words*

Some words are so common that almost every document contains them, such as “a” “the” “it”. It is useless to index them. They are called *stop words*. We can eliminate them from the original documents.

While accessing a term

While accessing a term



Solution 1: Search trees (*B- trees, B+ trees, Tries, ...*)

While accessing a term



Solution 1: Search trees (*B- trees, B+ trees, Tries, ...*)



Solution 2: Hashing

While accessing a term



Solution 1: Search trees (*B- trees, B+ trees, Tries, ...*)



Solution 2: Hashing

Discussion 3:

What are the pros and cons of using hashing, comparing to using search trees?

While accessing a term



Solution 1: Search trees (*B- trees, B+ trees, Tries, ...*)



Solution 2: Hashing

Discussion 3:

What are the pros and cons of using hashing, comparing to using search trees?

- 👉 faster for one word
- 👉 scanning in sequential order is not possible (e.g. range searches are expensive)

While not having enough memory

While not having enough memory

```
while ( read a document D ) {  
    while ( read a term T in D ) {  
        if ( Find( Dictionary, T ) == false )  
            Insert( Dictionary, T );  
        Get T's posting list;  
        Insert a node to T's posting list;  
    }  
}  
for ( i=0; i<BlockCnt; i++ )
```

While not having enough memory


```
BlockCnt = 0;
while ( read a document D ) {
    while ( read a term T in D ) {
        if ( out of memory ) {
            Write BlockIndex[BlockCnt] to disk;
            BlockCnt++;
            FreeMemory;
        }
        if ( Find( Dictionary, T ) == false )
            Insert( Dictionary, T );
        Get T's posting list;
        Insert a node to T's posting list;
    }
}
for ( i=0; i<BlockCnt; i++ )
```

While not having enough memory

```
BlockCnt = 0;
while ( read a document D ) {
    while ( read a term T in D ) {
        if ( out of memory ) {
            Write BlockIndex[BlockCnt] to disk;
            BlockCnt++;
            FreeMemory;
        }
        if ( Find( Dictionary, T ) == false )
            Insert( Dictionary, T );
        Get T's posting list;
        Insert a node to T's posting list;
    }
}
for ( i=0; i<BlockCnt; i++ )
    Merge( InvertedIndex, BlockIndex[i] );
```

While not having enough memory

```
BlockCnt = 0;
while ( read a document D ) {
    while ( read a term T in D ) {
        if ( out of memory ) {
            Write BlockIndex[BlockCnt] to disk;
            BlockCnt++;
            FreeMemory;
        }
        if ( Find( Dictionary, T ) == false )
            Insert( Dictionary, T );
        Get T's posting list;
        Insert a node to T's posting list;
    }
}
for ( i=0; i<BlockCnt; i++ )
    Merge( InvertedIndex, BlockIndex[i] );
```



Sorted

Distributed indexing (for web-scale indexing — don't try this at home!)

Distributed indexing (for web-scale indexing — don't try this at home!)

— **Each node contains index of a subset of collection**

Distributed indexing (for web-scale indexing — don't try this at home!)

— **Each node contains index of a subset of collection**

Distributed indexing (for web-scale indexing — don't try this at home!)

— Each node contains index of a subset of collection



Solution 1: Term-partitioned index

Distributed indexing (for web-scale indexing — don't try this at home!)

— Each node contains index of a subset of collection



Solution 1: Term-partitioned index



A~C



D~F

.....



X~Z

Distributed indexing (for web-scale indexing — don't try this at home!)

— Each node contains index of a subset of collection



Solution 1: Term-partitioned index



A~C



D~F

.....



X~Z



Solution 2: Document-partitioned index

Distributed indexing (for web-scale indexing — don't try this at home!)


— Each node contains index of a subset of collection



Solution 1: Term-partitioned index



Solution 2: Document-partitioned index



Dynamic indexing

Dynamic indexing

- 👉 **Docs come in over time**
 - postings updates for terms already in dictionary
 - new terms added to dictionary
- 👉 **Docs get deleted**

Dynamic indexing

👉 **Docs come in over time**


- postings updates for terms already in dictionary
- new terms added to dictionary

👉 **Docs get deleted**



Dynamic indexing

- 👉 **Docs come in over time**
 - postings updates for terms already in dictionary
 - new terms added to dictionary
- 👉 **Docs get deleted**





Dynamic indexing

👉 Docs come in over time

- postings updates for terms already in dictionary
- new terms added to dictionary


👉 Docs get deleted



Dynamic indexing

- 👉 **Docs come in over time**
 - postings updates for terms already in dictionary
 - new terms added to dictionary

- 👉 **Docs get deleted**




Dynamic indexing

👉 Docs come in over time

- postings updates for terms already in dictionary
- new terms added to dictionary

👉 Docs get deleted




Dynamic indexing

👉 Docs come in over time

- postings updates for terms already in dictionary
- new terms added to dictionary

👉 Docs get deleted



Very short answer (for exam-style):

1. When to re-index?

- When the auxiliary index grows beyond a size/segment threshold or at scheduled low-traffic times, often also when there are many deletions. Then merge auxiliary and main index (and possibly rebuild).

2. How to delete a document?

- Use **lazy deletion**: mark the document as deleted in a separate deletion list / bit-vector so queries ignore it, and **physically remove** its postings only when you rebuild/merge the index.

Compression

Compression

Term
a
arrive
damage
deliver
fire
gold
of
in
shipment
silver
truck


Compression

Term
a
arrive
damage
deliver
fire
gold
of
in
shipment
silver
truck

Compression




Compression




Compression




Compression




Compression




Compression




Compression



Compression



Compression



☞ Most gaps can be encoded with far fewer than 20 bits

Thresholding

Thresholding

阈值 != 阈值

Thresholding

 **Document:** only retrieve the top x documents where the documents are ranked by weight

Thresholding

👉 **Document:** only retrieve the top x documents where the documents are ranked by weight



Not feasible for Boolean queries



Can miss some relevant documents due to truncation

Thresholding

👉 **Document:** only retrieve the top x documents where the documents are ranked by weight

👎 Not feasible for Boolean queries

👎 Can miss some relevant documents due to truncation

👉 **Query:** Sort the query terms by their frequency in ascending order; search according to only some percentage of the original query terms

Thresholding

👉 **Document:** only retrieve the top x documents where the documents are ranked by weight



Not feasible for Boolean queries



Can miss some relevant documents due to truncation

👉 **Query:** Sort the query terms by their frequency in ascending order; search according to only some percentage of the original query terms

T1	T2	T3	T4	T5	T6	T7	T8	T9	T10
----	----	----	----	----	----	----	----	----	-----

Thresholding

👉 **Document:** only retrieve the top x documents where the documents are ranked by weight




Not feasible for Boolean queries



Can miss some relevant documents due to truncation

👉 **Query:** Sort the query terms by their frequency in ascending order; search according to only some percentage of the original query terms



Thresholding

👉 **Document:** only retrieve the top x documents where the documents are ranked by weight




Not feasible for Boolean queries



Can miss some relevant documents due to truncation

👉 **Query:** Sort the query terms by their frequency in ascending order; search according to only some percentage of the original query terms



Thresholding

👉 **Document:** only retrieve the top x documents where the documents are ranked by weight




Not feasible for Boolean queries



Can miss some relevant documents due to truncation

👉 **Query:** Sort the query terms by their frequency in ascending order; search according to only some percentage of the original query terms



Measures for a search engine

Measures for a search engine

- 👉 **How fast does it index**

- Number of documents/hour

- 👉 **How fast does it search**

- Latency as a function of index size

- 👉 **Expressiveness of query language**

- Ability to express complex information needs
 - Speed on complex queries

Measures for a search engine

- 👉 **How fast does it index**

- Number of documents/hour

- 👉 **How fast does it search**

- Latency as a function of index size

- 👉 **Expressiveness of query language**

- Ability to express complex information needs
 - Speed on complex queries



User happiness

Measures for a search engine

- 👉 **How fast does it index**

- Number of documents/hour

- 👉 **How fast does it search**

- Latency as a function of index size

- 👉 **Expressiveness of query language**

- Ability to express complex information needs
 - Speed on complex queries



User happiness ?

Measures for a search engine

- 👉 **How fast does it index**
 - Number of documents/hour
- 👉 **How fast does it search**
 - Latency as a function of index size
- 👉 **Expressiveness of query language**
 - Ability to express complex information needs
 - Speed on complex queries

👉 **User happiness ?**

- **Data Retrieval Performance Evaluation (after establishing correctness)**
 - > Response time
 - > Index space
- **Information Retrieval Performance Evaluation**
 - > + How *relevant* is the answer set?

Relevance measurement requires 3 elements:

1. A benchmark **document** collection
2. A benchmark suite of **queries**
3. A binary **assessment** of either Relevant or Irrelevant for each query-doc pair

Relevance measurement requires 3 elements:

1. A benchmark **document** collection
2. A benchmark suite of **queries**
3. A binary **assessment** of either **Relevant** or **Irrelevant** for each query-doc pair

	Relevant	Irrelevant
Retrieved	R_R	I_R
Not Retrieved	R_N	I_N

Relevance measurement requires 3 elements:

1. A benchmark **document** collection
2. A benchmark suite of **queries**
3. A binary **assessment** of either **Relevant** or **Irrelevant** for each query-doc pair

	Relevant	Irrelevant
Retrieved	R_R	I_R
Not Retrieved	R_N	I_N

$$\text{Precision } P = R_R / (R_R + I_R)$$

$$\text{Recall } R = R_R / (R_R + R_N)$$


Relevance measurement requires 3 elements:

1. A benchmark **document** collection
2. A benchmark suite of **queries**
3. A binary **assessment** of either Relevant or Irrelevant for each query-doc pair

	Relevant	Irrelevant
Retrieved	R_R	I_R
Not Retrieved	R_N	I_N

$$\text{Precision } P = R_R / (R_R + I_R)$$

$$\text{Recall } R = R_R / (R_R + R_N)$$




Relevance measurement requires 3 elements:

1. A benchmark **document** collection
2. A benchmark suite of **queries**
3. A binary **assessment** of either Relevant or Irrelevant for each query-doc pair

	Relevant	Irrelevant
Retrieved	R_R	I_R
Not Retrieved	R_N	I_N

$$\text{Precision } P = R_R / (R_R + I_R)$$

$$\text{Recall } R = R_R / (R_R + R_N)$$




Relevance measurement requires 3 elements:


1. A benchmark **document** collection
2. A benchmark suite of **queries**
3. A binary **assessment** of either Relevant or Irrelevant for each query-doc pair

	Relevant	Irrelevant
Retrieved	R_R	I_R
Not Retrieved	R_N	I_N


$$\text{Precision } P = R_R / (R_R + I_R)$$

$$\text{Recall } R = R_R / (R_R + R_N)$$






Returns relevant documents but
misses many useful ones too




Returns relevant documents but misses many useful ones too



Returns most relevant documents but includes lot of junk

Returns relevant documents but misses many useful ones too



Discussion 4:

How to improve the *relevancy* of
search results?

Discussion 4:

How to improve the *relevancy* of
search results?

↳ **Page Rank**

↳ **Semantic Web**

↳

Inverted File Index

- Inverted File Index
- Take-home messages

Take-Home Messages

- Inverted file index:
 - Mapping from items to posting lists of documents.
 - Sorted, frequency of appearance, positions.
 - Performance measures: precision, recall...

Thanks for your attention!
Discussions?

Reference

Stanford CS276: <https://web.stanford.edu/class/cs276/>

Schütze, Hinrich, Christopher D. Manning, and Prabhakar Raghavan. Introduction to information retrieval. Cambridge University Press, 2008. [Chap. 1, 4, 5, 6, 8](#)

Zobel, Justin, and Alistair Moffat. "Inverted files for text search engines." ACM computing surveys, 2006.